
Lecture 12

Indexes Continued

Summary – last week

Summar

- Indexes
 - Tree based indexes

This week:


- Indexes
 - R+ Tree
 - Bitmap indexes



Improved Version of R-Tree

- Where are R-trees inefficient?
 - They **allow overlapping** between neighboring MBRs
- R⁺-Trees (Sellis et al., 1987)
 - **Overlapping** of neighboring MBRs are prohibited
 - This may lead to identical leafs occurring more than once in the tree
 - Improve search efficiency, but similar scalability as R-trees

R⁺-Tree



- Overlaps are not permitted (A and B)
- Data rectangles are divided and may be present (e.g., G) in several leafs

Performance

Advantages

- The main advantage of R⁺-trees is to improve the search performance
- Especially for point queries, this saves 50% of access time

Disadvantages

- Drawback is the low occupancy of nodes resulting through many splits
- R⁺-trees often degenerate with the increasing number of changes

Bitmap Indexes

- Bitmap Indexes
 - Lets assume a relation Expenses with three attributes: Nr, Shop and Sum
 - A bitmap index for attribute Shop looks like this

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Value	Vector
P&C	001000
Real	010101
Saturn	100010

Bitmap Indexes (cont'd.)

- Handling modification
 - Assume record numbers are not changed
- Deletion
 - **Tombstone** replaces deleted record (6 doesn't become 5!)
 - Corresponding bit is set to 0

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Before

Value	Vector
P&C	001000
Real	010101
Saturn	100010



After

Value	Vector
P&C	001000
Real	010101
Saturn	100000

Bitmap Indexes (cont'd.)

- Inserted record is assigned the next record number
 - A bit of value 0 or 1 is appended to each bit vector
 - If new record contains a new value of the attribute, add one bit-vector
 - E.g., insert new record with REWE as shop

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80
7	REWE	23

Before		After	
Value	Vector	Value	Vector
P&C	001000	P&C	0010000
Real	010101	Real	0101010
Saturn	100010	Saturn	1000100
		REWE	0000001



Bitmap Indexes (cont'd.)

- Update

- Change the bit corresponding to the old value of the modified record to 0
- Change the bit corresponding to the new value of the modified record to 1
- If the new value is a new value of A, then insert a new bit-vector: e.g., replace Shop for record 2 to REWE

Nr	Shop	Sum
1	Saturn	150
2	REWE	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Before		After	
Value	Vector	Value	Vector
P&C	001000	P&C	001000
Real	010101	Real	000101
Saturn	100010	Saturn	100010
		REWE	010000



Bitmap Indexes (cont'd.)

- Select
 - Basic bit operations:
 - E.g., select the sums we have spent in Saturn and P&C

Saturn	OR	P&C	=	Result
1		0		1
0		0		0
0		1		1
0		0		0
1		0		1
0		0		0

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Value	Vector
P&C	001000
Real	010101
Saturn	100010

- Bitmap indexes should be used when selectivity is **high**

Bitmap Indexes (cont'd.)

- Advantages
 - Operations are efficient and easy to implement (directly supported by hardware)
- Disadvantages
 - For each new value of an attribute a new bitmap-vector is introduced
 - If we bitmap index an attribute like birthday (only day) we have 365 vectors: $365/8 \text{ bits} \approx 46 \text{ Bytes}$ for a record, just for that
 - Solution to such problems is multi-component bitmaps
 - Not fit for range queries where many bitmap vectors have to be read
 - Solution: range-encoded bitmap indexes

- Bitmap indexes
 - Great for indexing tables with set-like attributes e.g., Gender:Male/Female
 - Operations are efficient and easy to implement (directly supported by hardware)
 - Multi component reduce the storage while range encoded allow for fast range queries

Next Lecture

- Optimization
 - Partitioning
 - Joins
 - Materialized Views

