
Lecture 13

Optimizations

Summary – last week

Summary

- R+ Tree
 - Removes the overlaps in R Tree
- Bitmap indexes
 - Great for indexing tables with set-like attributes e.g., Gender:Male/Female
 - Operations are efficient and easy to implement (directly supported by hardware)



This week

Optimization

I. Partitioning



Why Partitioning?

Partitioning

- Breaking the data into several **physical units** that can be handled separately
- Granularity and partitioning are key to **efficient implementation of a warehouse**
- The question is not whether to use partitioning but how to do it



Partitioning (cont'd.)

- Why partitioning?
 - **Flexibility** in managing data
 - Smaller physical units allow
 - Easy restructuring e.g. adding weekly classification level
 - Cheap indexing
 - Sequential scans, if needed
 - Easy reorganization
 - Easy recovery
 - Easy monitoring



Partitioning (cont'd.)

- In DWs, partitioning is done to improve:
 - **Business query performance**, i.e., minimize the amount of data to scan
 - **Data availability**, e.g., back-up/restores can run at the partition level
 - **Database administration**, e.g., adding new columns to a table, archiving data, recreating indexes, loading tables

Partitioning (cont'd.)

- Possible approaches:
 - **Data partitioning**
where data is usually partitioned by
 - Date
 - Line of business
 - Geography
 - Organizational unit
 - Combinations of these factors
 - **Hardware partitioning**
 - Hardware partitioning makes data available to different processing nodes by ensuring that sub-processes are capable of running on the different nodes




Partitioning (cont'd.)

- At which **level** should data partitioning occur?

- Possibilities are


- on application
- or DBMS level



- Partitioning on DBMS level is clear, but it also makes sense to partition at **application level**
 - E.g., allows different definitions for each year
 - Important since DWs spans many years and as business evolves DWs change, too
 - Think for instance about changing tax laws

Data Partitioning

- Data partitioning, involves:
 - Splitting out the rows of a table into multiple tables i.e., **horizontal partitioning**
 - Splitting out the columns of a table into multiple tables i.e., **vertical partitioning**



Which partitioning type is good tax laws changing every year?

Data Partitioning (cont'd.)

- Horizontal partitioning
 - The set of tuples of a table is split among **disjoint** table parts
 - Definition: A set of Relations $\{R_1, \dots, R_n\}$ represent a **horizontal partitioning** of Master-Relation R if and only if $R_i \subseteq R$, $R_i \cap R_j = \emptyset$ and $R = \bigcup_i R_i$, for $1 \leq i, j \leq n$
 - According to the partitioning procedure we have different horizontal partitioning solutions
 - Range partitioning, List partitioning and Hash partitioning

Horizontal Partitioning

- Range Partitioning
 - Selects a partition by determining if the partitioning key is inside a certain **range**
 - A partition can be represented as a restriction on the master-relation
 - $R_i = \sigma_{P_i}(R)$, where P_i is the partitioning predicate. The partitioning predicate can involve more attributes
 - P_1 : Country = 'Germany' and Year = 2009
 - P_2 : Country = 'Germany' and Year < 2009
 - P_3 : Country \neq 'Germany'

Horizontal Partitioning (cont'd.)

- List Partitioning

- A partition is assigned a **list of values**. If the partitioning key has one of these values, the partition is chosen
 - For example all rows where the column Country is either Iceland, Norway, Sweden, Finland or Denmark could build a partition for the Scandinavian countries
- Is also expressed as a simple **restriction** on the master relation
 - The partitioning predicate involves just one attribute
 - P_1 : City IN ('Hamburg', 'Hannover', 'Berlin')
 - P_2 : City IN (DEFAULT) – represents tuples which do not fit to P_1

Horizontal Partitioning (cont'd.)

- Hash Partitioning
 - The value of a hash function determines membership in a partition. Assuming there are four partitions, the hash function could return a value from 0 to 3
 - For each tuple t , of the master-table R , the hash function will associate it to a partition table R_i
 - $R_i = \{t_1, \dots, t_m / t_j \in R \text{ and } H(t_j) = H(t_k) \text{ for } 1 \leq j, k \leq m\}$
 - This kind of partitioning is particularly used in parallel processing
 - The goal is to achieve an equal distribution of the data

Horizontal Partitioning (cont'd.)

- Horizontal partitioning in DataWarehousing partitions data by
 - Time dimension
 - Periods, such as week or month can be used or the data can be partitioned by the age of the data
 - E.g., if the analysis is usually done on last month's data the table could be partitioned into monthly segments
 - A dimension other than time
 - If queries usually run on a grouping of data: e.g., each branch tends to query on its own data and the dimension structure is not likely to change then partition the table on this dimension
 - On table size
 - If a dimension cannot be used, partition the table by a predefined size. If this method is used, metadata must be created to identify what is contained in each partition

Vertical Partitioning

- Vertical Partitioning
 - Involves creating tables with **fewer columns** and using additional tables to store the remaining columns
 - Different physical storage might be used e.g., storing infrequently used or very wide columns on a different device
 - Usually called **row splitting**
 - Row splitting creates a one-to-one relationship between the partitions

Vertical Partitioning (cont'd.)

- In DW, common vertical partitioning means
 - Moving **seldom used** columns from a highly-used table to another table
 - Creating a **view** across the two newly created tables restores the original table with a performance penalty
 - However, performance will increase when accessing the highly-used data e.g., for statistical analysis

Vertical Partitioning (cont'd.)

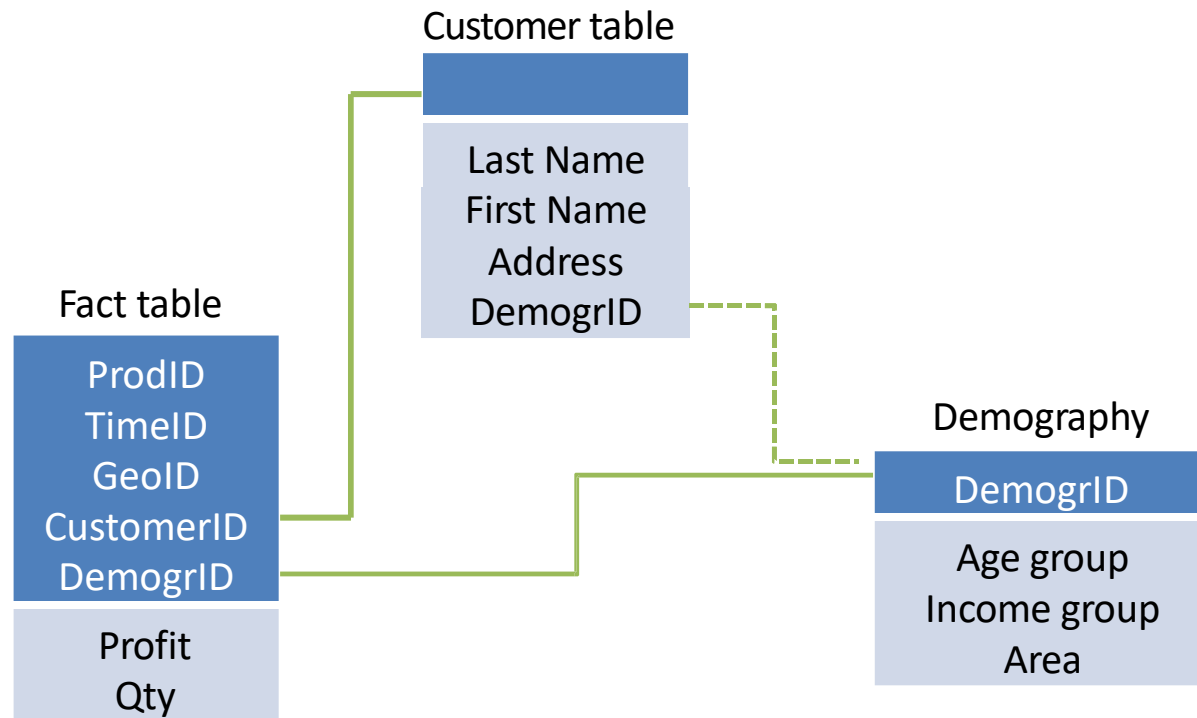
- In DWs with **very large** dimension tables (e.g., Amazon - a customer table with tens of millions of records) we have
 - Most of the attributes are rarely, if at all, queried
 - E.g., the address attribute is not as interesting as evaluating customers per age-group
 - But we must still maintain the link between the fact table and the **complete** customer dimension, which has **high performance costs!**

Vertical Partitioning (cont'd.)

- The solution is to use **Mini-Dimensions**, a special case of vertical partitioning
 - Many dimension attributes are used **very frequently** as browsing constraints
 - In big dimensions these constraints can be hard to find among the lesser used ones
 - Logical groups of often used constraints can be separated into **small dimensions** which are very well indexed and easily accessible for browsing


Vertical Partitioning (cont'd.)

- Mini-Dimensions, e.g., the **Demography** table



Vertical Partitioning (cont'd.)

- All variables in these mini-dimensions must be presented as **distinct classes**
- The key to the mini-dimension can be placed as a **foreign key** in both the fact table and dimension table from which it has been broken off
- Mini-dimensions, as their name suggests, should be kept small and compact




Partitioning

- Advantages
 - Records used together are grouped together
 - Each partition can be optimized for performance
 - Security, recovery
 - Partitions stored on different disks: contention
 - Take advantage of parallel processing capability
- Disadvantages
 - Slow retrieval across partitions (expensive joins)
 - Complexity

Partitioning (cont'd.)

- Use partitioning when:
 - A table is $> 2\text{GB}$ (from Oracle)
 - A Table is > 100 Million rows (praxis)
 - Think about it, if table is > 1 million rows
- Partitioning does not come for free!



Partitioning (cont'd.)

- Partitioning **m a n a g e m e n t**
 - Partitioning should be transparent outside the DBMS
 - The applications work with the Master-Table at logical level
 - The conversion to the physical partition tables is performed internally by the DBMS
 - It considers also data consistency as if the data were stored in just one table
 - Partitioning transparency is not yet a standard. Not all DBMS support it!

Partitioning Management

- Partitions in practice
 - Oracle supports Range-, List-, Hash-, Interval-, System-Partitions as well as combinations of these methods
 - E.g., partitioning in Oracle:
 - ```
CREATE TABLE SALES(
 ProdID NUMBER,
 GeoID NUMBER,
 TimeID DATE,
 Profit NUMBER)
PARTITION BY RANGE(timeID)(
 PARTITION before 2008
 VALUES LESS THAN (TO_DATE ('01-JAN-
2008', 'DD-MM-YYYY')),
 PARTITION 2008
 VALUES LESS THAN (TO_DATE ('01-JAN-
2009', 'DD-MM-YYYY'))
);
```





# Partitioning Management (cont'd.)

- Partitions in practice
  - In Oracle partitioning is performed with the help of the LESSTHAN function. How can we partition data in the current year?
    - ALTER TABLE Sales

ADD PARTITION after2008 VALUES LESSTHAN  
(MAXVALUE);



# Partitioning Management (cont'd.)

| RowID | ProdID | GeoID | TimeID  | Profit |
|-------|--------|-------|---------|--------|
| ...   | ...    | ...   | ...     | ...    |
| 121   | 132    | 2     | 05.2007 | 8K     |
| 122   | 12     | 2     | 08.2008 | 7K     |
| 123   | 15     | 1     | 09.2007 | 5K     |
| 124   | 14     | 3     | 01.2009 | 3K     |
| 125   | 143    | 2     | 03.2009 | 1,5K   |
| 126   | 99     | 3     | 05.2007 | 1K     |
|       | ...    | ...   | ...     | ...    |

| RowID | ProdID | GeoID | TimeID  | Profit |
|-------|--------|-------|---------|--------|
| ...   | ...    | ...   | ...     | ...    |
| 121   | 132    | 2     | 05.2007 | 8K     |
| 123   | 15     | 1     | 09.2007 | 5K     |
| 126   | 99     | 3     | 05.2007 | 1K     |

| RowID | ProdID | GeoID | TimeID  | Profit |
|-------|--------|-------|---------|--------|
| 122   | 12     | 2     | 08.2008 | 7K     |

| RowID | ProdID | GeoID | TimeID  | Profit |
|-------|--------|-------|---------|--------|
| 124   | 14     | 3     | 01.2009 | 3K     |
| 125   | 143    | 2     | 03.2009 | 1,5K   |

# Summary

*Summary*

- Partitioning: Horizontal or Vertical
  - Records used together are grouped together
  - However: slow retrieval across partitions
  - Mini-Dimensions

# Next Lecture

1. Joins
2. MaterializedViews

