
Lecture 14

Optimizations (Continued)

Summary – last week

Summary

Optimization

I. Partitioning



This week

Optimization

1. Joins
2. MaterializedViews



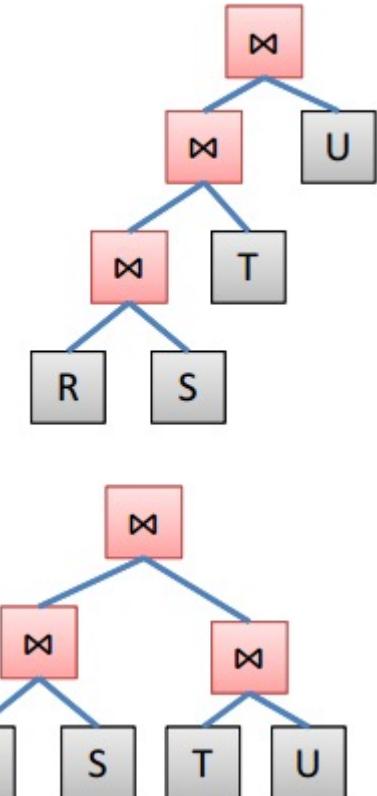
Join Optimization

- Often queries over **several** partitions are needed
 - This results in **joins** over the data
 - Though joins are **generally expensive operations**, the overall cost of the query may strongly differ with the chosen evaluation plan for the joins
- Joins are **commutative** and **associative**
 - $R \bowtie S \equiv S \bowtie R$
 - $R \bowtie (S \bowtie T) \equiv (S \bowtie R) \bowtie T$



Join Optimization (cont'd.)

- This allows to evaluate individual joins in **any order**
 - Results in **join trees**
 - Different join trees may show very different evaluation performance
 - Join trees have different **shapes**
 - Within a shape, there are different relation **assignments** possible
- Example: $R \bowtie S \bowtie T \bowtie U$



Join Optimization (cont'd.)

- Number of **possible join trees** grows rapidly with number of join relations
 - For n relations, there are $T(n)$ different tree shapes
 - $T(1) = 1$
 - $T(n) = \sum_{i=1}^{n-1} T(i)T(n-i)$
 - “Any number of $1 \leq i \leq n-1$ relations may be in the left subtree and ordered in $T(i)$ shapes while the remaining $n-i$ relations form the right subtree and can be arranged in $T(n-i)$ shapes.”

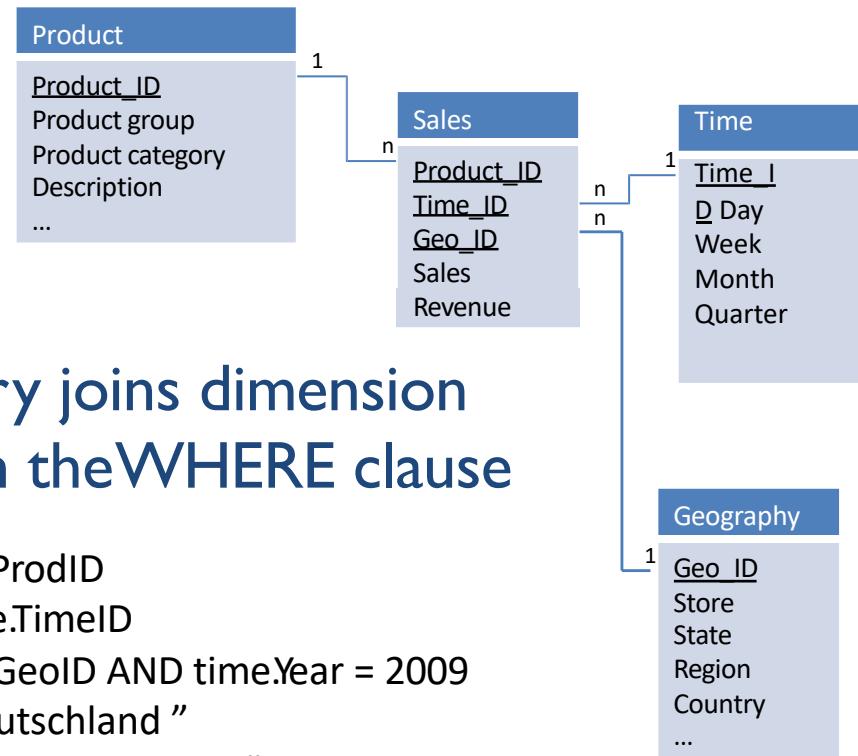
Join Optimization (cont'd.)

- Optimizer has 3 choices
 - Consider all possible join trees
 - Usually not possible
 - Consider a subset of all trees
 - i.e. restrict to trees of certain shapes
 - Use heuristics to pick a certain shape

Join Optimization in DW

- Relational optimization of **star-joins**
 - Star schema comprises a **big** fact table and many **small** dimension tables
 - An **OLAP SQL** query joins dimension and fact tables usually in the WHERE clause

```
sales.ProdID = product.ProdID  
AND sales.TimeID = time.TimeID  
AND sales.GeoID = geo.GeoID AND time.Year = 2009  
AND geo.Country = "Deutschland"  
AND product.group = "wash machines"
```

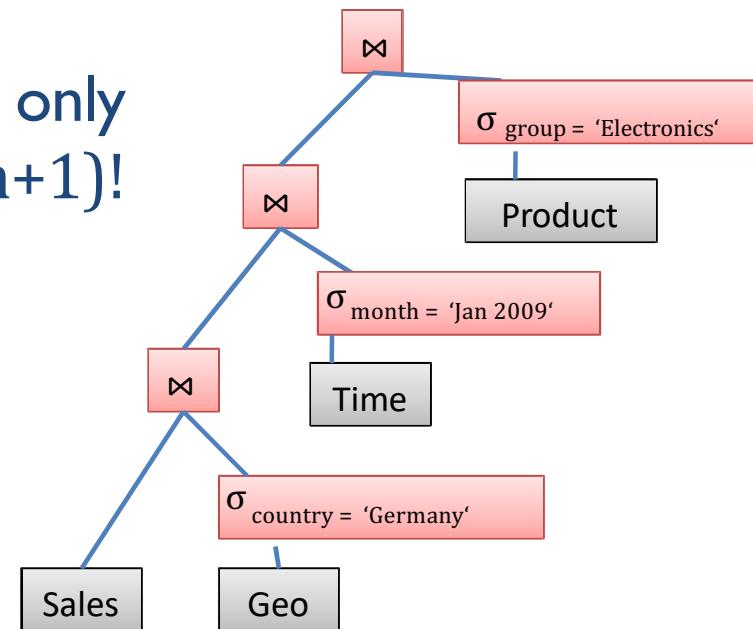


Join Optimization in DW (cont'd.)

- If the OLAP query specifies **restrictions** or **group by's** on n dimensions, an $n+1$ order join is necessary
 - Joins can be performed only pair-wise, resulting in $(n+1)!$ possible joining orders

1. $\sigma_{country=Germany} (Sales \bowtie Geo)$

2. $Sales \bowtie (\sigma_{country=Germany} Geo)$

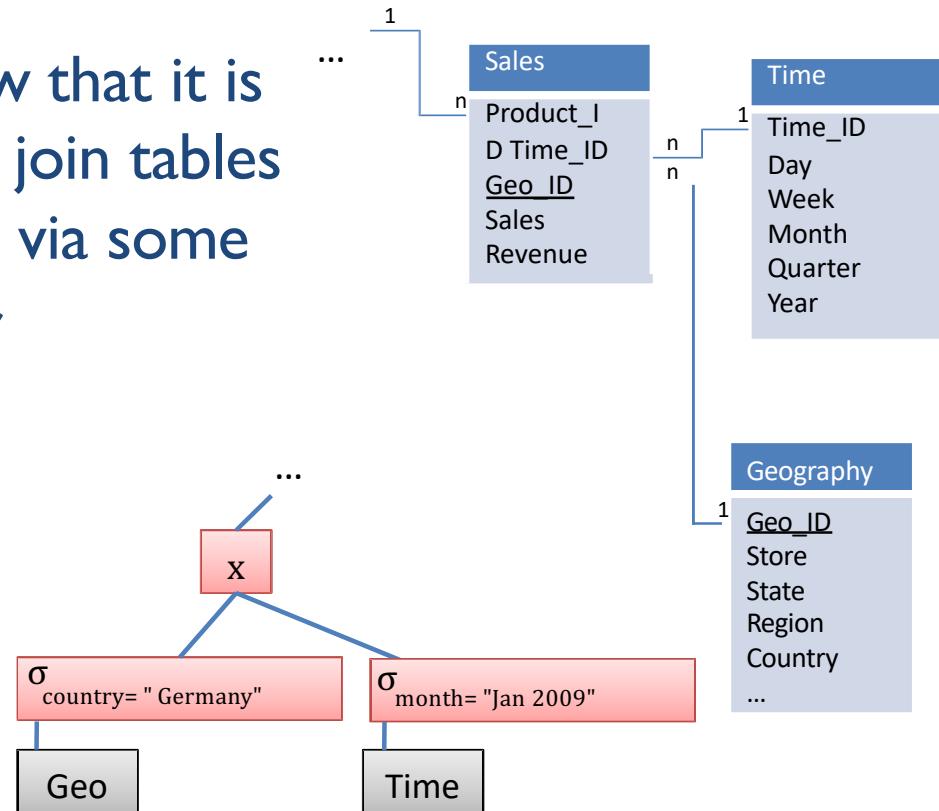


Join Heuristics

- To reduce the number of join-orders, **heuristics** are used

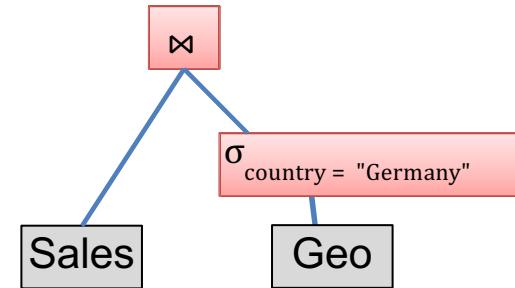
– In OLTP heuristics show that it is **not a good idea** to join tables that are not connected via some attribute to each other

- E.g., Geo with Time relation leads to a **Cartesian product**



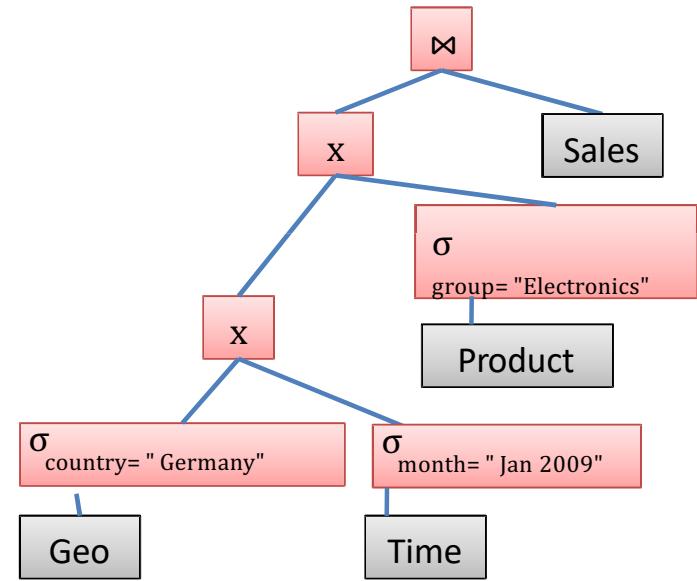
Join Heuristics (cont'd.)

- But this heuristic rule from OLTP is **not suitable** for DW!
 - E.g., join Sales with Geo in the following case:
 - Sales has 10 mil records, in Germany there are 10 stores, in January 2016 there were products sold in 20 days, and the Electronics group has 50 products
 - If 20% of our sales were performed in Germany, still we get 2mil intermediate results so index on Geo table would not help that much either.



Dimensional Cross Product

- In star-joins a **cross product** of the dimension tables is recommended
 - Geo dimension – 10 stores
 - Time dimension 20 days
 - Product dimension 50 products
 - $10 \times 20 \times 50 = 10\,000$ records after performing the cross product of the dimensions
 - The selectivity is in this case 0.1% which is fit for using an index

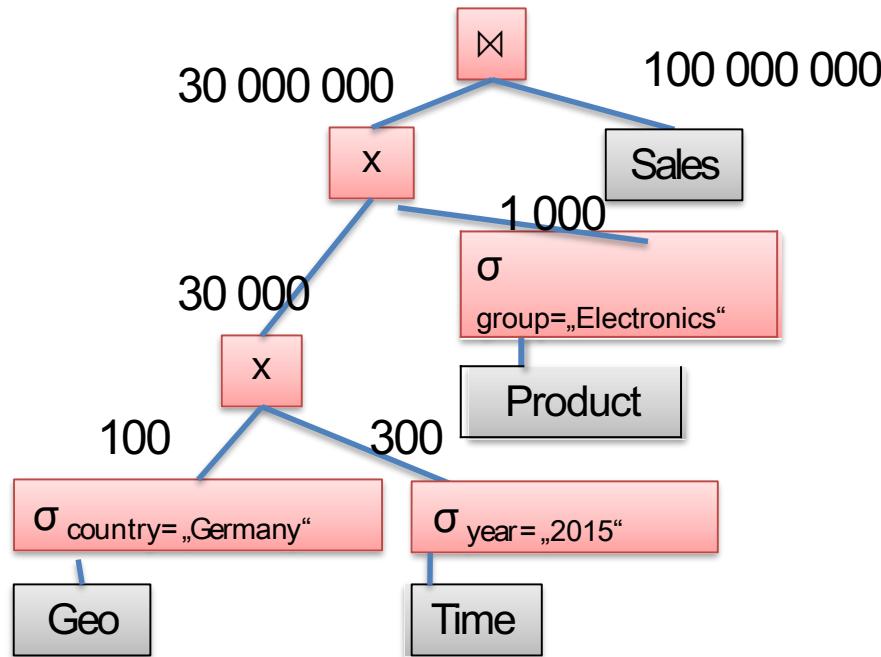


Dimensional Cross Product (cont'd.)

- But dimension cross products can also become **expensive**
 - If the restrictions on the dimensions are not restrictive enough or if there are many dimension tables
- E.g. query for the sales of all electronic products of a company in 2015:
 - The company has **100 stores** in Germany and it sells **1000 types of electronics products**
 - In 2015 it sold products in **300 working days**

Dimensional Cross Product (cont'd.)

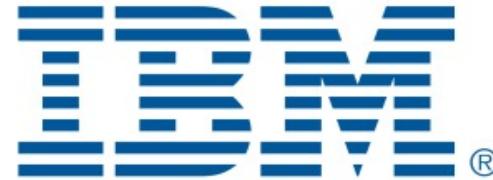
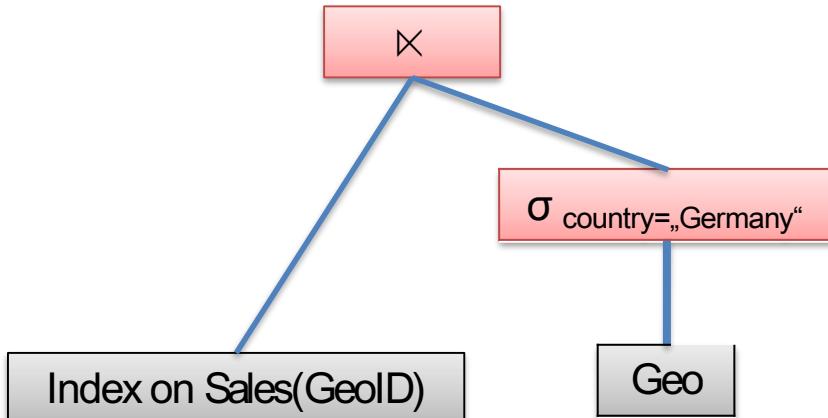
- $100 \text{ stores} * 300 \text{ days} * 1.000 \text{ products} = 30 \text{ mil records...}$



- Very **expensive** to compute

Star-join Optimization

- The **IBM DB2** solution for expensive dimension cross products is to build **semi-joins** of the dimension tables with indexes
 - A B*-Tree index will be on the fact table (**sales**) for each dimension



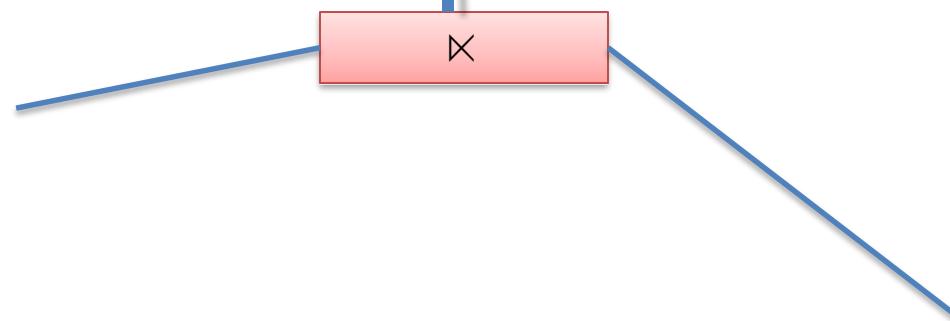
Keep all index entries for the Sales fact table for sales in Germany

Star-join Optimization (cont'd.)

Index on Sales
for GeoID

GeoID	ID
1	1
1	2
4	3
...	...

GeoID	ID
1	1
1	2



Ge

GeoID	Store	City	Country
1	S1	BS	German
2	S2	BS	Germany
3	S3	HAN	Germany
4	S4	Lyon	France

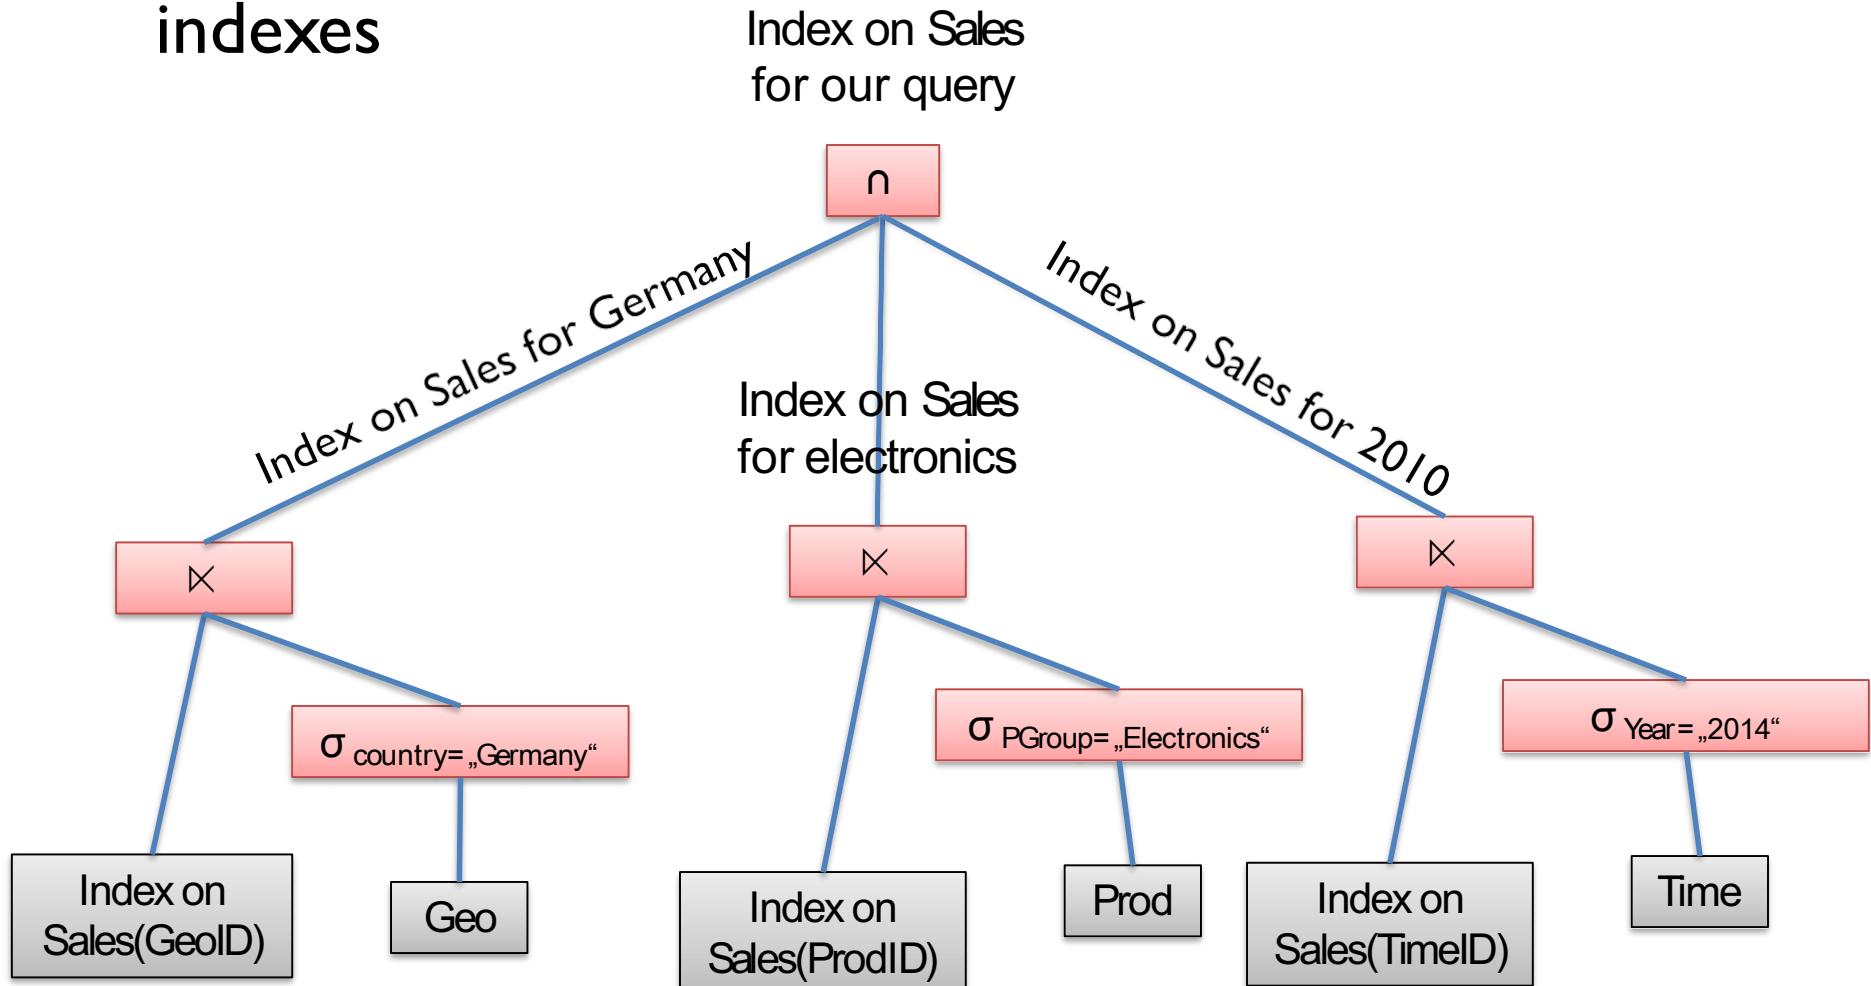
$\sigma_{\text{country}=\text{"Germany"}}$



GeoID	Store	City	Country
1	S1	BS	Germany
2	S2	BS	Germany
3	S3	HAN	Germany

Star-join Optimization (cont'd.)

- Reduces the fact table to what we need, based on indexes



Materialized Views

- MaterializedViews (MV)
 - Views whose tuples are **stored** in the database are said to be materialized
 - They provides fast access, like a (very high-level) cache
 - Need to maintain the view as the underlying tables change
 - Ideally, we want incremental view maintenance algorithms



Materialized Views (cont'd.)

- How can we use MV in DW?
 - E.g., we have queries requiring us to join the Sales table with another table and aggregate the result
 - `SELECT P.Categ, SUM(S.Qty) FROM Product P, Sales S WHERE P.ProdID=S.ProdID GROUP BY P.Categ`
 - `SELECT G.Store, SUM(S.Qty) FROM Geo G, Sales S WHERE G.Geoid=S.Geoid GROUP BY G.Store`
 -
 - There are more solutions to speed up such queries
 - Pre-compute the two joins involved (product with sales and geo with sales)
 - Pre-compute each query in its entirety
 - Or use an already materialized view

Materialized Views (cont'd.)

- Having the following view materialized
 - CREATE MATERIALIZED VIEW Totalsales (ProdID, GeoID, total) AS SELECT S.ProdID, S.GeoID, SUM(S.Qty) FROM Sales S GROUP BY S.ProdID, S.GeoID
- We can use it in our 2 queries
 - ```
SELECT P.Categ, SUM(T.Total) FROM Product P, Totalsales T WHERE P.ProdID=T.ProdID GROUP BY P.Categ
```
  - ```
SELECT G.Store, SUM(T.Total) FROM Geo G, Totalsales T WHERE G.GeoID=T.GeoID GROUP BY G.Store
```

Materialized Views (cont'd.)

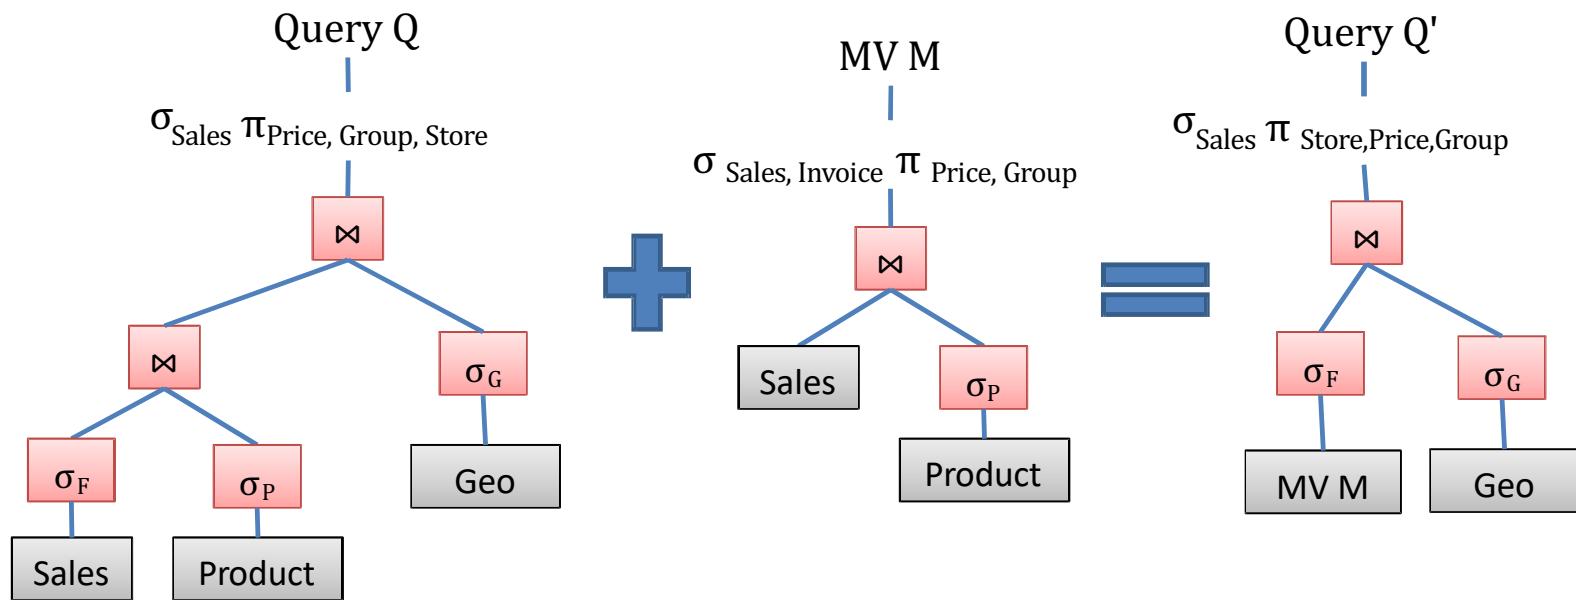
- MV issues
 - **Utilization**
 - What views should we materialize, and what indexes should we build on the pre-computed results?
 - **Choice of materialized views**
 - Given a query and a set of materialized views, can we use the materialized views to answer the query?
 - **Maintenance**
 - How frequently should we refresh materialized views to make them consistent with the underlying tables?
 - And how can we do this incrementally?

Utilization of MV

- Materialized views **utilization** has to be transparent
 - Queries are internally rewritten to use the available MVs by the **query rewriter**
 - The query rewriter performs integration of the MV based on the **query execution graph**

Utilization of MV (cont'd.)

- E.g., materialized views utilization, mono-block query



Integration of MV

- Integration of MV
 - **Valid replacement:** A query Q' represents a valid replacement of query Q by utilizing the materialized view M , if Q and Q' always deliver the same result set

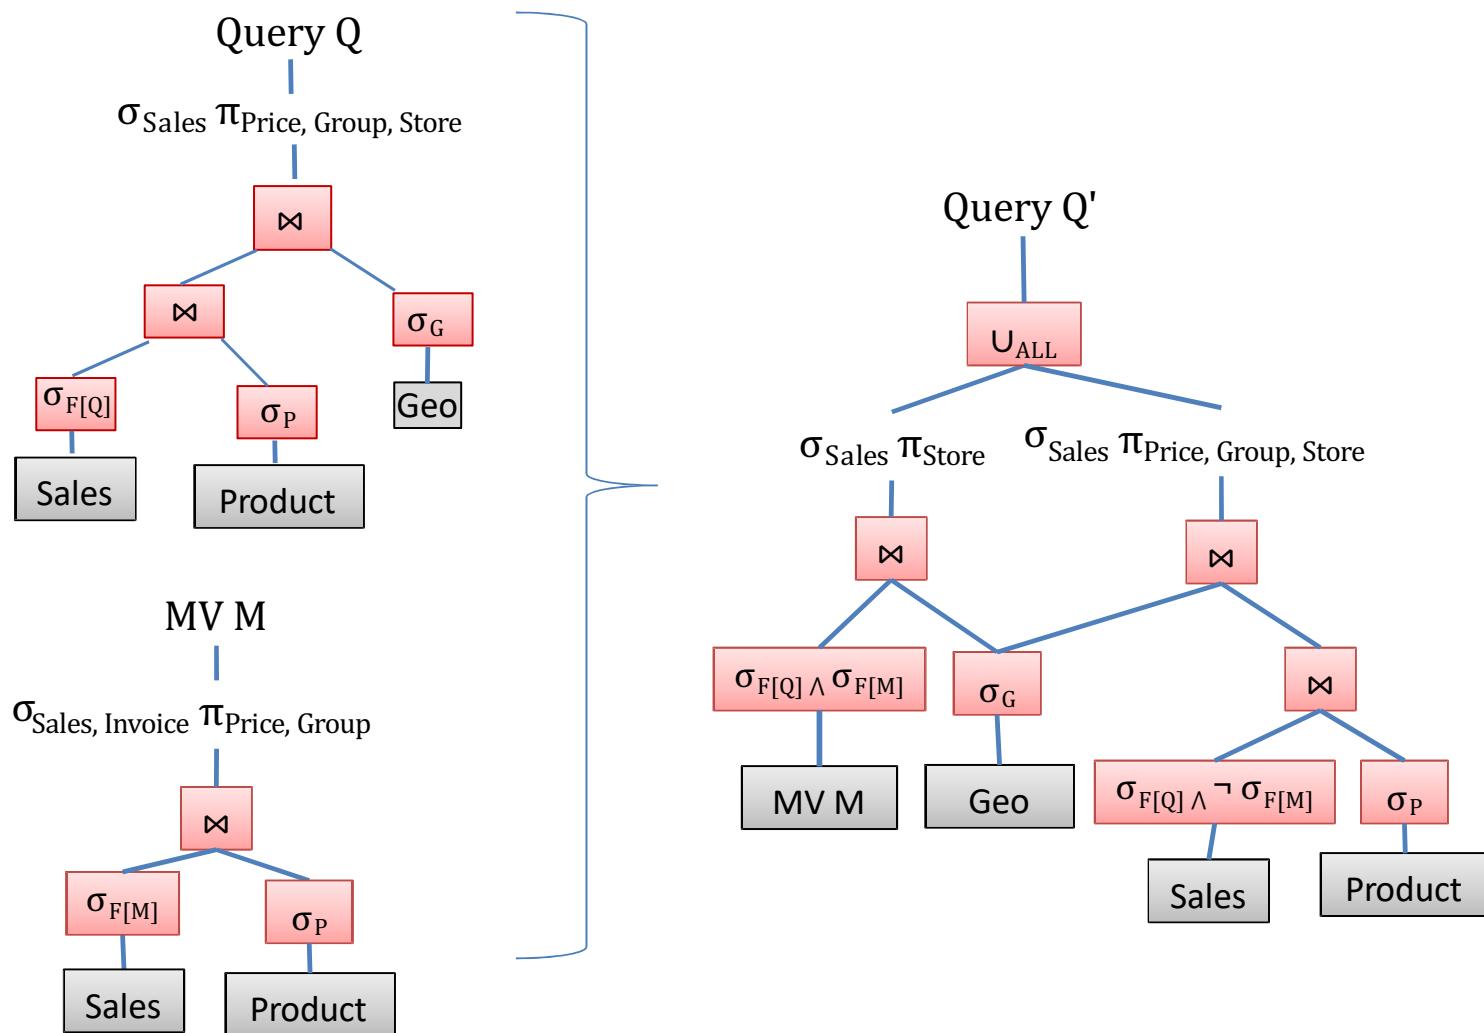
Integration of MV (cont'd.)

- In order to be able to integrate MV M in Q and obtain Q', the following **conditions** need to be respected
 - The selection condition in M cannot be more restrictive than the one in Q
 - The projection from Q has to be a subset of the projection from M
 - It has to be possible to derive the aggregation functions of $\pi(Q)$ from $\pi(M)$
 - Additional selection conditions in Q have to be possible also on M

Integration of MV (cont'd.)

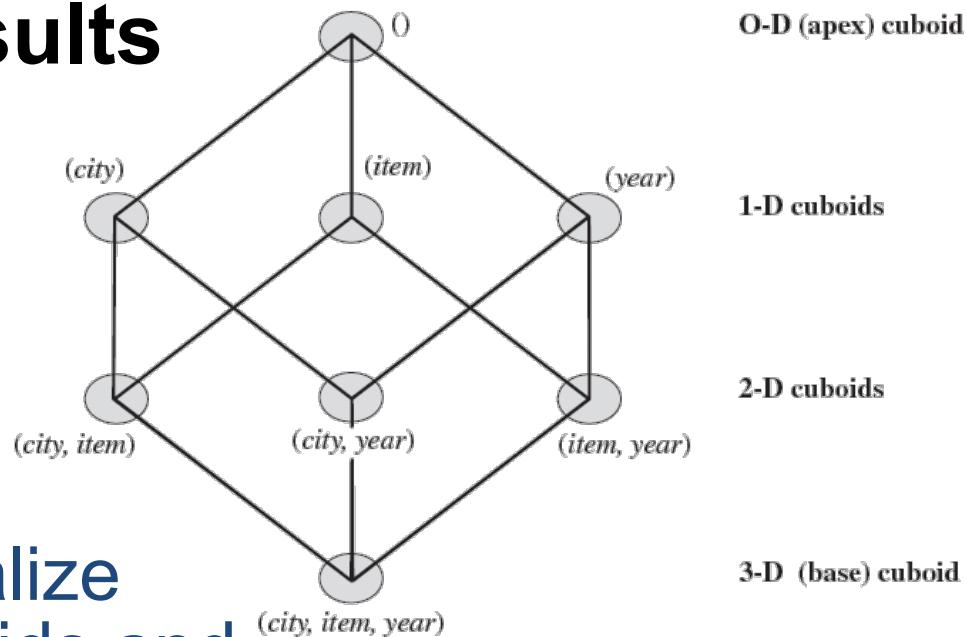
- How do we use MV even when there is no perfect match? (Multiblock queries)
- If the selection in M is more restrictive than the selection in Q
 - Split Q in Q_a and Q_b such that
 - $\sigma(Q_a) = (\sigma(Q) \wedge \sigma(M))$ and
 - $\sigma(Q_b) = (\sigma(Q) \wedge \neg \sigma(M))$

Integration of MV (cont'd.)



MVs in DWs

- In DW, materialized views are often used to store **aggregated results**
 - The number of nodes in the lattice of cuboids is
 - $|n| = \prod_{j=1}^n 2 = 2^n$
 - $n = 3$, $|n| = 8$ and we would need to materialize 2-D cuboids 1-D cuboids and 0D cuboids; in total 7 views
 - $n = 16$, $|n| = 65534$, ... too much to materialize
 - **What should we materialize?**

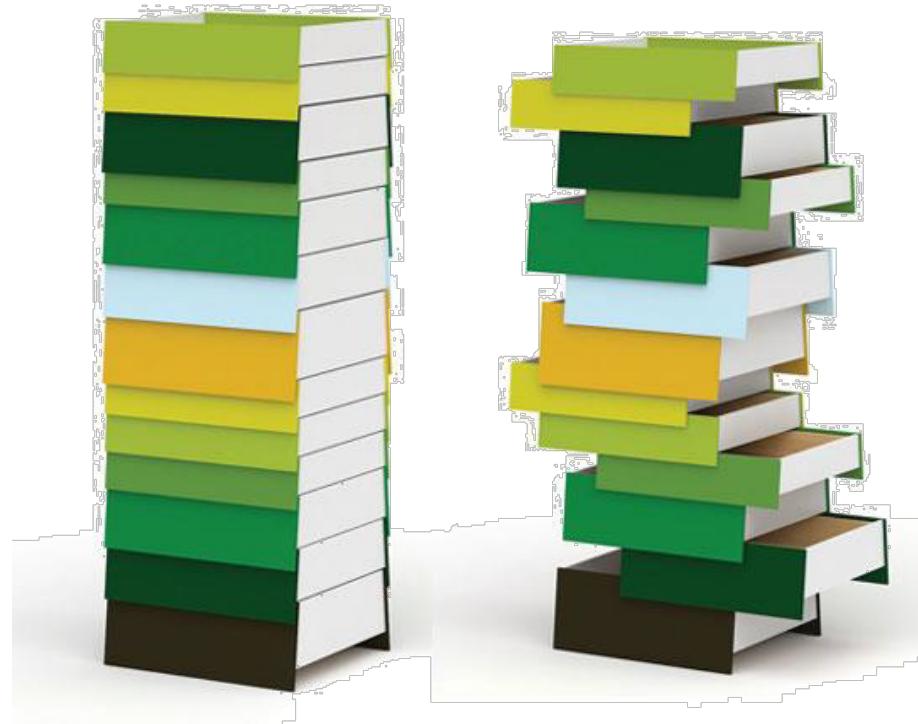


Maintenance of MV

- Maintenance of MV
 - Keeping a materialized view **up-to-date** with the underlying data
 - Important questions
 - **How** do we refresh a view when an underlying table is refreshed?
 - **When** should we refresh a view in response to a change in the underlying table?

How to Refresh a MV

- Materialized views can be maintained by re-computation on **every update**
 - – Not the best solution
- A better option is **incremental view maintenance**



When to Refresh a MV

- **Immediate**
 - As part of the transaction that modifies the underlying data tables
 - Advantage: materialized view is always consistent
 - Disadvantage: updates are slowed down
- **Deferred**
 - Some time later, in a separate transaction
 - Advantage: can scale to maintain many views without
 - slowing updates
 - Disadvantage: view briefly becomes inconsistent

Summary

Summary

- **Joins:** for DW it is sometimes better to perform cross product on dimensions first
- **MaterializedViews:** we can't materialize everything
 - Static or Dynamic choice of what to materialize
 - The benefit cost function is decisive

Next Lecture

- OLAP Queries!

