

Operating Systems

5. Synchronization

Synchronization

- Many processes/threads interact (shared memory / message passing)
 - Results of interactions not guaranteed to be deterministic
 - Concurrent writes to the same address
 - Result depends on the order of operations
- OS manage
 - Large number of resources
 - Common data structures
- Need to provide good coordination mechanisms
 - Access to the data structures
 - Ensure consistency
- Objective
 - Identify general [synchronization problems](#)
 - Provide [OS support for synchronization](#)

Example: Too Much Milk

Alice

3:00 Look in fridge - no milk

3:05 Leave for shop

3:10 Arrive at shop

3:15 Leave shop

3:20 Back home - put milk in fridge

3:25

Bob

Look in fridge - no milk

Leave for shop

Arrive at shop

Leave shop

Back home - put milk in fridge

Oooops!

Problem: Need to ensure that only one process is doing something at a time (e.g., getting milk)

Example: Money Flies Away ...

BALANCE: 2000 €

Bank thread A

Read $a := \text{BALANCE}$

$a := a + 500$

Write $\text{BALANCE} := a$

Bank thread B

Read $b := \text{BALANCE}$

$b := b - 200$

Write $\text{BALANCE} := b$

Ooops!! BALANCE: 1800 €!!!!

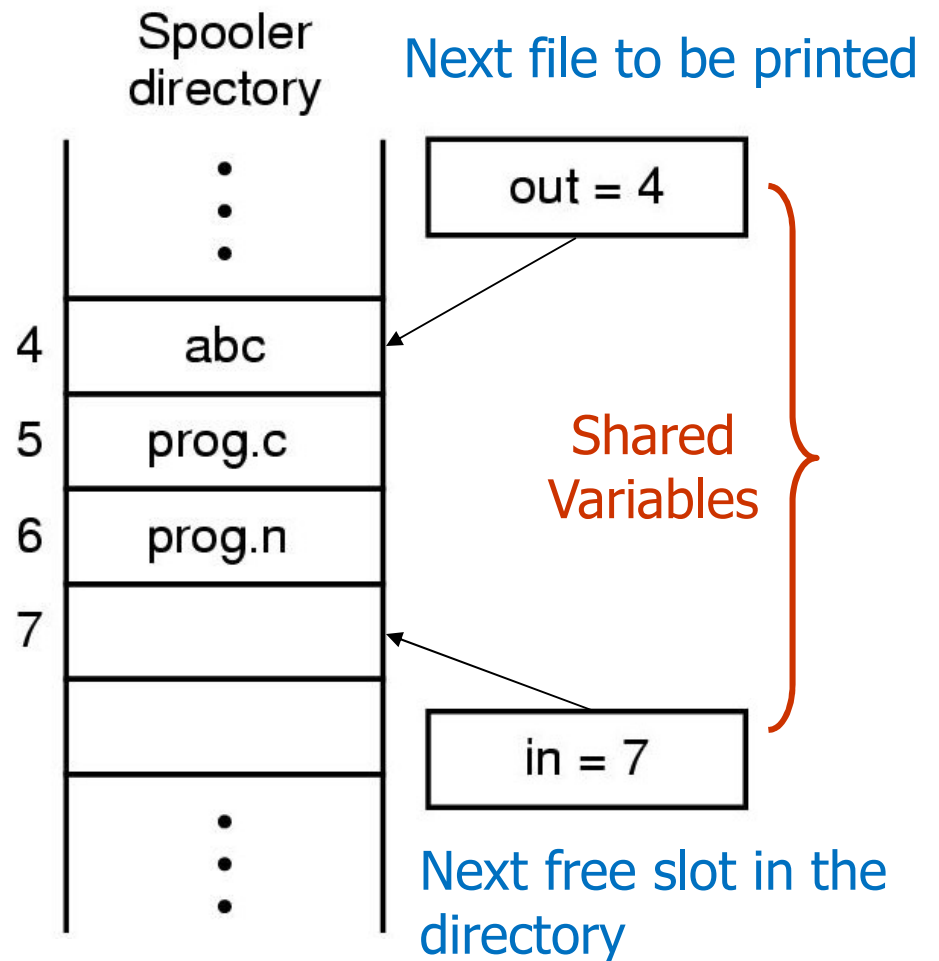
Problem: need to ensure that each process is executing its critical section (e.g updating BALANCE) exclusively (one at a time)

Example: Print Spooler

- If a process wishes to print a file
 - It adds its name in a **Spooler Directory**
- The **printer process**
 - Periodically checks the spooler directory
 - Prints a file
 - Removes its name from the directory

Example: Print Spooler

- If any process wants to print a file it will execute the following code
 1. Read the value of `in` in a local variable `next_free_slot`
 2. Store the name of its file in the `next_free_slot`
 3. Increment `next_free_slot`
 4. Store back in `in`



Example: Print Spooler

Let A and B be processes who want to print their files

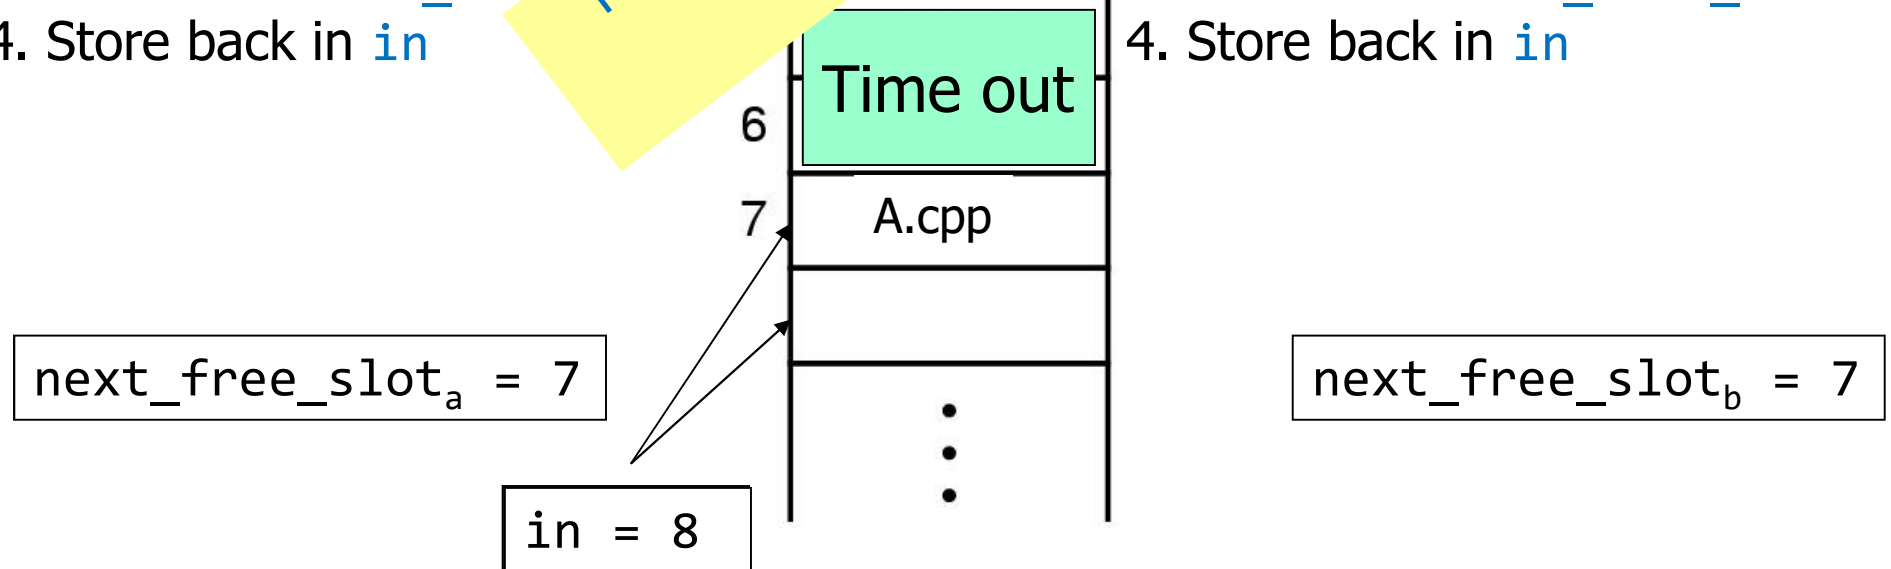
Process A

1. Read the value of `in` in a local variable `next_free_slot`
2. Store the name of its file in the `next_free_slot`
3. Increment `next_free_slot`
4. Store back in `in`

Process B

1. Read the value of `in` in a local variable `next_free_slot`
2. Store the name of its file in the `next_free_slot`
3. Increment `next_free_slot`
4. Store back in `in`

Process B will never receive any output



Example: Producer/Consumer Problem

- An integer count that keeps track of the number of full buffers
- Initially, count is set to 0
 - Count is incremented by the producer after it produces a new buffer
 - Count is decremented by the consumer after it consumes a buffer

Producer

```
while (true) {  
    /* produce an item and put in nextProduced */  
    while (counter == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```


Example: Producer/Consumer Problem

- An integer count that keeps track of the number of full buffers
- Initially, count is set to 0
 - Count is incremented by the producer after it produces a new buffer
 - Count is decremented by the consumer after it consumes a buffer

Consumer:

```
while (true) {  
    while (counter == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in nextConsumed */  
}
```

Producer/Consumer Problem – Race Condition

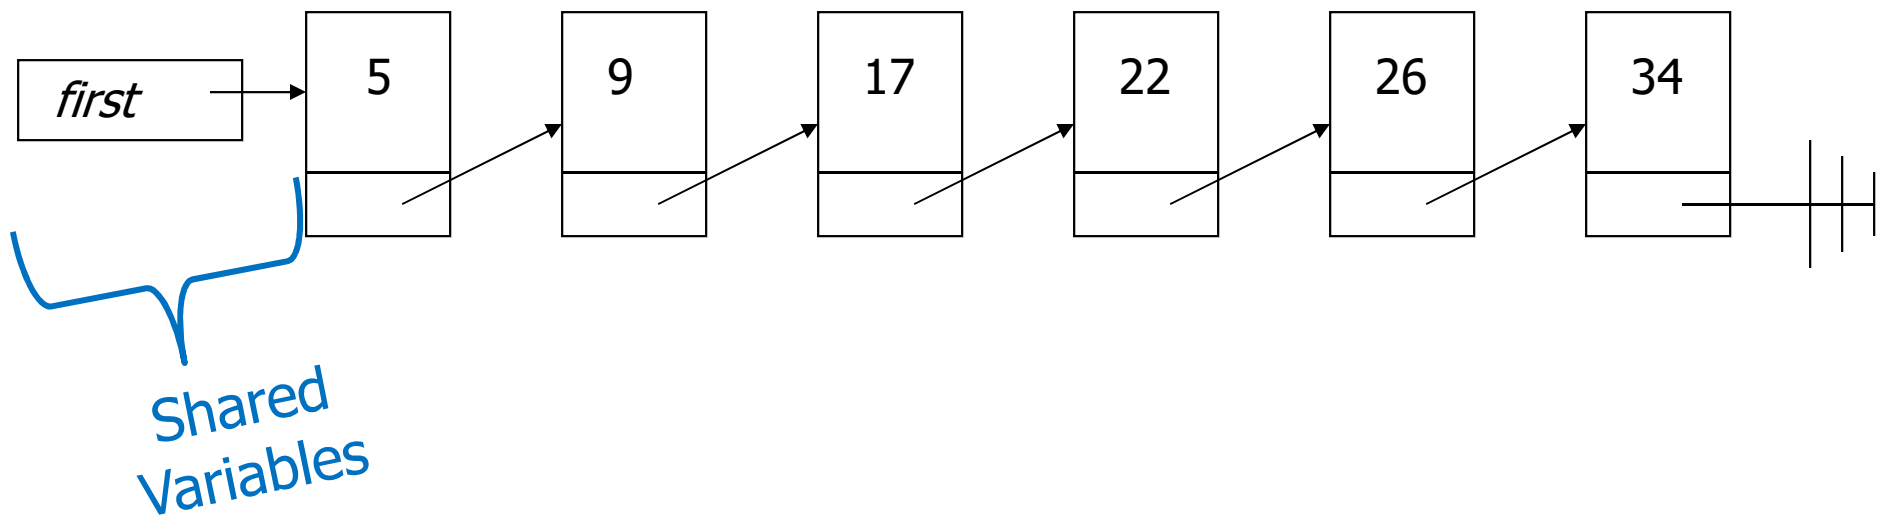
- `counter++` could be implemented as
 - `register1 = counter`
 - `register1 = register1 + 1`
 - `counter = register1`
- `counter--` could be implemented as
 - `register2 = counter`
 - `register2 = register2 - 1`
 - `count = register2`
- Consider this execution interleaving with “count = 5” initially:
 - S0: producer execute `register1 = counter` {`register1 = 5`}
 - S1: producer execute `register1 = register1 + 1` {`register1 = 6`}
 - S2: consumer execute `register2 = counter` {`register2 = 5`}
 - S3: consumer execute `register2 = register2 - 1` {`register2 = 4`}
 - S4: producer execute `counter = register1` {`count = 6`}
 - S5: consumer execute `counter = register2` {`count = 4`}

Race Condition

- Race condition
 - Several processes access and manipulate the same data concurrently
 - Outcome of the execution depends on the particular order in which the access takes place
 - Debugging is not easy → Most test runs will run fine
- At any given time a process is either
 - Doing internal computation → no race conditions
 - Or accessing shared data that can lead to race conditions
- Part of the program where the shared memory is accessed is called **Critical Region**
- Races condition can be avoided
 - If no two processes are in the critical region at the same time

Examples of Race Condition

- Adding node to a shared linked list



Critical-Section (CS) Problem

- n processes all competing to use some shared data / resource
- Each process has a code segment, called **critical section**, in which the shared data is accessed
- **Problem:** ensure that
 - Never two process are allowed to execute in their critical section at the same time
 - Access to the critical section must be an atomic action
- Structure of process P_i :

```
repeat
    entry section
        critical section
    exit section
        remainder section
until false;
```

Critical-Section (CS) – Example

```
void threadRoutine()
```

```
{
```

```
    int y, x, z;
```

```
    y = 3 + 5x;
```

```
    y = Global_Var;
```

```
    y = y + 1;
```

```
    Global_Var = y;
```

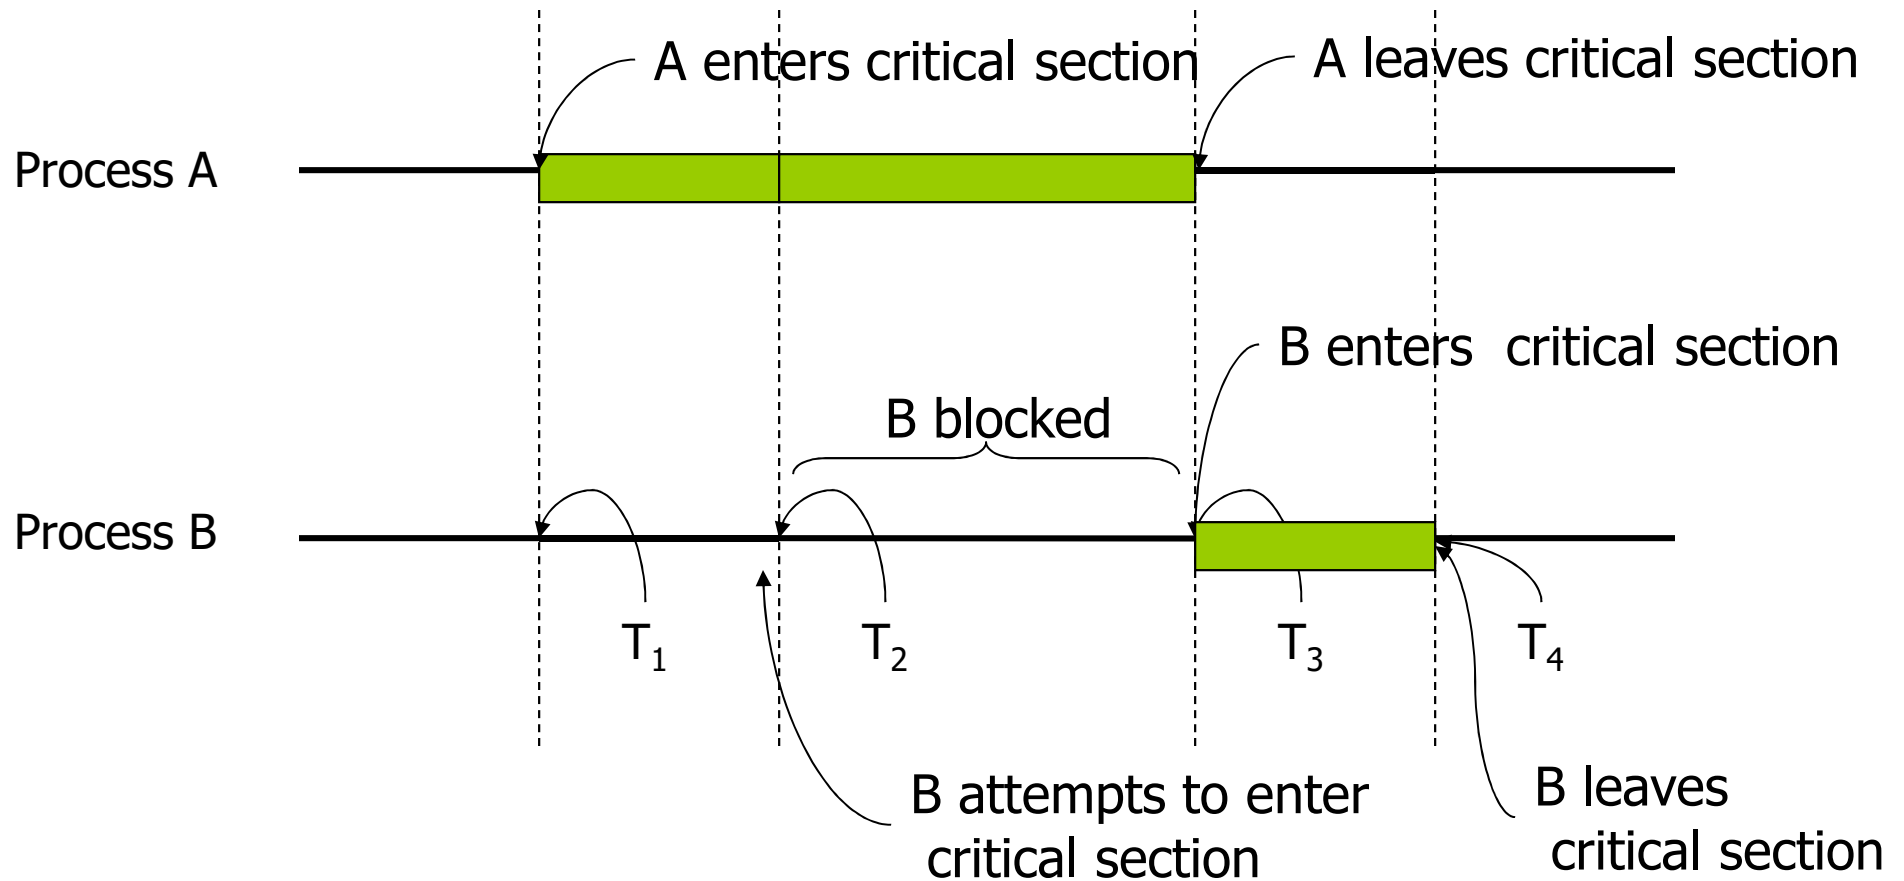


```
    ...
```

```
    ...
```

```
}
```

Critical-Section (CS)



Mutual Exclusion

At any given time, only one process is in the critical section

Requirements: Critical-Section (CS) Problem

Informally, the following are the requirements of CS problem

- No two processes may be simultaneously inside their critical sections
- No assumptions may be made about the speed or the number of processors
- No process running outside its critical section may block other processes
- No process should have to wait forever to enter its critical section

Requirements: Critical-Section Problem

- **Mutual exclusion**
 - Only one process at a time is allowed to execute in its critical section
 - **Progress**
 - Processes cannot be prevented forever from entering CS
-
- Progress: Distinguish between
 - **Fairness (no starvation):** Every process makes progress
 - **Bounded waiting**
 - **Deadlock freedom:** At least one process can make progress
 - Process running outside the critical section should not have influence
 - Assumptions
 - No process stays forever in CS
 - No assumption on the number of processors
 - No assumption on actual speed of processes

Requirements: Critical-Section Problem

1. Mutual exclusion
2. Progress
3. Bounded waiting
 - Before request by a process to enter its critical section is granted
 - A bound must exist on the number of times that other processes are allowed to enter their critical sections

Software Solution – Lock Variables

- Before entering a critical section a process should know if any other is already in the critical section or not
- Consider having a FLAG (also called lock)
 - FLAG = FALSE
 - A process is in the critical section
 - FLAG = TRUE
 - No process is in the critical section

```
// wait while someone else is in the
// critical region
1. while (FLAG == FALSE);
// stop others from entering critical region
2. FLAG = FALSE;
3. critical_section();
// after critical section let others enter
//the critical region
4. FLAG = TRUE;
5. noncritical_section();
```

Lock Variables

Process 1

```
1.while (FLAG == FALSE);  
2.FLAG = FALSE;  
  
3.critical_section();  
4.FLAG = TRUE;  
5.noncritical_section();
```

Process 2

```
1.while (FLAG == FALSE):  
1.while (FLAG == FALSE);  
2.FLAG = FALSE;  
3.critical_section();
```

FLAG = FALSE

Timeout

No two processes may be simultaneously inside
their critical sections

Process 2 's Program counter is at Line 2

Process 1 forgot that it was Process 2 turn

Lock Variables

```
// wait while someone else is in the
// critical region
1. while (FLAG == FALSE);
// stop others from entering critical region
2. FLAG = FALSE;
3. critical_section();
// after critical section let others enter
//the critical region
4. FLAG = TRUE;
5. noncritical_section();
```

- Algorithm does not satisfy critical section requirements
 - Progress is ok, but does NOT satisfy mutual exclusion

Software Solution – Strict Alternation

- We need to remember “Who’s turn it is?”
- If its Process 1’s turn then Process 2 should wait
- If its Process 2’s turn then Process 1 should wait

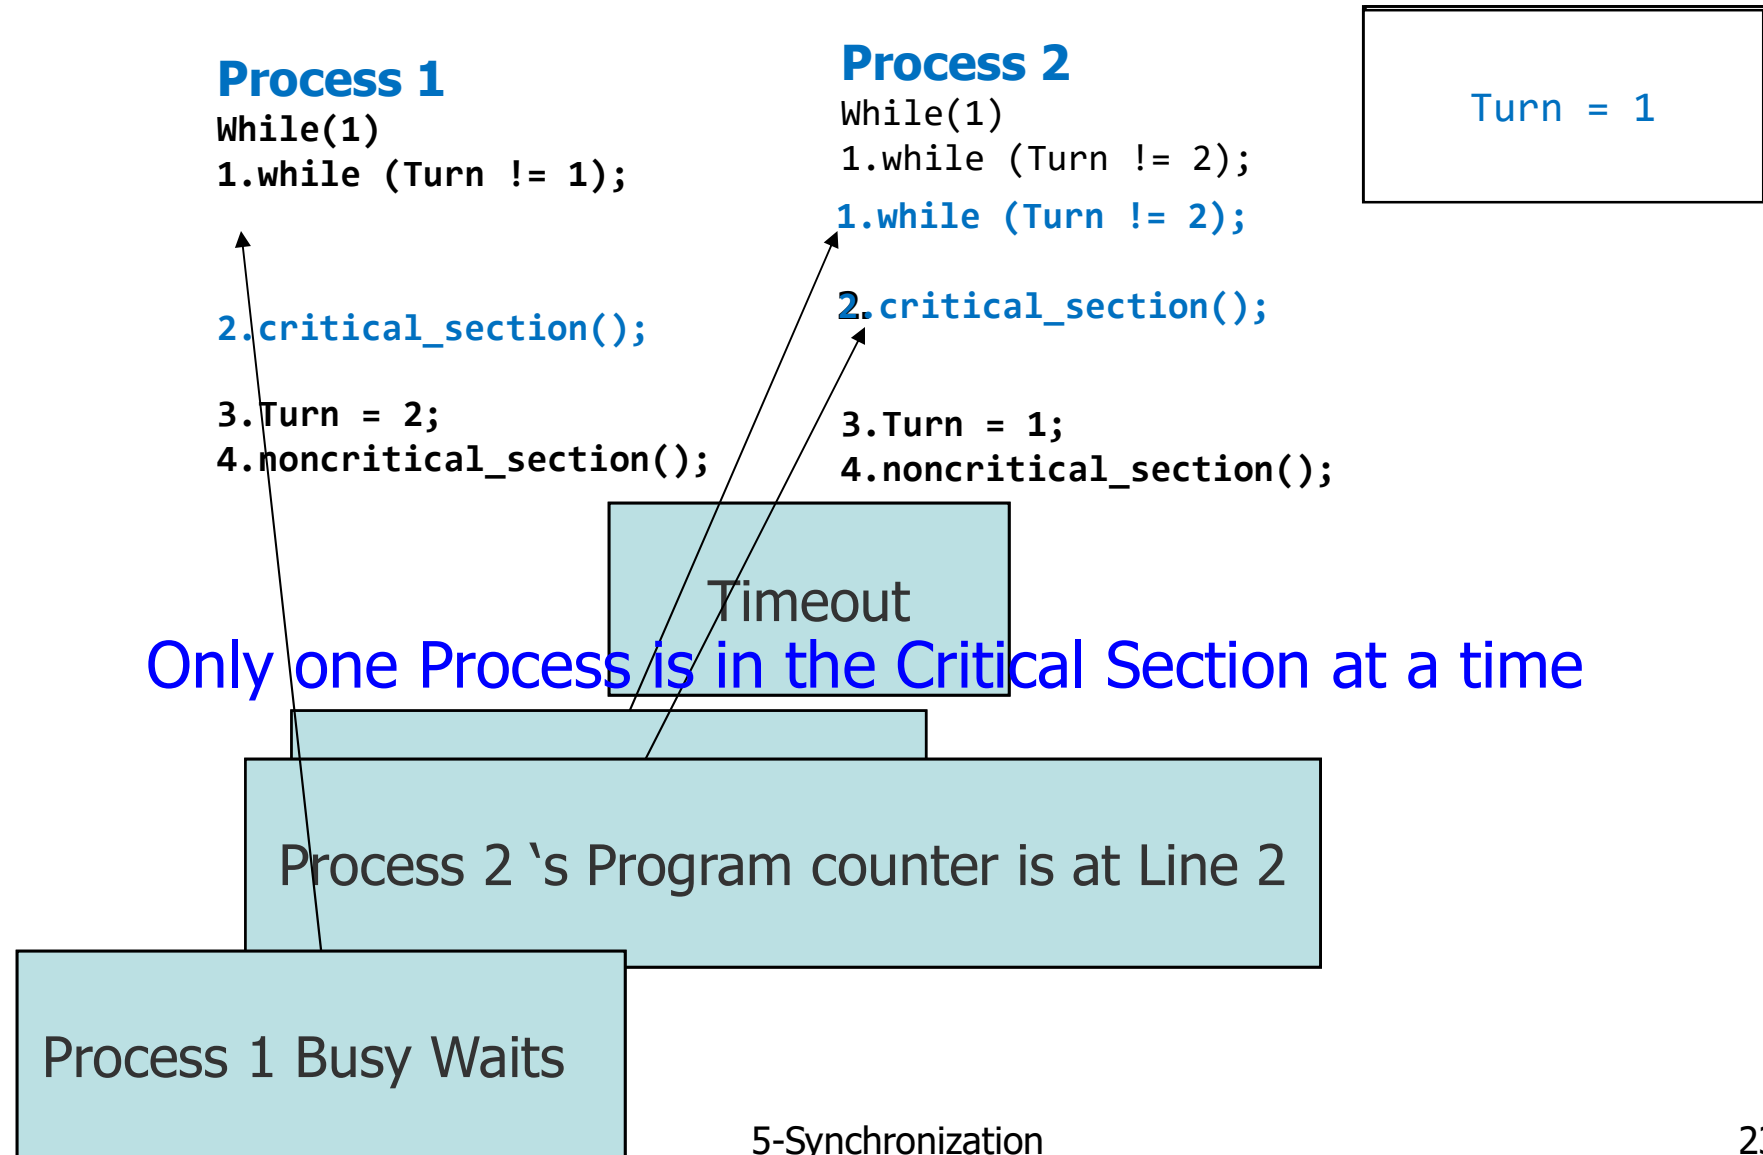
Process 1

```
while(TRUE)
{
    // wait for turn
    while (turn != 1);
    critical_section();
    turn = 2;
    noncritical_section();
}
```

Process 2

```
while(TRUE)
{
    // wait for turn
    while (turn != 2);
    critical_section();
    turn = 1;
    noncritical_section();
}
```

Strict Alternation



Strict Alternation

Process 1

```
while(TRUE)
{
    // wait for turn
    while (turn != 1);
    critical_section();
    turn = 2;
    noncritical_section();
}
```

Process 2

```
while(TRUE)
{
    // wait for turn
    while (turn != 2);
    critical_section();
    turn = 1;
    noncritical_section();
}
```

- What is the problem with strict alteration?
 - Satisfies mutual exclusion, but not progress

Strict Alternation

Turn = 1

Process 1

```
while(TRUE)
{
    // wait for turn
    while (turn != 1);
    critical_section();
    turn = 2;
    noncritical_section();
}
```

Process 2

```
while(TRUE)
{
    // wait for turn
    while (turn != 2);
    critical_section();
    turn = 1;
    noncritical_section();
}
```

- Process 1
 - Runs, Enters its critical section, Exits critical section, Sets turn to 2
- Process 1 is now in its non-critical section
 - Assume this non-critical procedure takes a long time
- Process 2, which is a much faster process, now runs
 - Once it has left its critical section, sets turn to 1
- Process 2 executes its non-critical section very quickly and returns to the top of the procedure

Strict Alternation

Turn = 1

Process 1

```
while(TRUE)
{
    // wait for turn
    while (turn != 1);
    critical_section();
    turn = 2;
    noncritical_section();
}
```

Process 2

```
while(TRUE)
{
    // wait for turn
    while (turn != 2);
    critical_section();
    turn = 1;
    noncritical_section();
}
```

- Process 1 is in its non-critical section
- Process 2 is waiting for turn to be set to 2
 - There is no reason why process 2 cannot enter its critical region as process 1 is not in its critical region

Strict Alternation

Process 1

```
while(TRUE)
{
    // wait for turn
    while (turn != 1);
    critical_section();
    turn = 2;
    noncritical_section();
}
```

Process 2

```
while(TRUE)
{
    // wait for turn
    while (turn != 2);
    critical_section();
    turn = 1;
    noncritical_section();
}
```

Problem with strict alteration

- **Violation:** No process running outside its critical section may block other processes
- This algorithm requires that the processes strictly alternate in entering the critical section
- Taking turns is **not a good** idea **if one of the processes is slower**

Yet Another Solution

- Strick alternation
 - Although it was Process 1's **turn**
 - But Process 1 was **not interested**
- Solution: **Remember whether a process is interested in CS or not**
 - Replace **int turn** with **bool interested[2]**
 - For example **interested[0] = FALSE** → Process 0 is not interested

Process 0

```
while(TRUE)
{
    interested[0] = TRUE;
    // wait for turn
    while(interested[1] != FALSE);
    critical_section();
    interested[0] = FALSE;
    noncritical_section();
}
```

Process 1

```
while(TRUE)
{
    interested[1] = TRUE;
    // wait for turn
    while(interested[0] != FALSE);
    critical_section();
    interested[1] = FALSE;
    noncritical_section();
}
```

Yet Another Solution

Process 0

```
while(TRUE)
{
    interested[0] = TRUE;

    while(interested[1] != FALSE);
```

Process 1

```
while(TRUE)
{
    interested[1] = TRUE;

    while(interested[0] != FALSE);
```

Timeout

DEADLOCK

Peterson's Algorithm

- Combine previous two algorithms
 - `int turn` with `bool interested`

```
Process Pi
while(TRUE)
{
    interested[i] = TRUE;
    turn = j;
    // wait
    while(interested[j]==TRUE && turn == j );
    critical_section();
    interested[i] = FALSE;
    noncritical_section();
}
```

Peterson's Algorithm

Process 0

```
while(TRUE)
{
    interested[0] = TRUE;
    turn = 1;
    // wait
    while(interested[0] == 1 && turn == 1 );
    critical_section();
    interested[0] = FALSE;
    noncritical_section();
}
```

Process 1

```
while(TRUE)
{
    interested[1] = TRUE;
    turn = 0;
    // wait
    while(interested[1] == 1 && turn == 0 );

    critical_section();
    interested[1] = FALSE;
    noncritical_section();
}
```



Timeout

Peterson's Algorithm

Process 0

```
while(TRUE)
{
    interested[0] = TRUE;
    turn = 1;

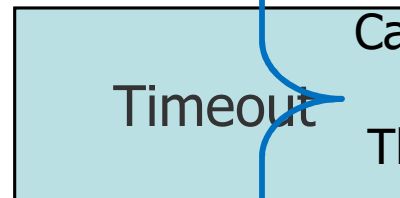
    // wait
    while(interested[0] == TRUE && turn == 1);

    critical_section();
    interested[0] = FALSE;
    noncritical_section();
}
```

Process 1

```
while(TRUE)
{
    interested[1] = TRUE;
    turn = 0;
    // wait
    while(interested[1] == TRUE && turn == 0);

    critical_section();
    interested[1] = FALSE;
    noncritical_section();
}
```



Can not be **TRUE** at the same time.
Thus used to break tie

Lamport's Bakery Algorithm (For n processes)

Idea

- Before entering its critical section, each process receives a number
- Holder of the smallest number enters the critical section
 - Fairness
 - No deadlock
- If processes P_i and P_j receive the same number
 - **if** $i < j$, then P_i is served first
 - **else** P_j is served first
- While processes are trying to enter CS
 - The numbering scheme generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5

Lamport's Bakery Algorithm

- Peterson's algorithm solves the critical-section problem for two processes
 - For multiple processes, bakery algorithm is used

```
Shared var choosing: array [0..n - 1] of boolean (init false);
        number: array [0..n - 1] of integer (init 0),

Process Pi
1: repeat
2:   choosing[i] := true;
3:   number[i] := max(number[0], number[1], ..., number [n - 1]) + 1;
4:   choosing[i] := false;
5:   for j := 0 to n - 1 do begin
6:     while choosing[j] do no-op;
7:     while number[j] ≠ 0 and (number[j], j) < (number[i], i) do
8:       no-op;
9:   end;
10:  critical section
11:  number[i] := 0;
12:  remainder section
13: until false;
```

CS Problem: Software-based Solutions

- Peterson and Bakery algorithms solve critical section problem
- These algorithms are not used in practice
 - Mutual exclusion depends on the order of steps in the software based algorithms
 - However, modern Compilers often reorder instructions to enhance performance

Mutual Exclusion: Hardware Support

- Interrupt Disabling

- A process runs until it invokes an operating-system service or until it is interrupted
- Disabling interrupts guarantees mutual exclusion

```
while(true){  
    /* disable interrupts */  
    critical section  
    /* enable interrupts */  
    remainder section  
}
```

- Disadvantage

- Processor is limited in its ability to interleave programs
- Multiprocessors: Disable interrupts on one processor will not guarantee mutual exclusion
- User programs cannot directly disabled interrupts

Mutual Exclusion: Other Hardware Support

- Special Machine Instructions
 - Performed in a single instruction cycle: Reading and writing in one atomic step
- Not subject to interference from other instructions
 - Uniprocessor system: Executed without interrupt
 - Multiprocessor system: Executed with locked system bus

Test and Set (TSL)

Test and Set Instruction

```
1: boolean testset (int & i) {  
2:     if (i == 0) {  
3:         i = 1;  
4:         return false;  
5:     }  
6:     else return true;  
7: }
```

Test and Set Instruction

```
1: boolean testset (Boolean *target) {  
2:     boolean rv = *target;  
3:     *target = true;  
4:     return rv;  
5: }
```

- Atomic processor instructions
- Idea: Only one process at a time succeeds to set lock=1

Test and Set (TSL)

Test and Set Instruction

```
1: boolean testset(int& i) {  
2:     if (i == 0) {  
3:         i = 1;  
4:         return false;  
5:     }  
6:     else return true;  
7: }
```

```
1: void P(int i){  
2:     while (true) {  
3:         while (testset(bolt));  
4:         critical section  
5:         bolt = 0;  
6:         remainder section  
7:     }  
8: }
```

```
9: void main()  
10: {  
11:     bolt =0;  
13:     parbegin(P(1), ..., P(N));  
14: }
```

- Atomic processor instructions
- Idea: Only one process at a time succeeds to set lock=1

Exchange (Swap)

Exchange Instruction

```
1: void exchange(int& mem1,int& mem2)
2: {
3:     int temp = mem1;
4:     mem1 = mem2;
5:     mem2 = temp;
6: }
```

```
1: void P(int i){
2:     while (true) {
3:         int key = 1;
4:         while(key) {
5:             exchange(bolt,key);
6:         }
7:         critical section
8:         exchange(bolt,key);
9:         remainder section
10:    }
11:}

12: void main(){
13:     bolt =0;
14:     parbegin(P(1), ..., P(N));
15: }
```

- Atomic processor instructions
- Idea: Only one process at a time will be able to set key = 0

Mutual Exclusion Using Machine Instructions

Advantages

- Applicable to any number of processes on single or multiple CPUs sharing main memory
- It is simple and therefore easy to verify

Disadvantages

- **Busy-waiting** consumes processor time
- **Starvation** is possible
 - A process leaves a CS
 - More than one process is waiting
- Deadlock possible **if used in priority-based scheduling systems**: Ex. scenario
 - Low priority process has the critical region
 - Higher priority process needs it
 - **The higher priority process will obtain the CPU to wait for the critical region**

Avoiding Starvation

- Bounded waiting with **Test and Set instruction**

```
do {
    waiting[i] = TRUE;
    key = TRUE;
    while (waiting[i] && key)
        key = TestAndSet(&lock);
    waiting[i] = FALSE;
    // critical section
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = FALSE;
    else
        waiting[j] = FALSE;
    // remainder section
} while (TRUE);
```

Priority Inversion

- Consider three processes
 - Process L with low priority that requires a resource R
 - Process H with high priority that also requires a resource R
 - Process M with medium priority
- If H starts after L has acquired resource R
 - H has to wait to run until L relinquishes resource R
- Now when M starts
 - M is higher priority unblocked process, it will be scheduled before L
 - Since L has been preempted by M, L cannot relinquish R
 - M will run until it is finished
 - Then L will run - at least up to a point where it can relinquish R
 - Then H will run
- **Priority inversion:** Process with the medium priority ran before a process with high priority

Priority Inheritance

- If high priority process has to wait for some resource shared with an executing low priority process
- Low priority process is temporarily assigned the priority of the highest waiting priority process
- Medium priority processes are prevented from pre-empting (originally) low priority process avoiding priority inversion

Semaphores

- Higher-level synchronization construct
 - Designed by Edsger Dijkstra in 1960's



Semaphores



- Synchronization variable accessed by
 - System calls of the operating system
 - Associated with a queue of **blocked processes**
- Accessible via atomic `wait()` and `signal()` operations
 - Originally called `P()` and `V()`

`wait()`: Process blocks until a signal is received

`signal()`: Another process can make progress: Process in the queue, next process calling `wait()`

- Different semaphore semantics possible
 - **Binary:** Only one process at a time makes progress
 - **General (counting):** n processes at a time make progress

Binary Semaphore

- **Wait:** Process blocks if state of semaphore is 0
- **Signal:** Unblocks a process or sets the state of the semaphore to 1

```
wait () / P() {  
    while S <= 0 ; // no-op  
    S--;  
}
```

Atomic

An atomic operation that waits for semaphore to become positive then decrements it by 1

```
signal () / V() {  
    S++;  
}
```

Atomic

An atomic operation that increments semaphore by 1

How to avoid busy waiting?

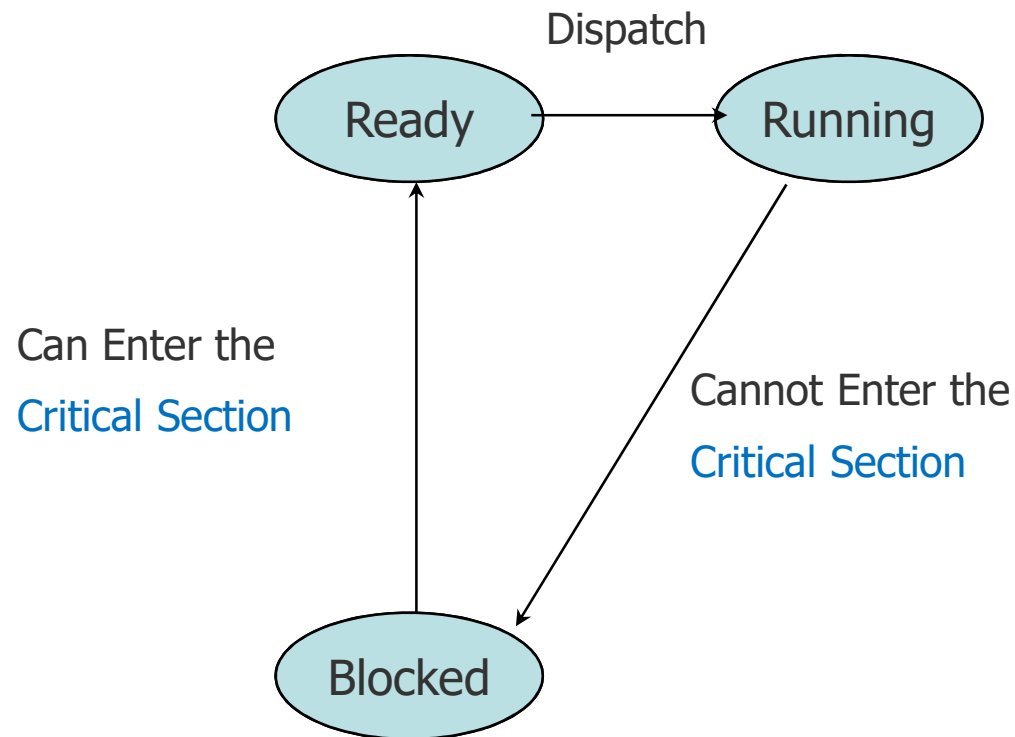
Busy-waiting

- Problem with hardware and software based implementations
 - If one thread already has the lock, and another thread T wants to acquire the lock
 - Acquiring thread T waste CPU's time by executing the Idle loop

To avoid CPU time

- If a thread (process) T cannot enter critical section
 - T's status should be changed from **Running** to **Blocked**
 - T should be removed from the **ready list** and entered in a **block list**
- When T enters critical section
 - T's status should be changed from **Blocked** to **Ready/Running**
 - T should be entered in the **ready list**

Busy-waiting



No CPU time wastage

Sleep/Block and Wakeup Implementation

- `sleep()/block()` and `wakeup()` **can** be two system calls
- When a process P_i determines that it can not enter the CS it calls `sleep()/block()`
 - P_i entered in the block list
- When the other process finishes with the CS, it wakes up P_i , `wakeup(P_i)`
 - P_i entered in the ready list
- `wakeup()` takes a single input parameter, which is the ID of the process to be unblocked

Binary Semaphore

- **Wait:** Process blocks if state of semaphore is 0
- **Signal:** Unblocks a process or sets the state of the semaphore to 1

```
struct bin_sem_t {
    enum {zero, one} value;
    queueType queue;
    /* Can be linked list of PCB */
};

int wait(bin_sem_t s) {
    if (s.value==1)
        s.value=0;
    else{
        /* block and place current
           process in s.queue */
        state(current) = BLOCKED
        append(s.queue, current)
    }
}
```

```
int signal(bin_sem_t s) {
    if (isempty(s.queue))
        s.value=1;
    else{
        /* unblock and remove first
           process from s */
        p=pop(s.queue);
        state(p) = READY/RUNNING
    }
}
```

Counting Semaphore

- **Count:** Determines no. of processes that can still pass semaphore
- **Wait:** Decreases count and blocks if count is < 0
- **Signal:** Unblocks a process or increases count

```
struct sem_t {
    int count;
    queueType queue;
};

int wait(sem_t s) {
    s.count--;
    if (s.count < 0)
        /* block and place current
           process in s.queue */
        state(current) = BLOCKED
        append(s.queue, current)
    }
}
```

```
int signal(sem_t s) {
    s.count++;
    if (s.count ≤ 0){
        /* unblock and remove
           first process from s */
        p=pop(s.queue);
        state(p) = RUNNING/READY
    }
}
```

How to ensure atomicity?

Implementing Semaphores

- **Problem:** Concurrent calls of signal and wait
- **Solution:** Use hardware synchronization!
- **Properties:** Still limited busy waiting

```
int wait(sem_t s)
{
    /* mode switch */
    while (testset(s.flag));
    s.count--;
    if (s.count < 0) {
        /* block and place current
           process in s.queue */
        state(current) = BLOCKED
        append(s.queue, current)
    }
    s.flag=0;
}
```

```
int signal(sem_t s) {
    /* mode switch */
    while (testset(s.flag));
    s.count++;
    if (s.count ≤ 0) {
        /* unblock and remove
           first process from s */
        p=pop(s.queue);
        state(p) = RUNNING
    }
    s.flag=0;
}
```

Critical Section Using Semaphores

Shared var mutex of semaphore (init =1)

Process P_i :

```
1: repeat
2:   wait(mutex)
3:   critical section
4:   signal(mutex);
5:   remainder section
6: until false;
```

Properties:

- Again simple
- No busy waiting
- Fair if queues of semaphores are implemented fair

Bad Use of Semaphores Can Lead to Deadlocks!

- Let S and Q be two semaphores initialized to 1

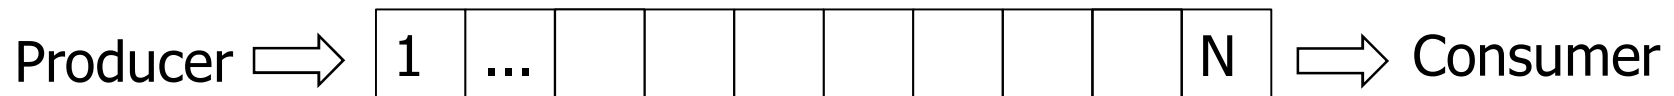
P_0	P_1
1 wait(S);	2 wait(Q);
3 wait(Q);	4 wait(S);
⋮	⋮
signal(S);	signal(Q);
signal(Q);	signal(S);

Synchronization Problems

- Critical section synchronization very common
- **Problems**
 - Synchronization on many variables becomes increasingly complicated
 - Deadlock because of bad programming likely to happen
- **Idea:** Understand useful synchronization patterns that match IPC patterns
 - Consumers/Producer problem
 - High degree of concurrency by decoupling of tasks
 - Dining Philosophers
 - Coordinate access to resources
 - Readers Writers Problem
 - Coordinate access to shared data structures

Producer/Consumer Problem Revisited

- Also known as **bounded buffer problem**



- Overall problem description
 - A buffer can hold N items
 - Producer pushes items into the buffer
 - Consumer pulls items from the buffer
 - Producer needs to wait when buffer is full
 - Consumer needs to wait when the buffer is empty

Producer/Consumer Problem Revisited

```
int BUFFER_SIZE = 100; int count = 0;
```

```
void producer(void) { int item;
```

```
    while(TRUE) {
```

```
        produce_item(&item);
```

```
        if(count == BUFFER_SIZE)
```

```
            sleep ();
```

```
        enter_item(item);
```

```
        count++;
```

```
        if(count == 1)
```

```
            wakeup(consumer);
```

```
    }
```

```
}
```

```
void consumer(void) { int item;
```

```
    while(TRUE) {
```

```
        if(count == 0)
```

```
            sleep ();
```

```
        remove_item(&item);
```

```
        count--;
```

```
        if(count == BUFFER_SIZE - 1)
```

```
            wakeup(producer);
```

```
        consume_item(&item);
```

```
    }
```

```
}
```

5-Synchronization

1 Producer must wait for consumer to empty buffers

2 Consumer must wait for producer to fill buffers

3 A thread can manipulate buffer queue at a time

Producer/Consumer Problem Revisited

- Synchronization between producer and consumer
 - Use of three separate semaphores
- 1. Producer must wait for consumer to empty buffers, if all full
 - Semaphore **empty** initialized to the value N
- 2. Consumer must wait for producer to fill buffers, if none full
 - Semaphore **full** initialized to the value 0
- 3. Producer/consumer can manipulate buffer queue at a time
 - Semaphore **mutex** initialized to the value 1

Producer/Consumer Problem Revisited

- Semaphore **empty** initialized to the value N
- Semaphore **full** initialized to the value 0
- Semaphore **mutex** initialized to the value 1

Producer:

```
void producer(void) {  
    int item;  
    while(TRUE) {  
        produce_item(&item);  
        wait (empty);  
        wait(mutex);  
        enter_item(item);  
        signal(mutex);  
        signal(full);  
    }  
}
```


Consumer:

```
void consumer(void) {  
    int item;  
    while(TRUE) {  
        wait(full);  
        wait(mutex);  
        remove_item(&item);  
        signal(mutex);  
        signal(empty);  
        consume_item(&item);  
    }  
}
```

Changing Order of Semaphores


Producer:

```
void producer(void) {  
    int item;  
    while(TRUE) {  
        produce_item(&item);  
        wait (mutex);  
        wait(empty);  
        enter_item(item);  
        signal(mutex);  
        signal(full);  
    }  
}
```



Consumer:

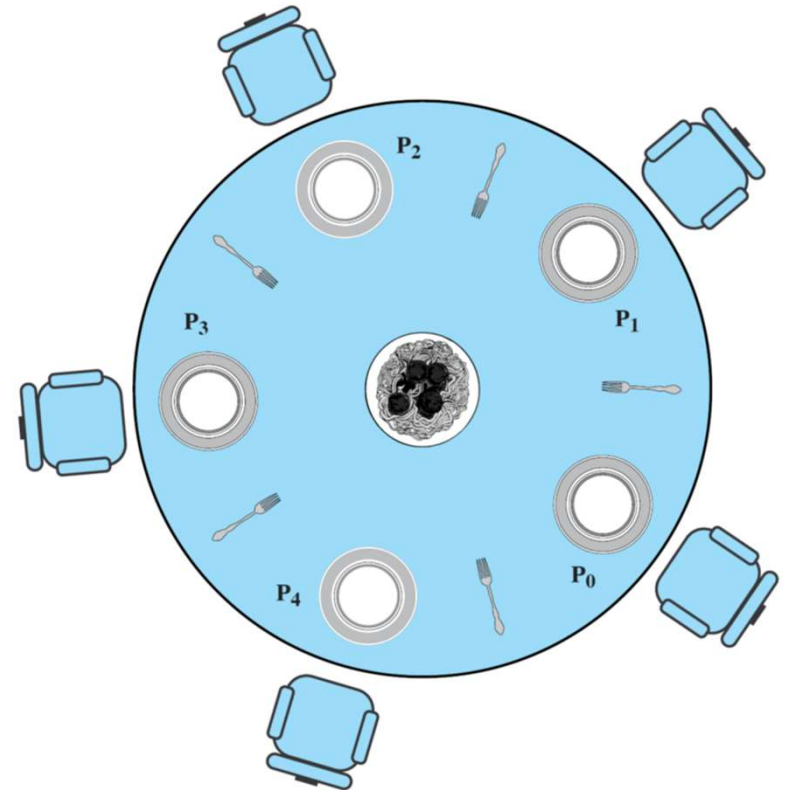
```
void consumer(void) {  
    int item;  
    while(TRUE) {  
        wait(full);  
        wait(mutex);  
        remove_item(&item);  
        signal(mutex);  
        signal(empty);  
        consume_item(&item);  
    }  
}
```



DEADLOCK

Dining Philosophers Problem [E.W. Dijkstra, 1965]

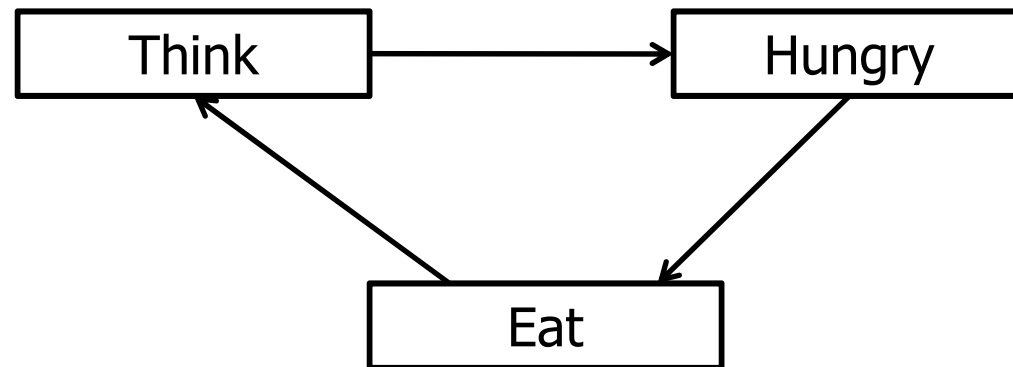
- Simplest form of the resource allocation problem
- **Intuition:**
 - **Philosophers** ~ Processes which aim to execute code
 - **Food** ~ code which requires exclusive access to resources
 - **Forks** ~ conflicts on shared resources
- **Hungry** philosophers
 - Try to **exclusively acquire** forks shared with neighbors
- **Eat** after they have acquired both forks
 - **Release** forks after eating
- **Think** if not interested to eat



Dining Philosophers Problem: States and Properties

Goal

- No deadlocks
- Hungry philosophers will eventually be able to eat



States of a philosopher

Dining Philosophers Problem Using Semaphores

- Represent each fork with a semaphore
 - Semaphore `fork[5]` initialized to 1

Structure of Philosopher *i*

```
do {  
    wait ( fork[i] );  
    wait ( fork[ (i + 1) % 5] );  
    // eat  
    signal ( fork[i] );  
    signal (fork[ (i + 1) % 5] );  
    // think  
} while (TRUE);
```

- Solution guarantees that no two neighbors eat simultaneously
 - Deadlock can be occur

Semaphores Summary

- Semaphores can be used to solve any of traditional synchronization problems
- Semaphore have some **drawbacks**
 - They are essentially shared global variables
 - Can potentially be accessed anywhere in program
 - No connection between the semaphore and the data being controlled by the semaphore
 - Used both for critical sections (mutual exclusion) and coordination (scheduling)
 - No control or guarantee of proper usage
- Sometimes hard to use and prone to bugs
 - Another approach: Use programming language support

Monitors

Goal: Avoid problems of managing multiple semaphores

Monitor is a programming language construct (not OS)

- Controls access to shared data
- Synchronization code added by compiler, enforced at runtime

```
monitor name
{
    // shared variable

    procedure P1 ( ... ) {
        ...
    }
    ...
    procedure PN (...){
        ...
    }
    initialization (...) {
        ...
    }
}
```

Monitors

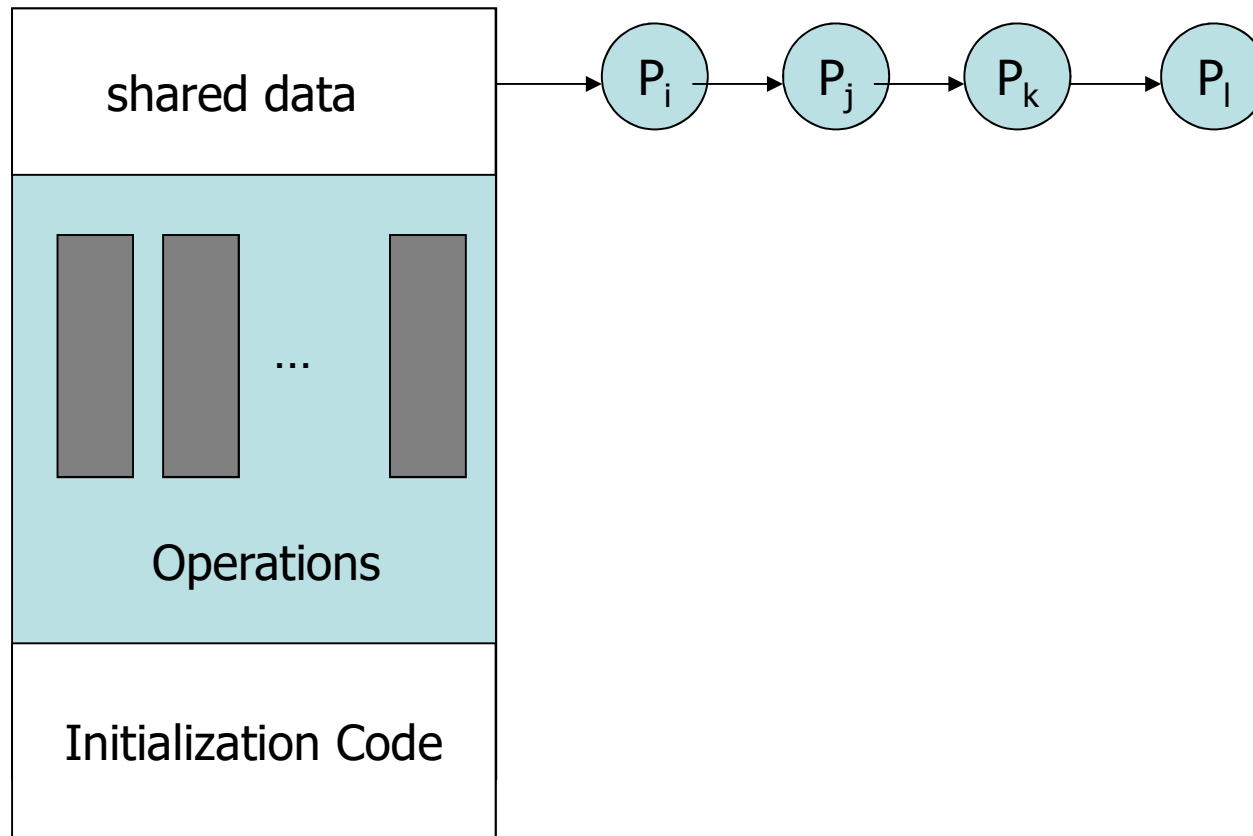
A monitor guarantees **mutual exclusion**

- Only **one thread** can **execute** any monitor **procedure at any time**
 - The thread is “in the monitor”
- If a **second thread invokes** a monitor **procedure** when a **first thread** is already executing one, it **blocks**
 - So the monitor has to have a **wait queue**...
- If a thread within a monitor blocks, another one can enter

```
monitor name
{
    // shared variable

    procedure P1 ( ... ) {
        ...
    }
    ...
    procedure PN (...){
        ...
    }
    initialization (...) {
        ...
    }
}
```

Monitors



Monitors

- A monitor has four components
 - Initialization
 - Share (private) data
 - Monitor procedures
 - Monitor entry queue
- A monitor looks like a class with
 - Constructors, private data and methods
 - Only major difference is that classes do not have entry queues

Mutual Exclusion With a Monitor

```
monitor mutual_exclusion{  
    procedure CS();  
        Critical Section  
    end;  
}
```

```
Process  $P_i$ :  
1: REPEAT  
2:     CS();  
3:     Remainder Section  
4: UNTIL false;
```

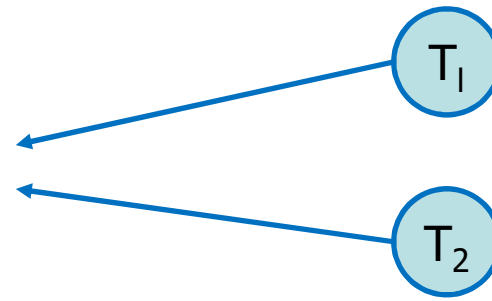
Monitors Example

```
Monitor account {
```

```
    double balance;
```

```
    double withdraw(amount) {  
        balance = balance - amount;  
        return balance;  
    }
```

```
}
```



- Threads are blocked while waiting to get into monitor
- When first thread exists, another can enter

Condition Variables in Monitors

- Similar concept to semaphores
- Condition variable associated with a queue

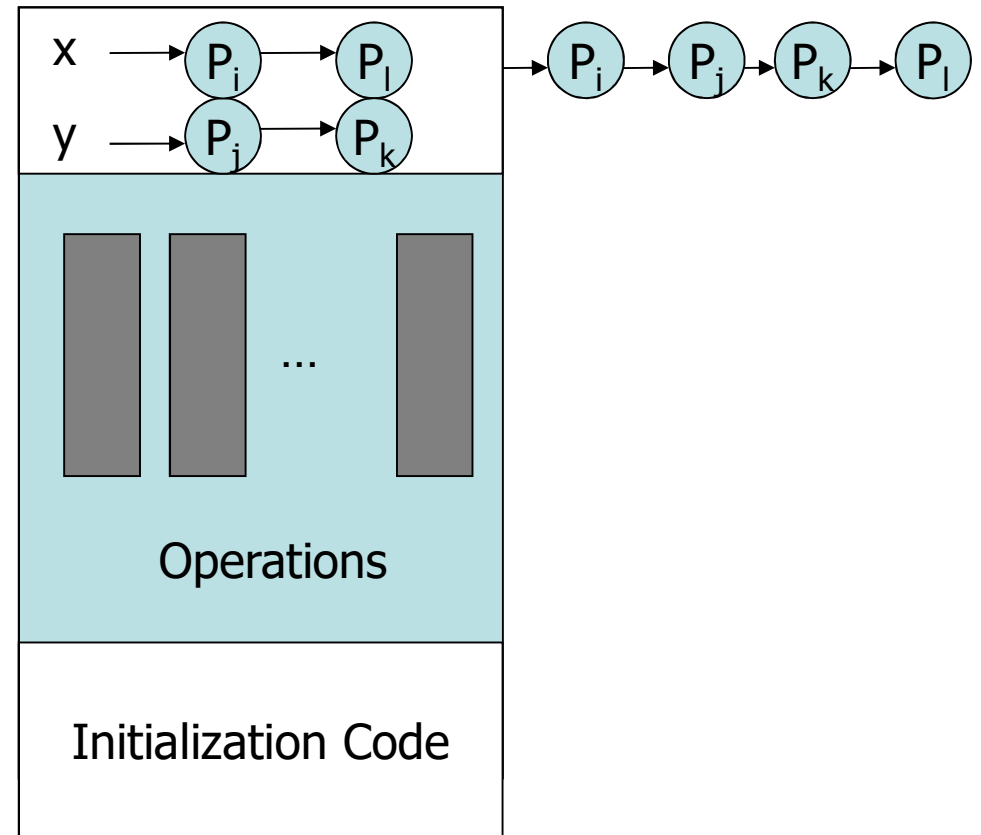
condition x; // declaration

x.wait()

- Process which invokes this operation is blocked

x.signal()

- A blocked process can resume execution
 - If no process issue x.wait () on variable x, then it has no effect on the variable



Producer/Consumer Problem Revisited

```
monitor ProducerConsumer {  
    int itemCount;  
    condition full, empty;  
  
    procedure initialization( ){  
        itemCount = 0;  
    }  
    procedure add(item) {  
        if (itemCount == BUFFER_SIZE)  
            full.wait( );  
  
        putItemIntoBuffer(item);  
        itemCount = itemCount + 1;  
  
        if (itemCount == 1)  
            empty.signal();  
    }  
    procedure remove() {  
        if (itemCount == 0)  
            empty.wait( );
```

```
        item = removeItemFromBuffer();  
        itemCount = itemCount - 1;  
  
        if (itemCount == BUFFER_SIZE - 1)  
            full.signal( );  
        return item;  
    }  
}  
  
procedure producer() {  
    while (true) {  
        item = produceItem();  
        ProducerConsumer.add(item);  
    }  
}  
  
procedure consumer() {  
    while (true) {  
        item = ProducerConsumer.remove();  
        consumeItem(item);  
    }  
}
```

Signal Semantics

What happens if P executes `x.signal()` & has still code to execute?

Two possibilities

- **Hoare monitors** (Concurrent Pascal, original)
 - Signal and wait: P either waits for Q to leave the monitor or waits for some other condition
 - The condition that Q was anticipating is guaranteed to hold
 - `if (empty) wait(condition);`
- **Mesa monitors** (Mesa, Java)
 - Signal and continue: Q either waits until P leaves the monitor or for some other condition
 - Condition is not necessarily true when Q runs again
 - `while (empty) wait(condition);`

Condition Variables vs. Semaphores

Semaphore

- Can be used anywhere in a program, but should not be used in a monitor
- `wait()` does not always block the caller (i.e., when the semaphore counter is greater than zero)
- `signal()` either releases a blocked thread, if there is one, or increases the semaphore counter
- If `signal()` releases a blocked thread, the caller and the released thread both continue

Condition variables

- Can only be used in monitors
- `wait()` always blocks the caller
- `signal()` either releases a blocked thread, if there is one, or the signal is lost as if it never happens
- If `signal()` releases a blocked thread. Only one of the caller or the released thread can continue, but not both

Dining Philosophers Problem Revisited

```
monitor DiningPhilosophers {  
    enum {THINK,HUNGRY,EAT) state[5];  
    condition self [5];  
  
    void pickup (int i) {  
        state[i] = HUNGRY;  
        test(i);  
        if (state[i] != EATING)  
            self [i].wait();  
    }  
    void putdown (int i) {  
        state[i] = THINKING;  
        // test left & right neighbors  
        test((i + 4) % 5);  
        test((i + 1) % 5);  
    }  
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
}
```

```
void test (int i) {  
    if ((state[(i+4)%5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i+1)%5] != EATING)) {  
        state[i] = EATING ;  
        self[i].signal() ;  
    }  
}
```

Function of Philosopher i

```
Void philosopher(int i) {  
    while (true) {  
        DiningPhilosophers.pickup (i);  
        // EAT  
        DiningPhilosophers.putdown (i);  
        // THINK  
    }  
}
```

No deadlock but starvation is possible

Condition Variables and Locks

- Condition variables are also used without monitors in conjunction with blocking locks

```
procedure( ) {  
    Acquire(lock);  
    /* Do something including signal and wait*/  
    Release(lock);  
}
```

- A monitor is “just like” a module whose state includes a condition variable and a lock
 - Difference is syntactic; with monitors, compiler adds the code
- It is “just as if” each procedure in the module calls `acquire()` on entry and `release()` on exit
 - But can be done anywhere in procedure, at finer granularity

Condition Variables and Locks – Example

```
Lock lock;
Condition cond;
int AddToQueue() {
    lock.Acquire();        // Lock before using shared data
    put item on queue;     // Ok to access shared data
    cond.signal();
    lock.Release();        // Unlock after done with shared data
}
int RemoveFromQueue() {
    lock.Acquire();
    while nothing on queue
        cond.wait(&lock); //Atomically release lock and block until signal
                          //Automatically re-acquire lock before it returns
                          // Difference in syntax, e.g., wait(cond, lock)
    remove item from queue;
    lock.Release();
    return item;
}
```

Any Question So Far?

