


Operating Systems

6. Deadlocks

Deadlock

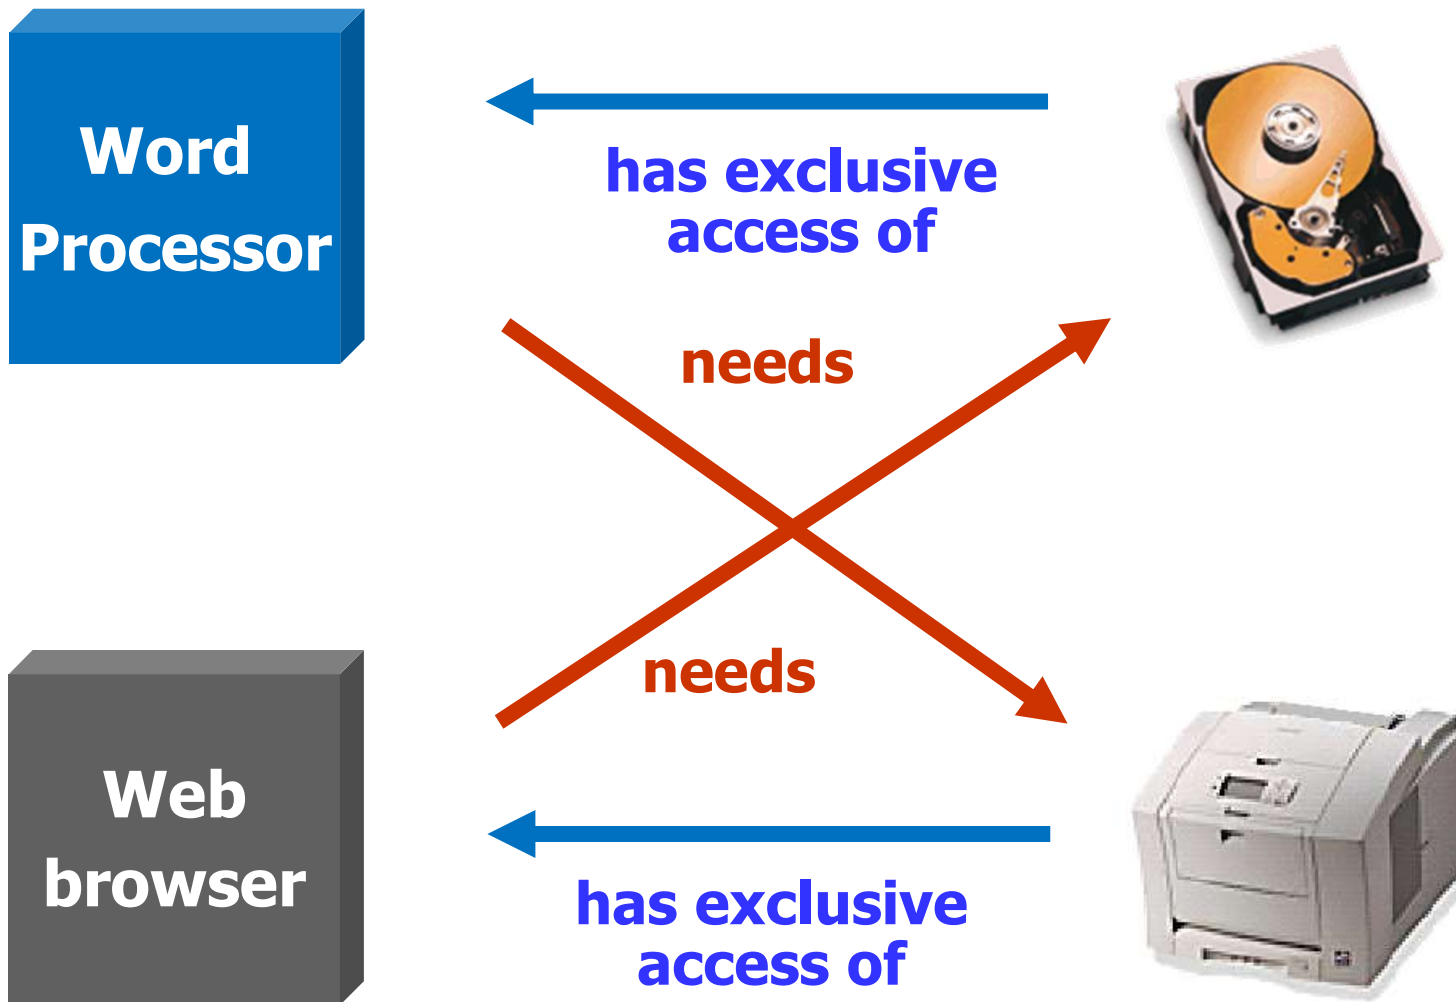
- A set of **blocked processes** each **waiting for an event** that **only another process in the set can cause**
- Example of a possible event
 - Resource to become available

Bridge Crossing Example



- Traffic only in one direction
- Each “half” of the bridge can be viewed as a resource
- If a **deadlock** occurs, it can be **resolved** if cars back up (**preempt resources and rollback**)
 - Several cars may have to be backed up
 - Starvation is possible

Deadlocks in the Computer World



Deadlocks in Resource Allocation – Examples

- **Semaphores** A and B, initialized to 1 (or: system has 2 tape drives; P0 and P1 each hold one tape drive and each needs another one)

P0	P1
wait (A);	wait(B);
wait (B);	wait(A);

- 200Kbytes **memory-space** is available

P0	P1
request (80Kbytes);	request (80Kbytes);
...	...
request (70Kbytes);	request (70Kbytes);

- Deadlock might occur if both processes proceed to the second request

Deadlocks in Resource Allocation – Examples

- Message-passing with blocking receive

P0	P1
receive(P1);	receive(P0)
send(P0, M1);	send(P1 , M0);

Conditions for Deadlock

Four conditions must hold simultaneously for a deadlock to occur
[Coffman-et al 1971]

1. **Mutual exclusion:** Only one process at a time can use a resource
2. **Hold and wait:** A process holding some resource can request additional resources and wait for them if they are held by other processes
3. **No preemption:** A resource can only be released voluntarily by the process holding it, after that process has completed its task
4. **Circular wait:** There exists a circular chain of 2 or more blocked processes, each waiting for a resource held by the next process in the chain

Resource Allocation & Handling of Deadlocks

- **Many possible options dealing with deadlocks in an OS!**
- **Deadlock Prevention:** Structurally restrict the way in which processes request resources (restriction on Conditions 1-3)
- **Deadlock Avoidance:** Processes are required to specify info in advance about their resource usage
 - Info: (max) resources required in running state
 - OS schedules processes in a way that deadlock is avoided
 - **No restrictions** on Conditions 1-3 !
- **Deadlock Detection:** Deadlock state allowed followed by recovery
- **Deadlock Ignorance:** Pretend that deadlocks never occur in the system (can be a “solution” sometimes?!...)

Resource Allocation with Deadlock Prevention

- Restrain the ways requests can be made
 - Attack at least one of the 4 conditions, so that deadlocks are not possible to happen
- 1. **Mutual Exclusion**
 - Cannot do much here ...
- 2. **Hold and Wait**
 - Must guarantee that when a process requests a resource, it does not hold any other resources
 - Require process to request and be allocated **all its resources at once**
 - Or allow process to request resources only when the process has none
 - Low resource utilization; starvation possible

Resource Allocation with Deadlock Prevention

3. No Preemption


- If a process holds some resources, requests another resource that cannot be immediately allocated
- It release all held resources and request them again
 - Risk for starvation!

4. Circular Wait

- Impose total ordering of all resource types
- Require that each process requests resources in an increasing order of enumeration
 - For example, e.g., first the tape, then the disk

Resource Allocation: System Model


- Resource types R_1, R_2, \dots, R_m
 - For example, CPU, memory space, I/O devices, files
 - Each resource type R_i has W_i instances
- Each process utilizes a resource as follows (as discussed for the dining philosophers):
 - Request (hungry)
 - Use (eat)
 - Release



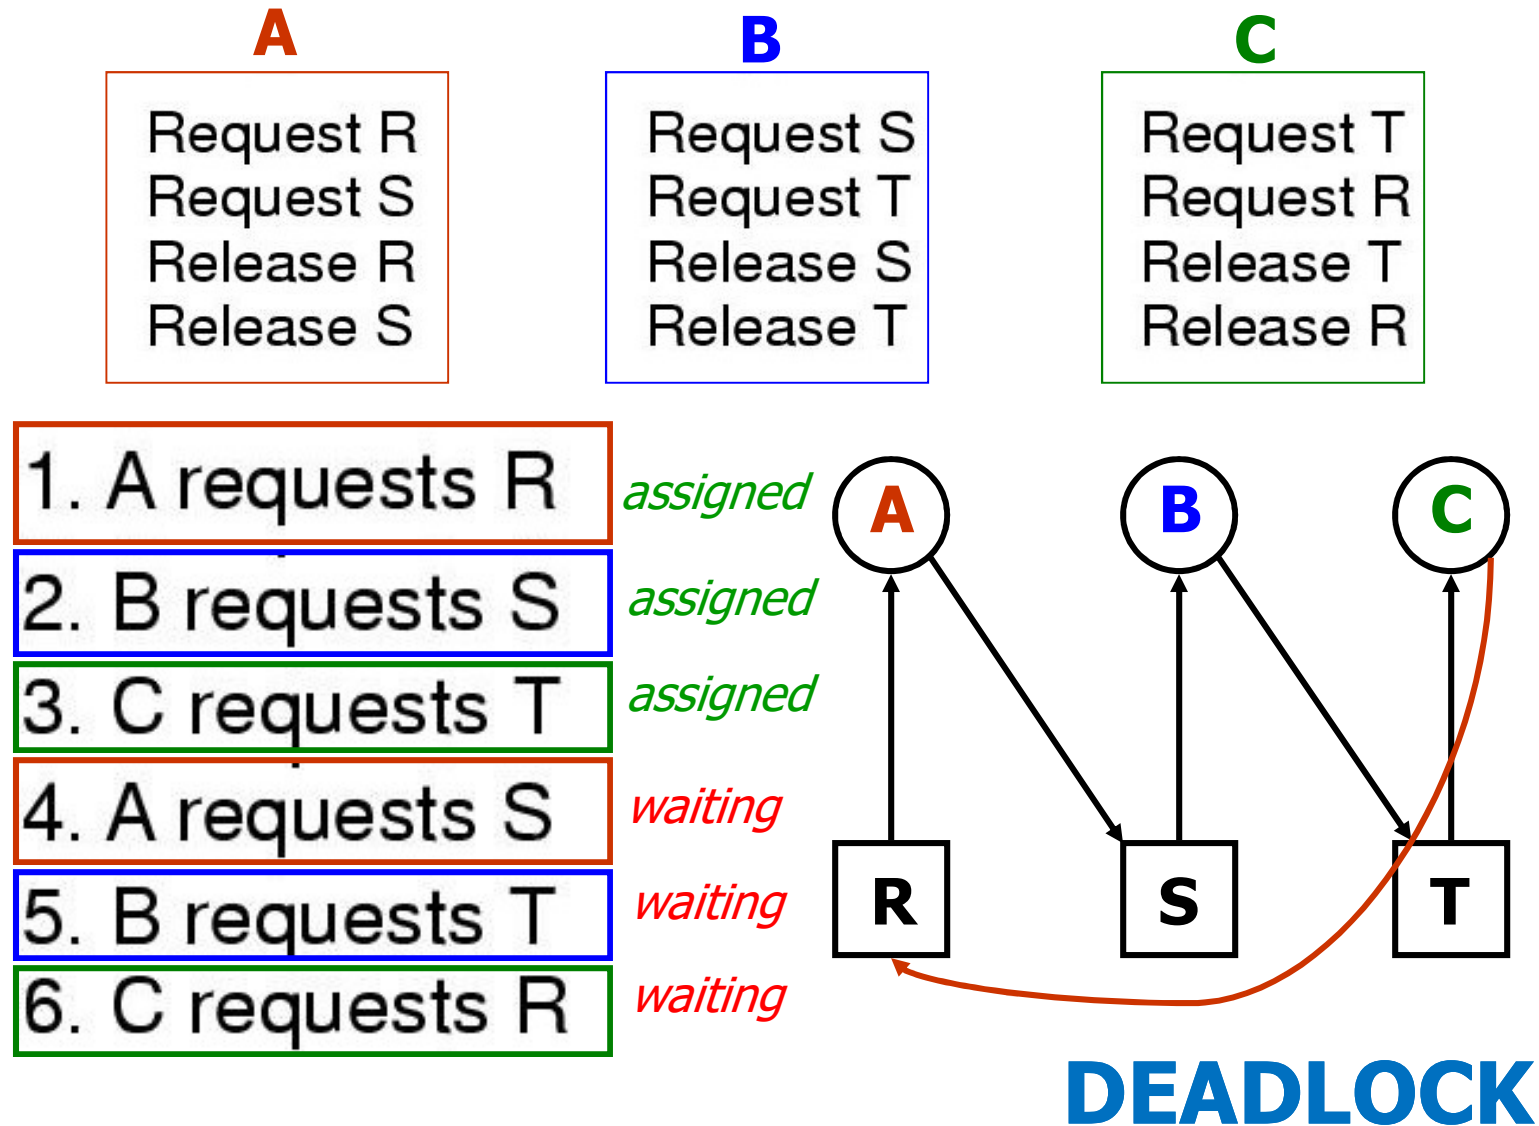
Resource Allocation Graph

A set of vertices V and a set of edges E

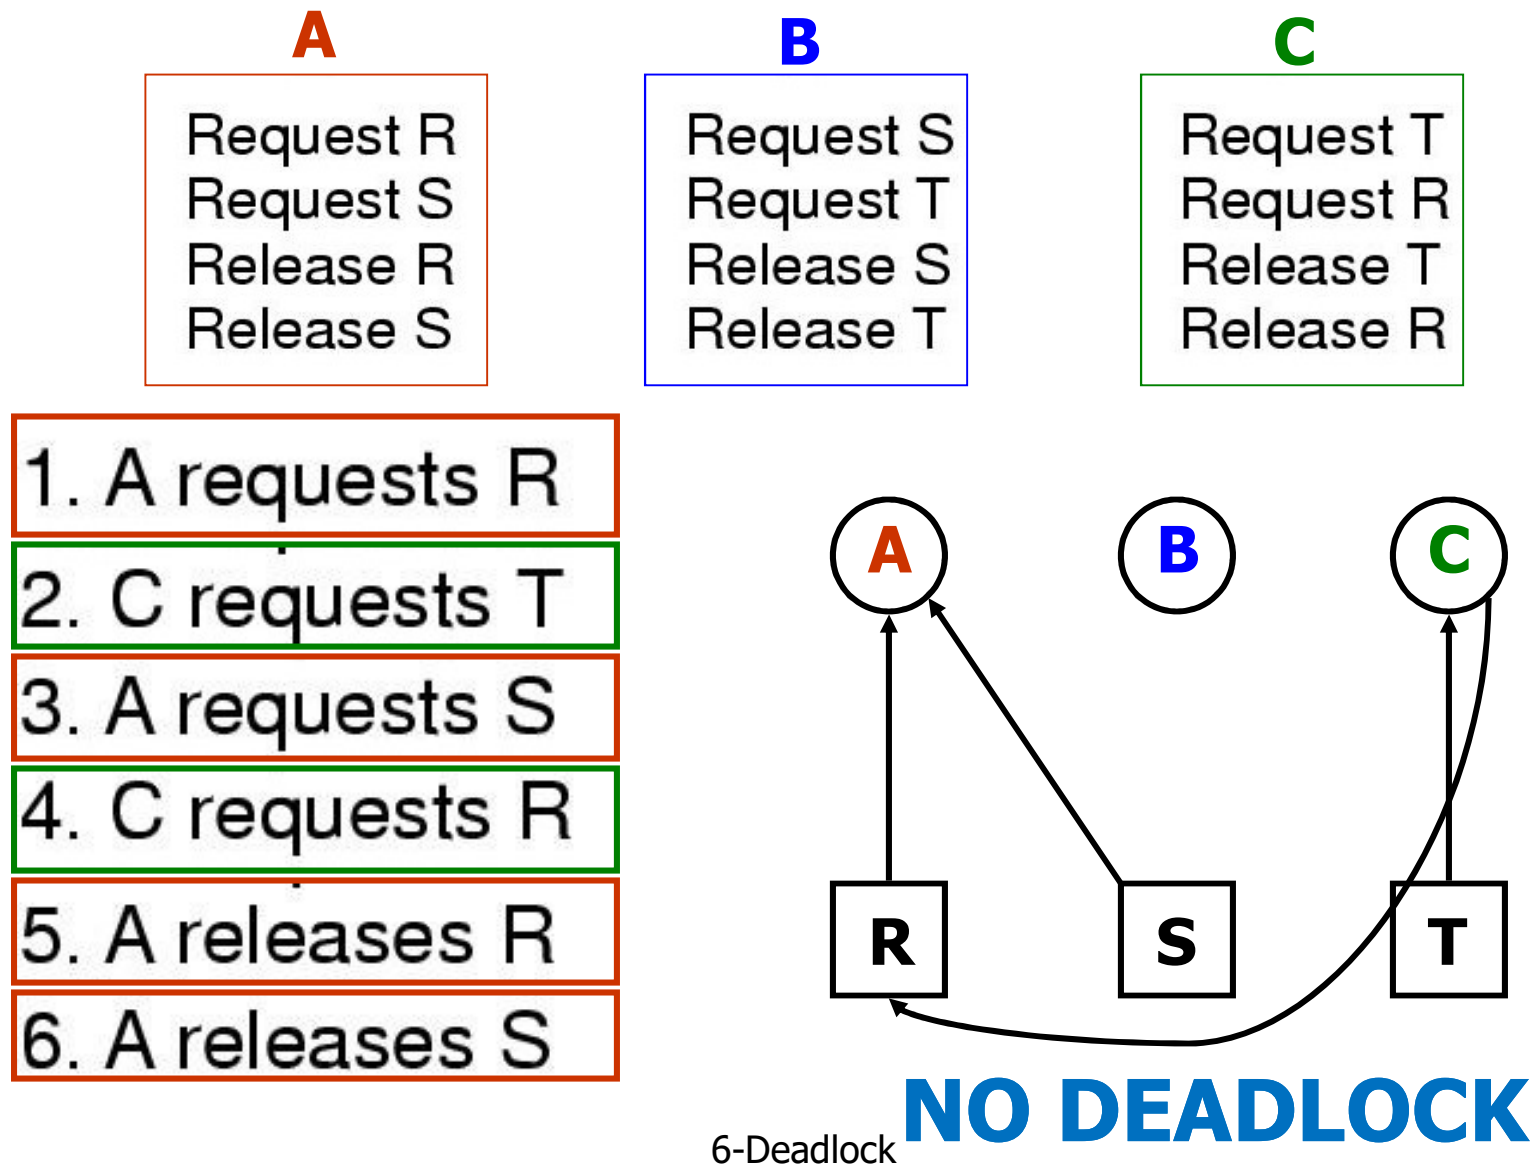
- V is partitioned into two sets
 - $P = \{P_1, P_2, \dots, P_n\}$ the set of **processes**
 - $R = \{R_1, R_2, \dots, R_m\}$ the set of **resource types**
- Request edge: $P_i \rightarrow R_j$
- Assignment edge: $R_j \rightarrow P_i$



Resource Allocation Graph – Example



Resource Allocation Graph – Example



Basic Facts


- Graph contains **no cycles** \Rightarrow **no deadlock**
(i.e., cycle is always a necessary condition for deadlock)
- If graph contains **a cycle** \Rightarrow
 - if \exists **one instance per resource type**, then **deadlock**
 - if \exists **several instances per resource type**, then **possibility of deadlock**

Theorem:

Assume we apply the immediate-allocation-method

If the corresponding graph contains a **knot**, i.e. a strongly connected component and no outgoing edges (sinks), we encounter a deadlock


Resource Allocation Graph: Deadlock



Strongly connected component and no outgoing edge

6-Deadlock


Resource Allocation Graph: Cycle But No Deadlock



Strongly connected component with an outgoing edge

6-Deadlock

Resource Allocation Graph: Cycle But No Deadlock




6-Deadlock

Save, Unsafe, and Deadlock States

Possible global OS states depending on the current resource allocation of processes

- **Safe** state
 - No deadlocks in the horizon
 - OS can find a schedule which avoids a deadlock
- **Unsafe** state
 - Deadlock might occur (later...)
 - For example, future resource requests

Note: Safe states require information on future behavior of processes!



Resource Allocation with Deadlock Avoidance

- **Requires** a priori information available
 - For instance, each **process** declares **maximum number of resources** of each type that it **may need** (e.g., memory/disk pages)

Deadlock-avoidance Algorithm

- Examines the **resource-allocation state**...
 - Available and allocated resources
 - Maximum possible demands of the processes
- **Avoidance** = Ensure that system will not enter an unsafe state

Idea: If **satisfying a request** will result in an **unsafe state**, **requesting process is suspended until enough resources are freed** by processes that will terminate in the meanwhile


Safety Checking: Determining a Safe State

- **Safe state** = \exists a **safe sequence** $\langle P_1, P_2, \dots, P_n \rangle$ of terminating all processes
 - For each P_i , the requests can still be granted by
 - Currently available resources
 - Resources held by P_1, P_2, \dots, P_{i-1}
- The system can **schedule the processes** as follows
 - If P_i 's need for resources cannot be satisfied immediately, then it can
 - Wait until all P_1, P_2, \dots, P_{i-1} have finished
 - Obtain needed resources, execute, release resources, terminate
 - Then the next process can obtain its needed resources, and so on

Simple Example (Snapshot 1)

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B	A B	A B	A B
P1	1 0	1 1	0 1	0 1
P2	0 0	1 1	1 1	

- P2 currently requests resource A (hold by P1)
 - The system is in a safe state since the sequence $\langle P1, P2 \rangle$ satisfies safety criteria




Simple Example (Snapshot 2)

- Consider P_2 also requests B

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B	A B	A B	A B
P1	1 0	1 1	0 1	0 0
P2	0 1	1 1	1 0	

- Allocating B to P_2 leaves the system in an **unsafe state** since there is no sequence that satisfies safety criteria (Available vector is 0 !)



Banker's Algorithm for Deadlock Avoidance

Data Structures

- **Max:** $n \times m$ matrix (n = no. of processes, m = no. of resource types)
 - $\text{Max}[i,j] = k$: P_i may request max k instances of resource type R_j
- **Allocation:** $n \times m$ matrix
 - $\text{Allocation}[i,j] = k$: P_i is currently allocated k instances of R_j
- **Available:** length m vector
 - $\text{Available}[j] = k$: k instances of resource type R_j available
- **Need:** $n \times m$ matrix
 - $\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$: potential max request by P_i for resource type R_j

RECALL: **Avoidance** = Ensure that system will not enter unsafe state

Banker's Algorithm – Example

- Processes P0 through P4
- Resource types A (10 instances), B (5 instances), C (7 instances)
- Snapshot at time T0

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	7 4 3	3 3 2
P1	2 0 0	3 2 2	1 2 2	
P2	3 0 2	9 0 2	6 0 0	
P3	2 1 1	2 2 2	0 1 1	
P4	0 0 2	4 3 3	4 3 1	

- How to check whether the system is in safe state?

Banker's Algorithm: Safety Check

- Work and Finish: Auxiliary vectors of length m and n, respectively

init:

Work := Available

Finish[i] = false for $i = 1, 2, \dots, n$.

while ($\exists i$ s.t. (Finish [i] = false) && ($\text{Need}_i \leq \text{Work}$)) **do**

Work := Work + Allocation_i

Finish[i] := true

if ($\forall i$ Finish[i] = true) **then**

system is in a safe state

else

state is unsafe

Banker's Algorithm – Example

- Processes P0 through P4
- Resource types A (10 instances), B (5 instances), C (7 instances)
- Snapshot at time T0

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	7 4 3	3 3 2
P1	2 0 0	3 2 2	1 2 2	
P2	3 0 2	9 0 2	6 0 0	
P3	2 1 1	2 2 2	0 1 1	
P4	0 0 2	4 3 3	4 3 1	

- The system is in a safe state since the sequence < P1, P3, P4, P2, P0 > satisfies safety criteria

Banker's Algorithm: Resource Allocation

```
For each new Requesti do
  Assert Requesti ≤ Needi
  if Requesti > Available ⇒ Pi must wait
  else
    remember the current resource-allocation state S;
    Available := Available - Requesti;
    Allocationi := Allocationi + Requesti;
    Needi := Needi - Requesti;
    if (safety-check is OK) ⇒
      commit allocation of resources to Pi
    else ( unsafe ) ⇒
      Pi must wait
      restore resource-allocation state S;
```

Banker's Algorithm – Example

- Processes P1 request (1,0,2)

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	7 4 3	3 3 2
P1	2 0 0	3 2 2	1 2 2	
P2	3 0 2	9 0 2	6 0 0	
P3	2 1 1	2 2 2	0 1 1	
P4	0 0 2	4 3 3	4 3 1	

- Check $REQUEST_1 \leq Need_1$ (that is $(1,0,2) \leq (1,2,2)$)
- Check that $Request_1 \leq Available$ (that is, $(1,0,2) \leq (3,3,2)$) \Rightarrow true

Banker's Algorithm – Example

- Processes P1 request (1,0,2)

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	7 4 3	2 3 0
P1	3 0 2	3 2 2	0 2 0	
P2	3 0 2	9 0 2	6 0 0	
P3	2 1 1	2 2 2	0 1 1	
P4	0 0 2	4 3 3	4 3 1	

- Executing safety algorithm shows that sequence <P1, P3, P4, P0, P2> satisfies safety requirement

Banker's Algorithm – Example

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
P0	0 1 0	7 5 3	7 4 3	2 3 0
P1	3 0 2	3 2 2	0 2 0	
P2	3 0 2	9 0 2	6 0 0	
P3	2 1 1	2 2 2	0 1 1	
P4	0 0 2	4 3 3	4 3 1	

- Can request for (3,3,0) by P4 be granted?
- Can request for (0,2,0) by P0 be granted?

Deadlock Detection & Recovery

- Allow system to enter a deadlock state
- Detection algorithm
 - Using Banker's algorithm idea
 - Resources can have multiple instances
 - Using resource allocation graph
 - Resources have single instance
- Recovery scheme

Deadlock Detection Using Banker's Algorithm

- Similar as detecting unsafe states using Banker's algorithm
 - How is similarity explained?
 - If they cost the same why not use avoidance instead of detection & recovery?

Data structures

- **Available:** Vector of length m : number of available resources of each type
- **Allocation:** $n \times m$ matrix: number of resources of each type currently allocated to each process
- **Request:** $n \times m$ matrix: current request of each process
 - Request $[i \ j] = k$: P_i is requesting k more instances of resource type R_j

Deadlock Detection Using Banker's Algorithm

- Work and Finish: Auxiliary vectors of length m and n, respectively

init:

Work := Available

Finish[i] = false for $i = 1, 2, \dots, n$.

while ($\exists i$ s.t. (Finish [i] = false) && (Request_i ≤ Work)) **do**

Work := Work + Allocation_i

Finish[i] := true

if ($\exists i$ Finish[i] = false) **then**

system is in a deadlock state

Example of Detection Algorithm

- Five processes P_0 through P_4
- Three resource types
 - A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $\text{Finish}[i] = \text{true}$ for all i

Example of Detection Algorithm


- P_2 requests an additional instance of type C
- Snapshot at time T_0

	Allocation	Request	Available
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 1	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	


- State of system?
 - Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes' requests
 - Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4

Deadlock Detection Using Graphs


- Deadlock detection using **wait-for graph**
 - Variant of resource allocation graph
- **Wait-for graph** only contains **edges between processes**
 - Edge $P_i \rightarrow P_j$ in wait-for graph implies
 - Process P_i is waiting for P_j to release a resource that P_i needs
- Wait-for graph is **obtained from resource allocation graph**
 - Removing the resource nodes
 - Collapsing the appropriate edges
- Edge $P_i \rightarrow P_j$ exists in wait-for graph if and only if
 - Resource allocation graph contains two edges
 - $P_i \rightarrow R_q$ and $R_q \rightarrow P_j$ for some resource R_q



Deadlock Detection Using Graphs



Resource-Allocation Graph



Wait-for graph

Deadlock Detection Using Graphs

- **Deadlock** exists if and only if **wait-for graph contains cycle**
- Periodically invoke an algorithm that searches for a cycle in graph
 - Which algorithm can be used for deadlock detection?
- An algorithm to detect a cycle in a graph requires an order of n^2 operations
 - Where n is the number of vertices in the graph

Detection Algorithm Usage

- Note: cost for detection algorithm is the same as that for the Bankers algorithm!
- When, and how often, to invoke depends on
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
- If algorithm is invoked arbitrarily, there may be many cycles in the resource graph
 - We would not be able to tell which of the many deadlocked processes “caused” the deadlock

How to Recover From Deadlocks

- Once there is a deadlock situation
 - No way that all processes can proceed with the current allocation state
- Two main techniques to recover
 - Process Termination
 - Resource Preemption

Recovery From Deadlock: Process Termination

- **Abort all** deadlocked processes
 - Most common
 - Long computations may be lost
- **Abort one process at a time** until deadlock is eliminated
 - Need to run after every abort another instance of the deadlock detection algorithm
- In which order should we choose to abort? Criteria?
 - Least amount of processor time consumed so far
 - Least amount of output produced so far
 - Most estimated time remaining
 - Least total resources allocated so far
 - Lowest priority

Recovery From Deadlock: Resource Preemption

- Select a **victim**
 - Which resources to preempt from which processes
 - As in process termination, determine order of preemption to minimize cost
- **Rollback**
 - Return to some safe state, restart process from that state
 - Checkpointing necessary
- **Watch for starvation**
 - Resources may always be preempted from same process
 - A process must only be picked as victim only a (small) finite no. of times
 - Number of rollbacks are a cost factor

Combined Approach to Deadlock Handling

Combine the three basic approaches

- Prevention, avoidance, detection
- Can use the optimal approach for each type of resources in the system

Approach

- **Group resources** into number of different resource classes
- Use **linear ordering to prevent deadlocks between resource classes** (i.e., prevent circular wait)
- Use most **appropriate technique for handling deadlocks within each class**, e.g.,
 - Swap space (e.g., blocks in disk): **Prevention by preallocation**
 - Process resources (e.g., files): **Avoidance by knowing max needs**
 - Main memory: **Prevention by preemption**
 - Internal (e.g., I/O channels): **Prevention by ordering**

Summary

- Studied conditions when deadlocks occur
- Discussed options to deal with deadlocks
- Most operating systems completely ignore the problem
 - Unix, Windows, ...
- None of the approaches suitable for all resources
 - Possibly need to select optimal approach for each class of resource

Any Question So Far?

