


Operating Systems

7. Memory Management

Memory Management

- Ideally programmers want memory that is
 - Large
 - Fast
 - Non volatile
- Memory hierarchy
 - Small amount of fast, expensive memory – cache
 - Some medium-speed, medium price main memory
 - Gigabytes of slow, cheap disk storage
- Memory manager handles the memory hierarchy



Requirements for Memory Management Unit (MMU)

- Relocation
 - Change the physical placement of a process
- Protection
 - Restrict processes access to physical memory
- Sharing
 - Allow processes to share memory location
- Logical Organization
 - Support organization of computer programs etc.
- Physical Organization
 - Efficient utilization of hardware
 - Memory hierarchy

Background

Program

- Typically resides on disk
- Must be brought into memory to be executed

Address binding of a program


- **Compile time**
 - Program must always be loaded to the same memory location
- **Load time**
 - Memory location determined when program is loaded
- **Run time/ Execution time**
 - No absolute memory locations
 - Memory location determined by memory management unit (MMU)
 - Swapping of processes and relocation

Background

Memory

- Can be subdivided to accommodate multiple processes
- Needs to be allocated efficiently to pack as many processes into memory as possible

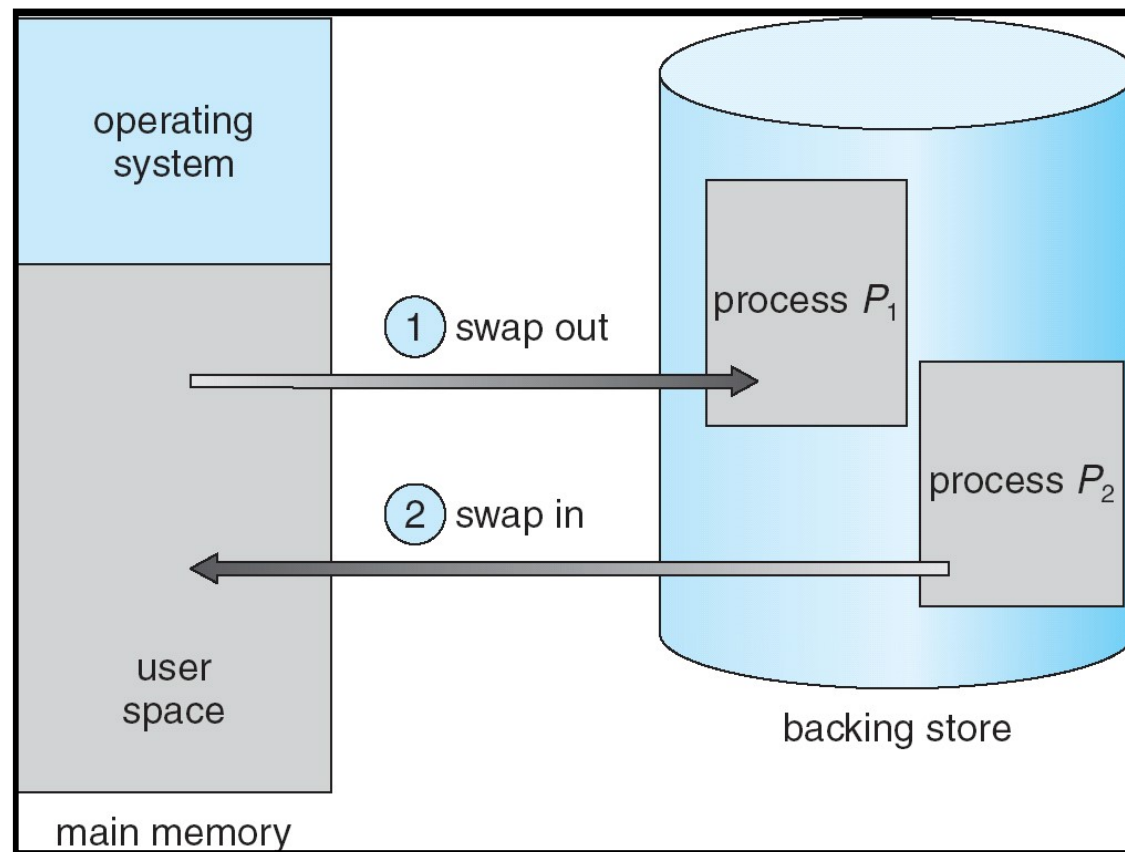
Position and Function of the MMU



- **Logical address:** Generated by the CPU
 - Also referred to as virtual address or relative address
- **Physical address:** Address seen by the memory unit
 - Also referred to as absolute address

Sharing Memory Through Swapping

- A process can be **swapped out** of memory to a backing store
 - Swap device
- Later brought back (**swap-in**) into memory for continued execution



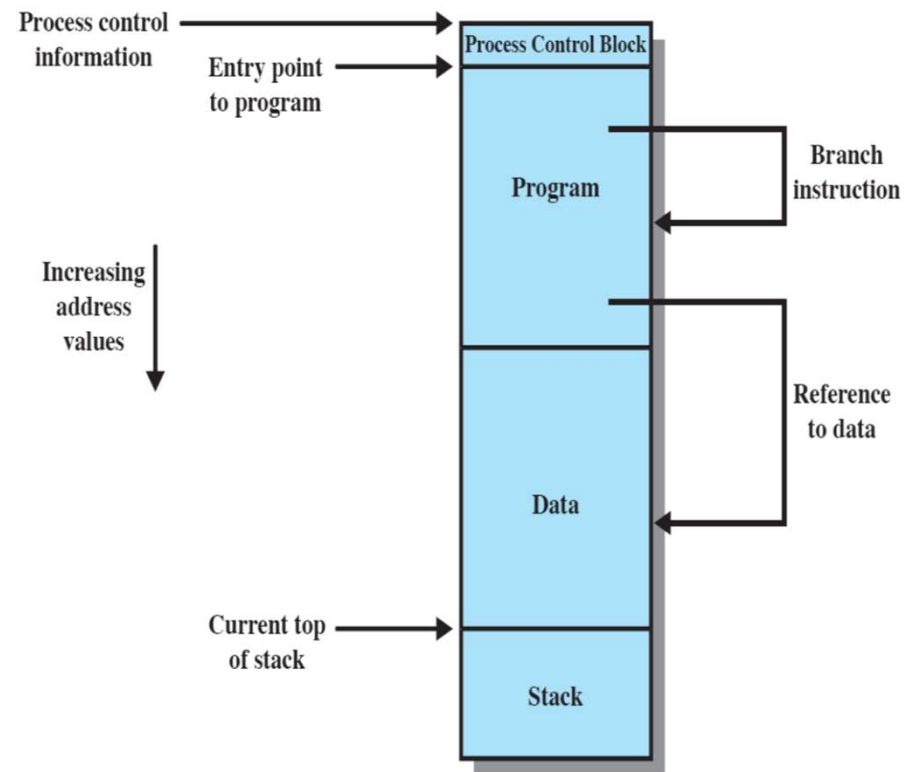
Relocation and Protection

Relocation

- Swapping
 - Program must be relocated
- Cannot be sure about the location in memory
- Addressing
 - No absolute addressing feasible
 - For example, locations of variables, code routines

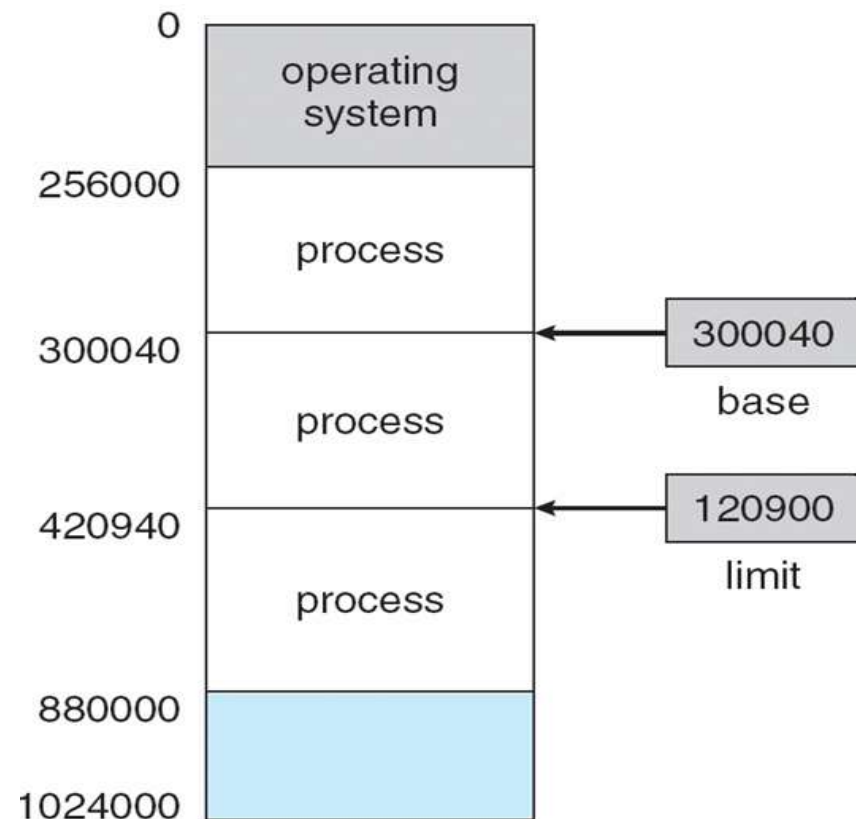
Protection

- Must keep a program out of other processes' partitions



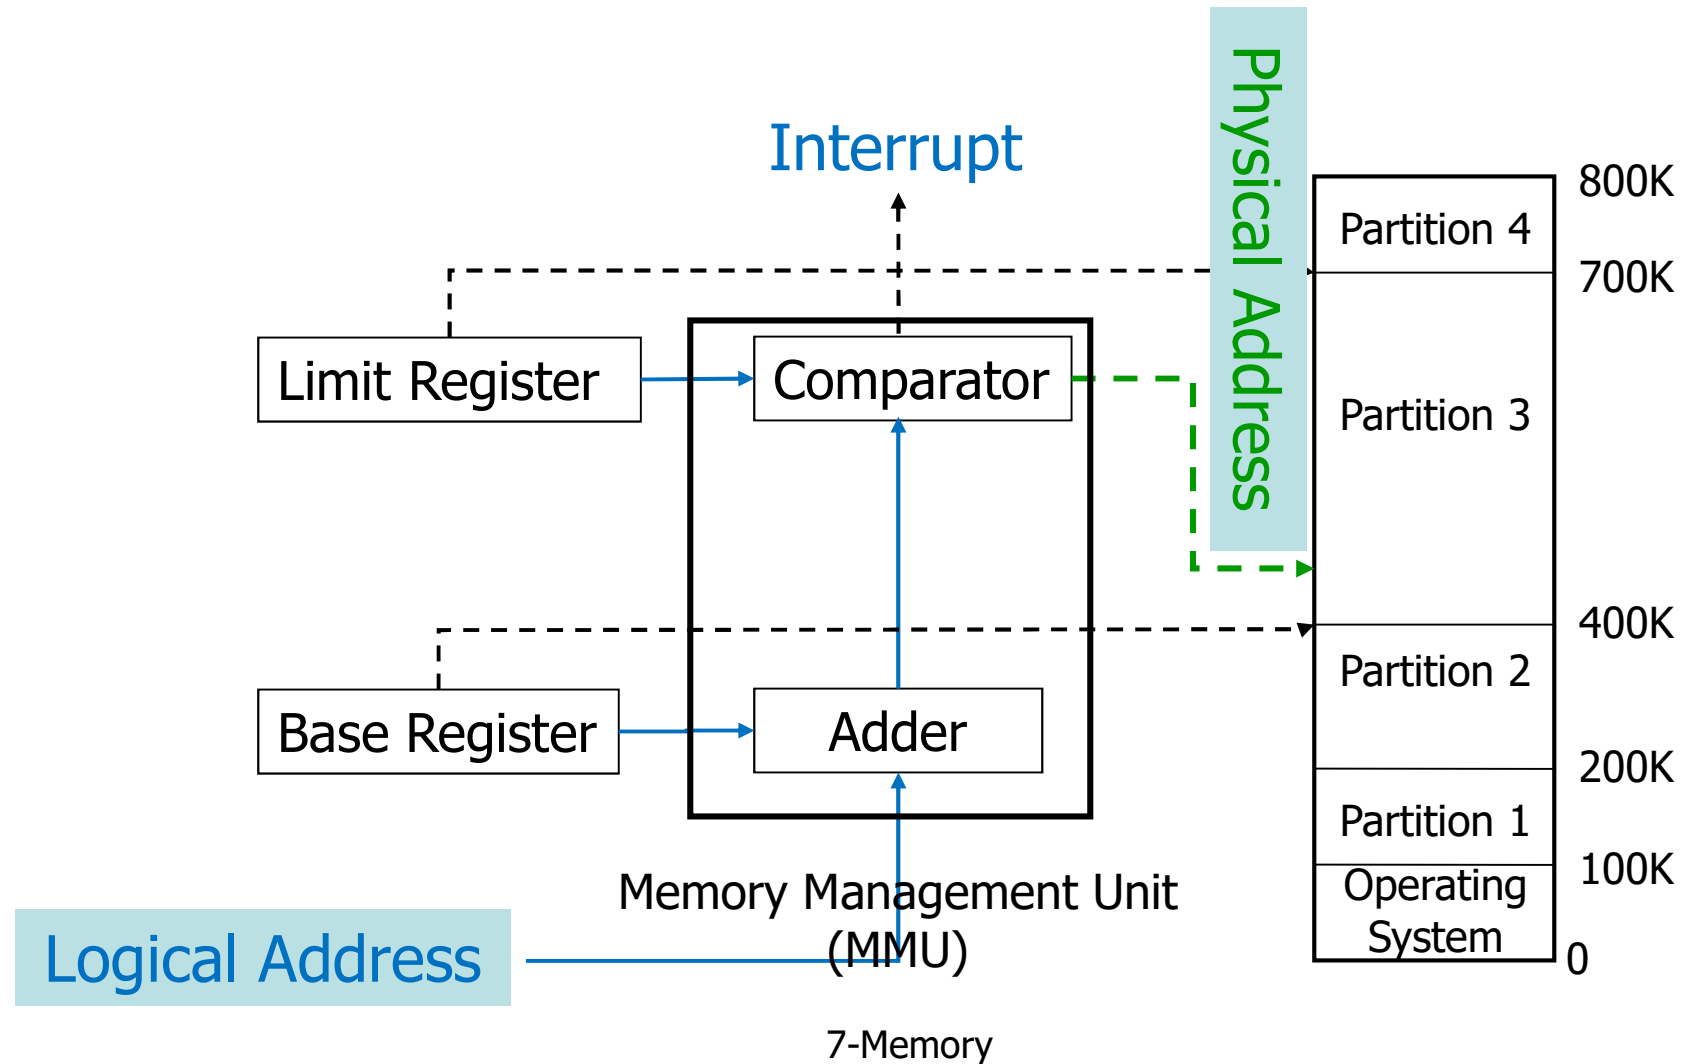
Hardware Support For Relocation and Protection

- Base, bounds (limit) registers
 - Define the range which a process can access
 - Set when the process is executing



Hardware Support For Relocation and Protection

Addition and comparison has to be performed on every address




Contiguous Allocation of Memory: Fixed Partitioning

- Any program, no matter how small, occupies an entire partition
- Causes **internal fragmentation**
 - Allocated memory can be larger than requested memory
 - Difference internal to a partition, but not being used by the process



(a) Equal-size partitions

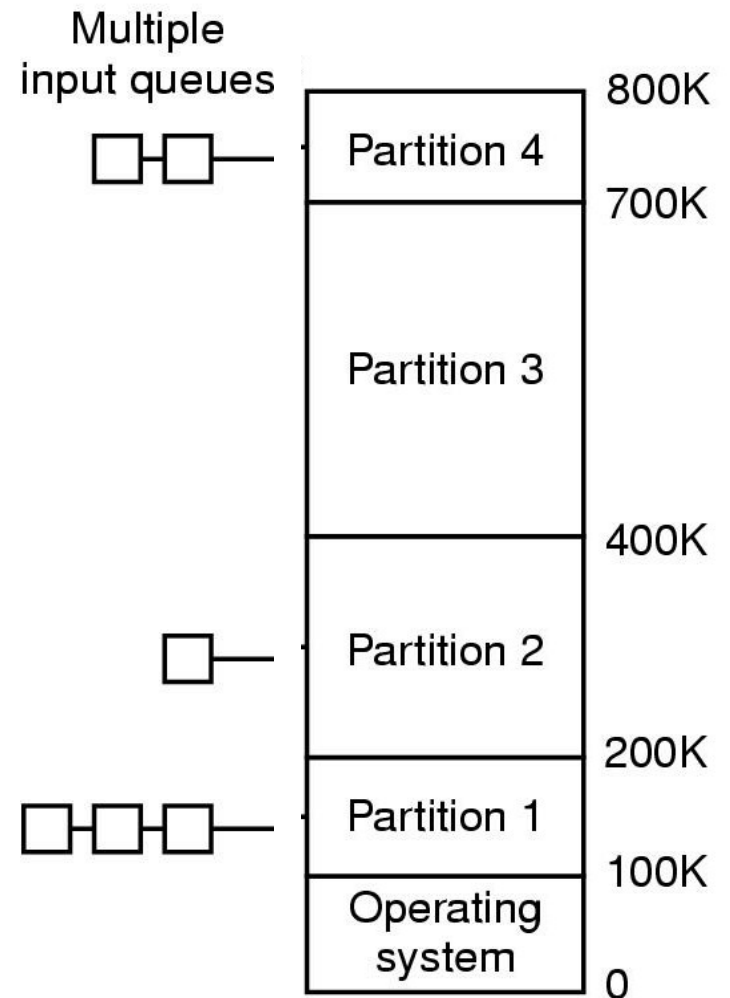


(b) Unequal-size partitions

Fixed Partitioning

Multiple Queues


- Small processes have to wait, even though plenty of memory is free
 - When a large partition is empty
 - Queues for small partition is full
- Alternative: Maintain a **single Queue**



Fixed Partitioning

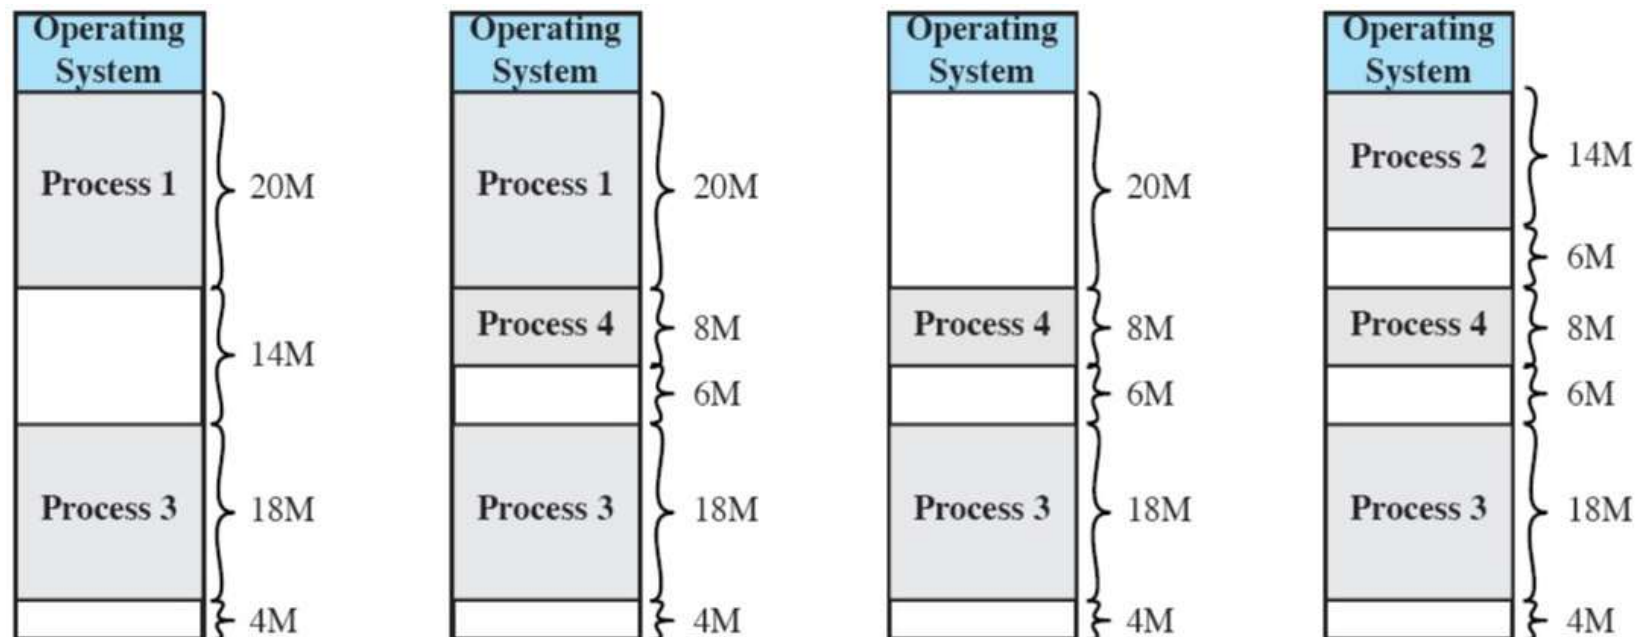
Single Queue

- Whenever a partition becomes free, a process is selected
 - Closest to the front of the queue
 - Smaller than the partition size
- Undesirable to waste a large partition for smaller process
 - Search the queue, find the largest process that fits it
- Unfair for smaller processes
 - A process may not be skipped more than k times




Contiguous Allocation: Dynamic Partitioning

- Process is allocated exactly as much memory as required
- Eventually holes in memory: **External fragmentation**
 - Total memory space exists to satisfy a request
 - But it is not contiguous
- Must use **compaction** to shift processes (**defragmentation**)




Memory Management With Bitmaps

- Memory is divided up into allocation units
 - Few words
 - Or several kilobytes



Memory Management With Bitmaps

- Size of allocation unit is important
- Smaller allocation unit
 - Larger bitmap required
- Larger allocation unit
 - Smaller bitmap required
 - More memory will be wasted
 - If the process size is not an exact multiple of the allocation unit




Memory Management With Bitmaps

- To bring a k unit process in memory
- Search for a k run consecutive 0 bits in the map
- Search can be slow
- Since, k run may cross word boundaries

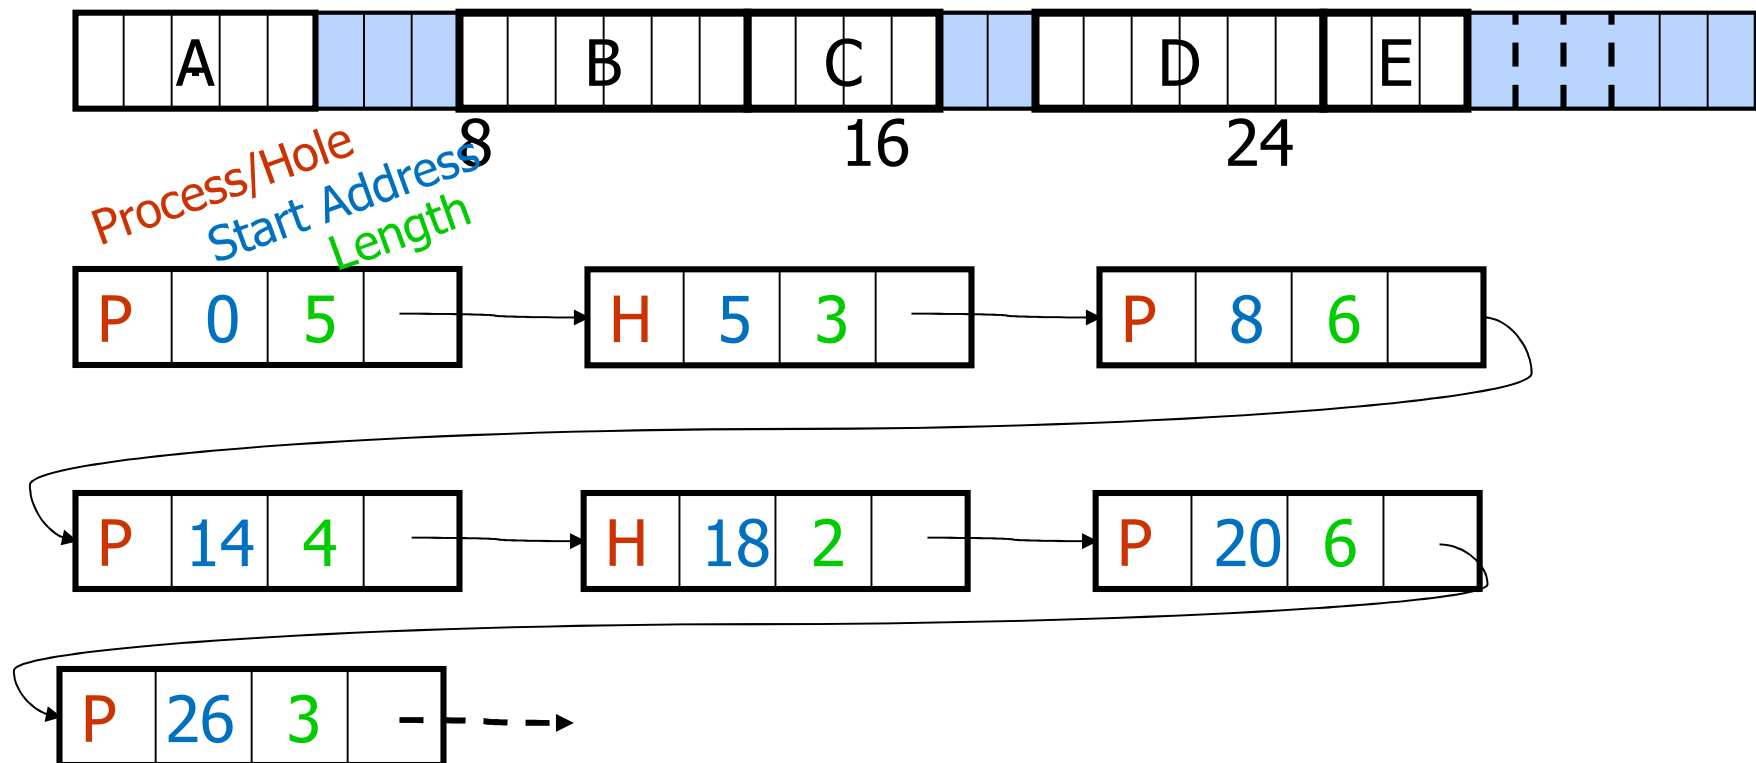
Find run of length = 3

Find run of length = 4



Memory Management with Linked Lists


- Linked list of memory segments
 - Free segments → Holes
 - Allocated segments → Processes



Memory Management with Linked Lists

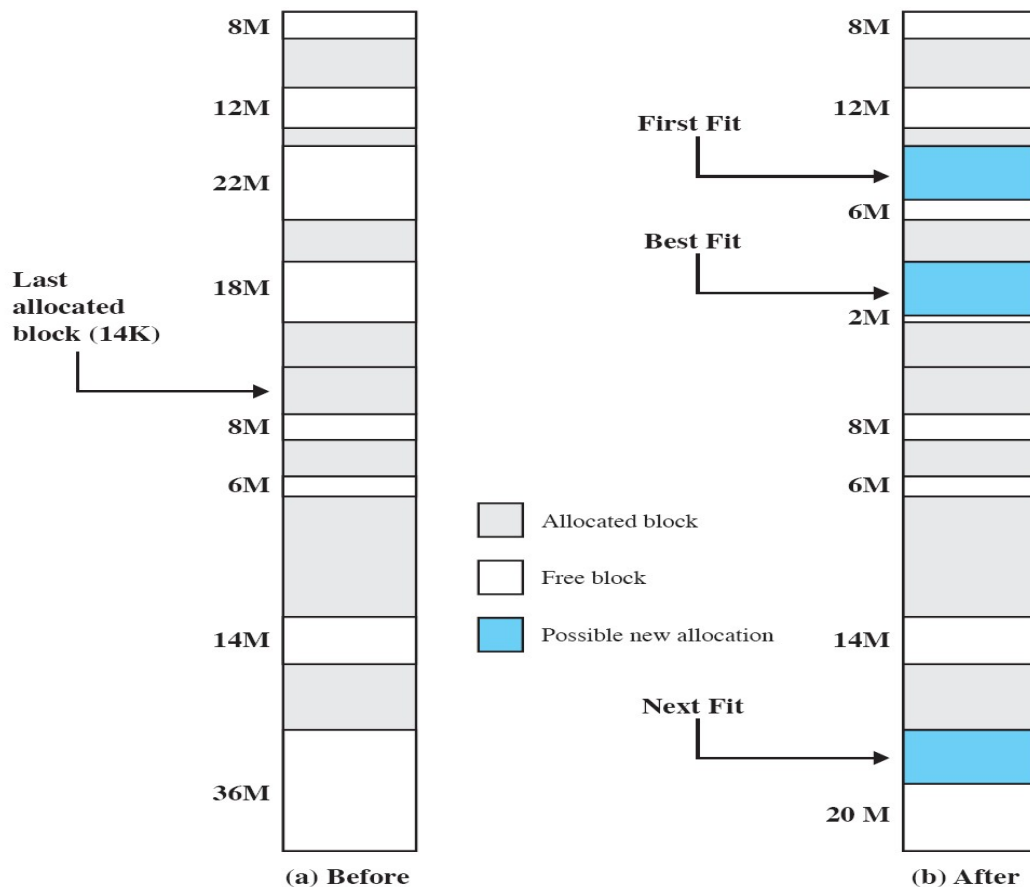
- Segment list is sorted by addresses
- Sorting helps in updating the list, when a process is swapped out or exits
- A process usually has two neighbors

Updating the list requires



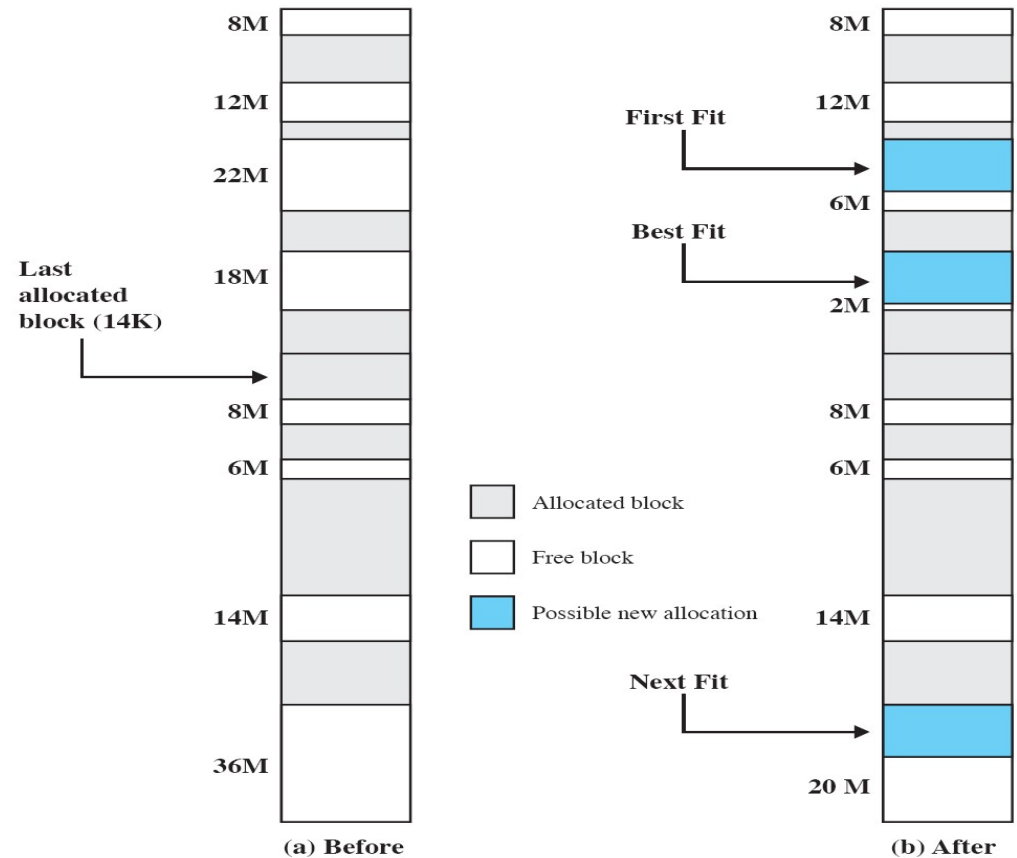
Dynamic Partitioning: Placement Algorithms

- **First-fit:** Use the first block that is big enough
 - Simple and fast



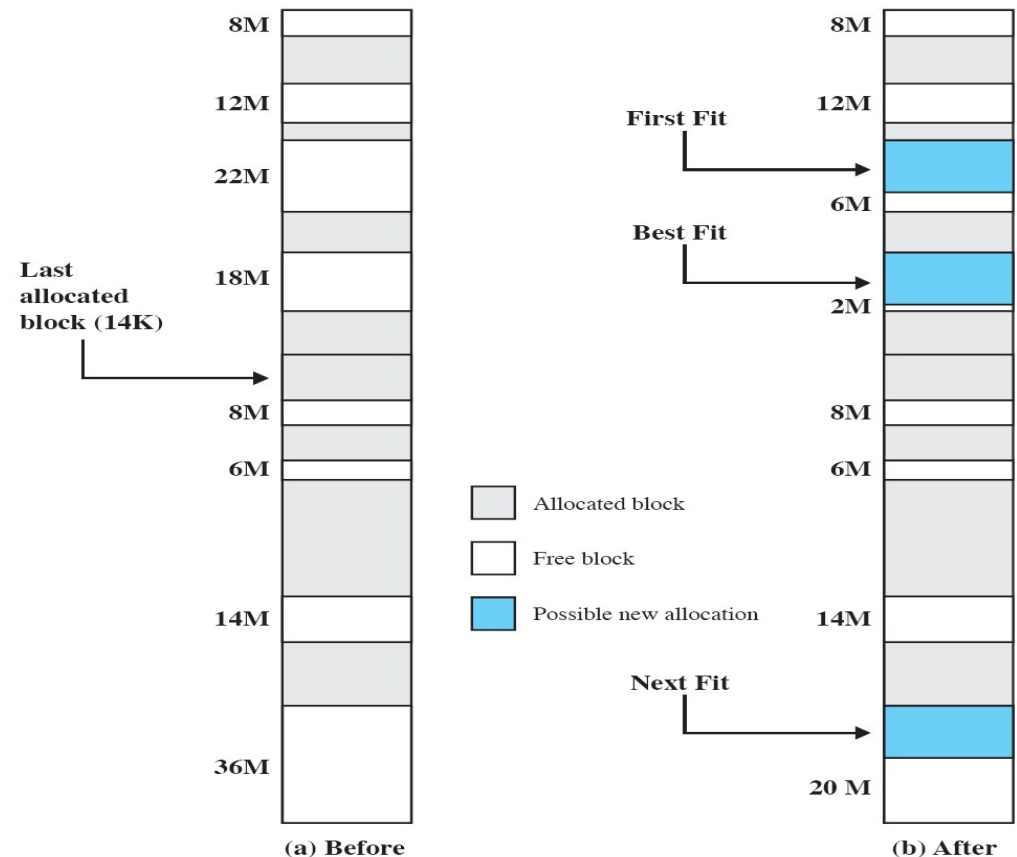
Dynamic Partitioning: Placement Algorithms

- **Next-fit:** Use the next block that is big enough
 - Tends to eat-up the large block at the end of the memory
 - Gives no better performance than First-fit



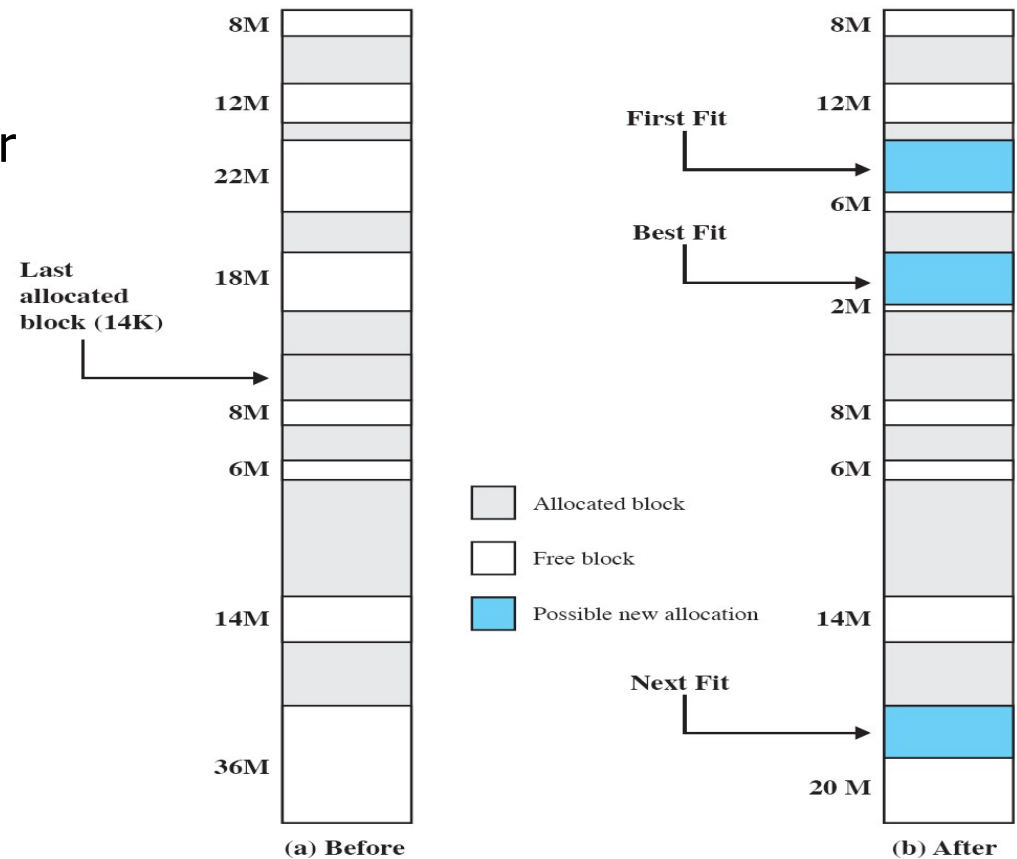
Dynamic Partitioning: Placement Algorithms

- **Best-fit:** Use the smallest block that is big enough
 - Must search entire list (unless free blocks are ordered by size)
 - Produces the smallest leftover hole
 - Tends to fill up memory, with tiny useless holes
 - First Fit generates larger holes on the average



Dynamic Partitioning: Placement Algorithms

- **Worst-fit:** Use the largest block
 - Must also search entire list
 - Produces the largest leftover hole...
 - ... but eats-up big blocks

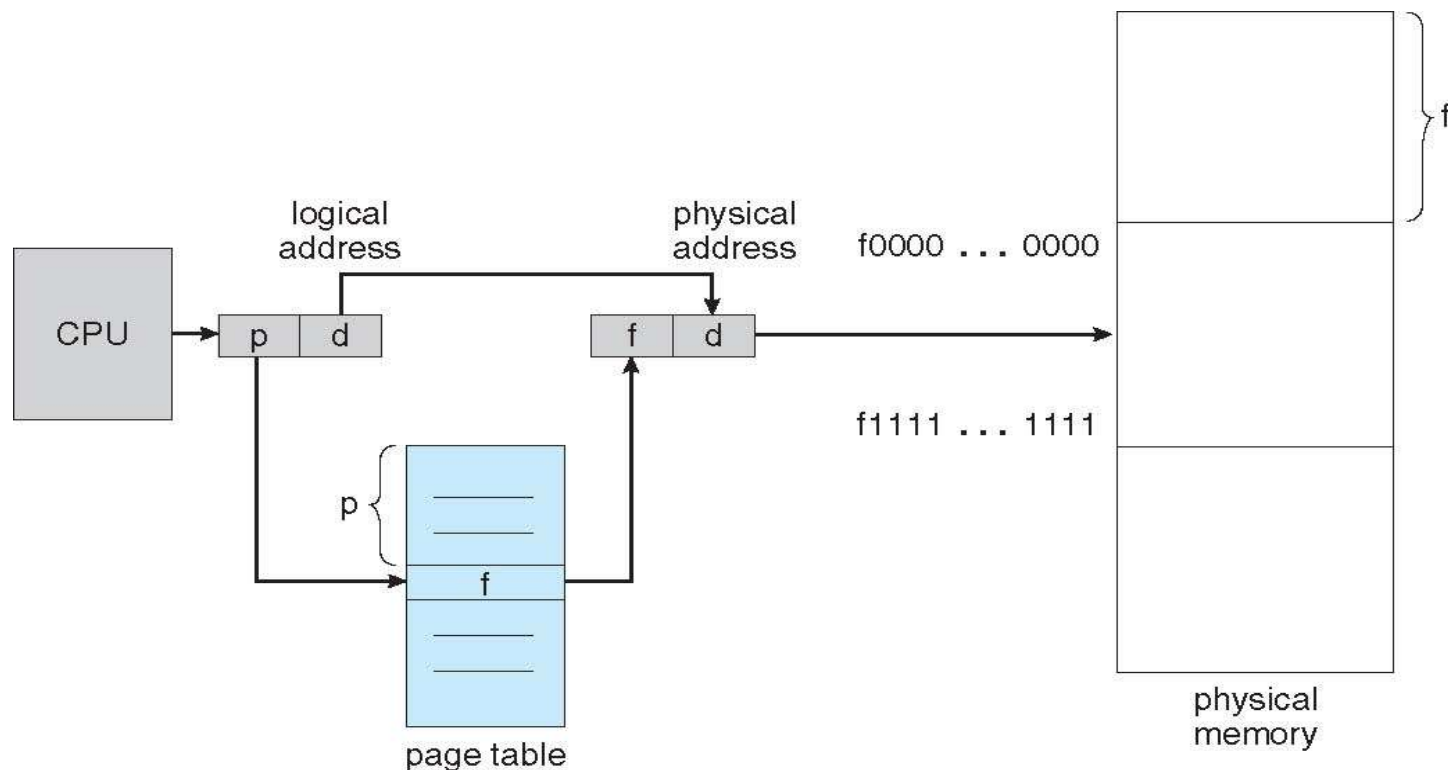


Memory Management with Linked Lists

- Search time of all four placement algorithms can be improved by
 - Keeping separate lists for Process and Holes
 - While allocating memory only have to search the lists of holes
 - The Hole list can be kept sorted by size
 - As soon as a hole that fits is found, no more searching is required
- Drawback
 - Problem while de-allocating memory
 - A node from the process list has to be inserted in the hole list

To Avoid External Fragmentation: Paging

- Partition memory into small equal-size chunks (**frames**)
- Divide each process into the same size chunks (**pages**)
- OS maintains a **page table** for each process
 - Contains the frame location for each page in the process
 - Memory address = (page number, offset within page)



Page Example: Allocation of Frames

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(d) Load Process C

Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B


Main memory

0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load Process D

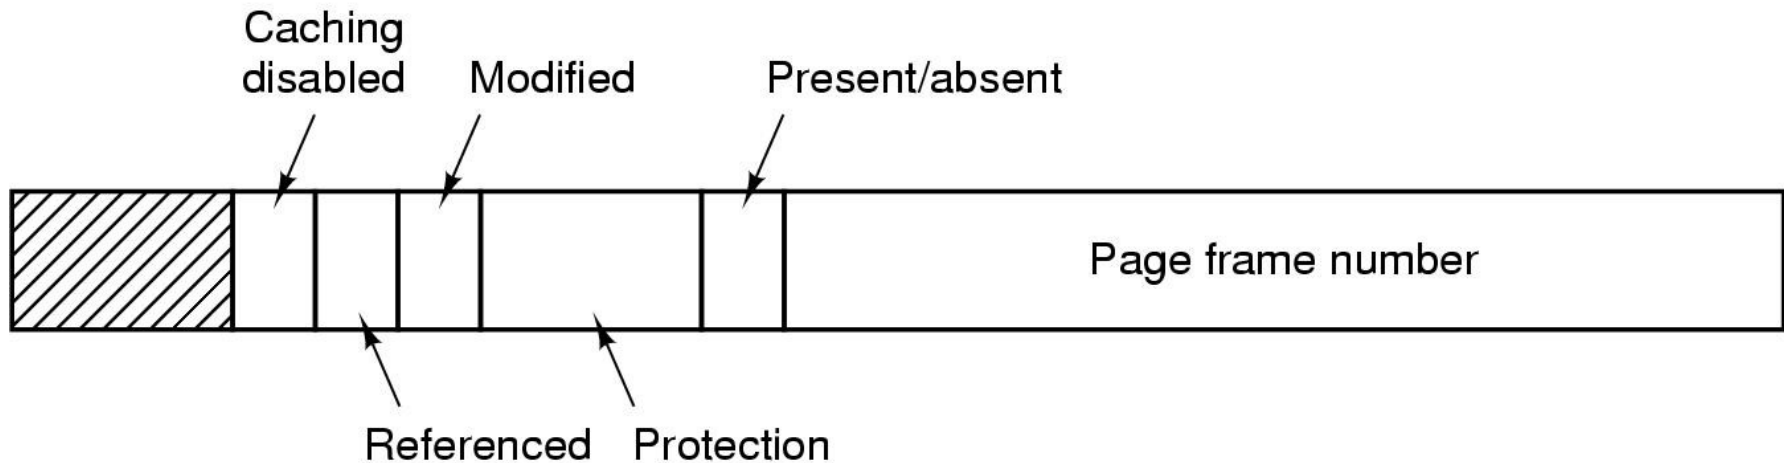
Paging Example

- Pages of a process form the logical memory
- Maintain for each process a page table



- **Question:** Do we avoid fragmentation completely?

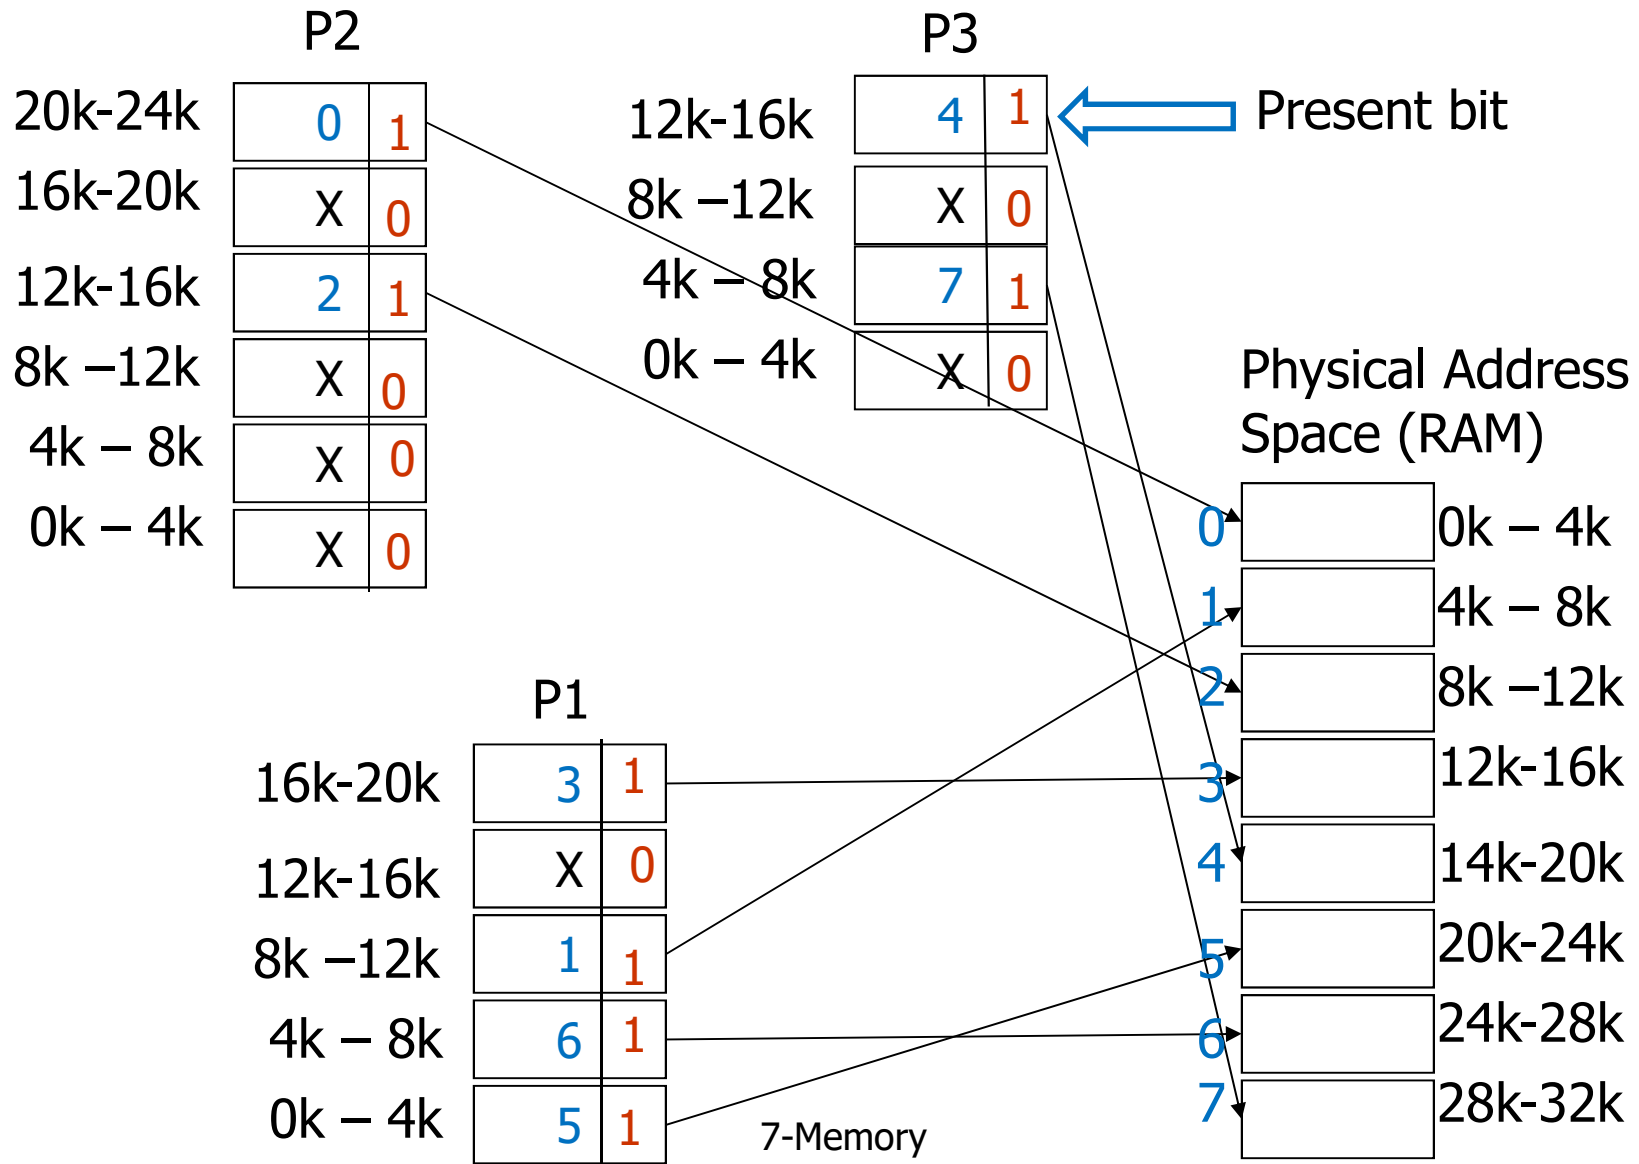
Typical Page Table Entry



Exact layout is hardware dependent

- **Present/Valid:** value valid to be used, otherwise page fault
- **Protection:** defines accesses which are permitted
- **Referenced:** help to decide which frames can be reclaimed
- **Caching:** Needed if value is read from I/O device
- **Modified:** Needed if page frames are reclaimed
 - **Dirty:** Needs to be written to disk first


Paging Example



Address Translation Scheme

Address generated by CPU is divided into two parts

- Page number (p) – used as an **index into a page table** which contains base address of each page in physical memory
- Page offset (d) – combined with base address to define the **physical memory address** that is sent to the memory unit



- For given logical address space 2^m and page size 2^n

Address Translation Scheme

- Size of logical address space = 2^m
 - # of pages $(2^{m-n}) \times$ page size (2^n)
- Example: 16 X 4096
 - $2^m = 2^4 \times 2^{12}$
 - $2^m = 2^{16}$
 - $m = 16$ bits

page number	page offset
p	d
m - n	n

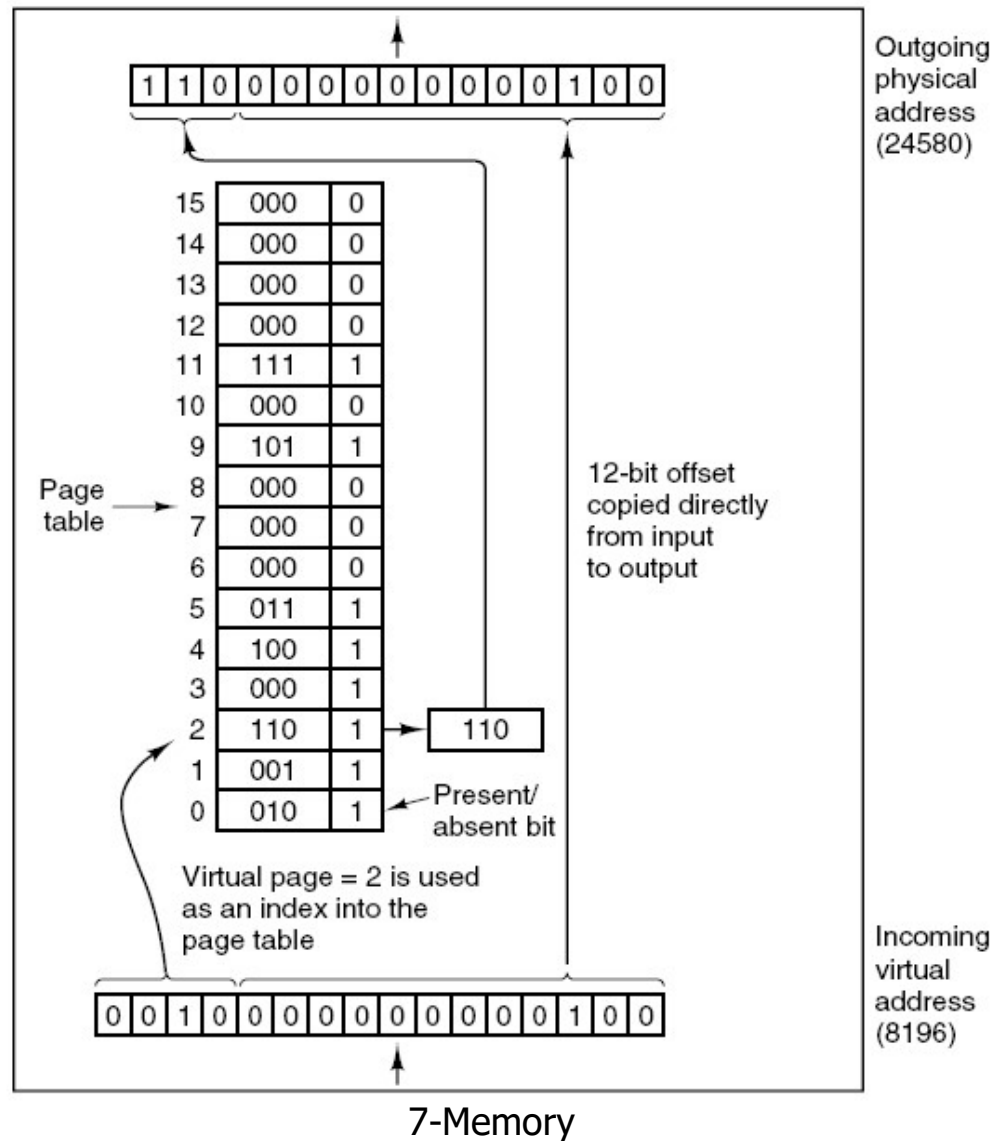
Address Translation Scheme: Example

Given $m=16$ and $n=12$

- Page number = $m-n$ bits = Leftmost 4 bits
 - 4 bit page number can represent 16 pages
- Page offset = n bits = Rightmost 12 bits
 - 12 bit offset can address all 4096 bytes within a page (Frame size) !!
- Consider logical address 8196 = 0010000000000100
 - Page number is 2 and offset is 4

page number	page offset
p	d
$m - n$	n

Address Translation Scheme: Example



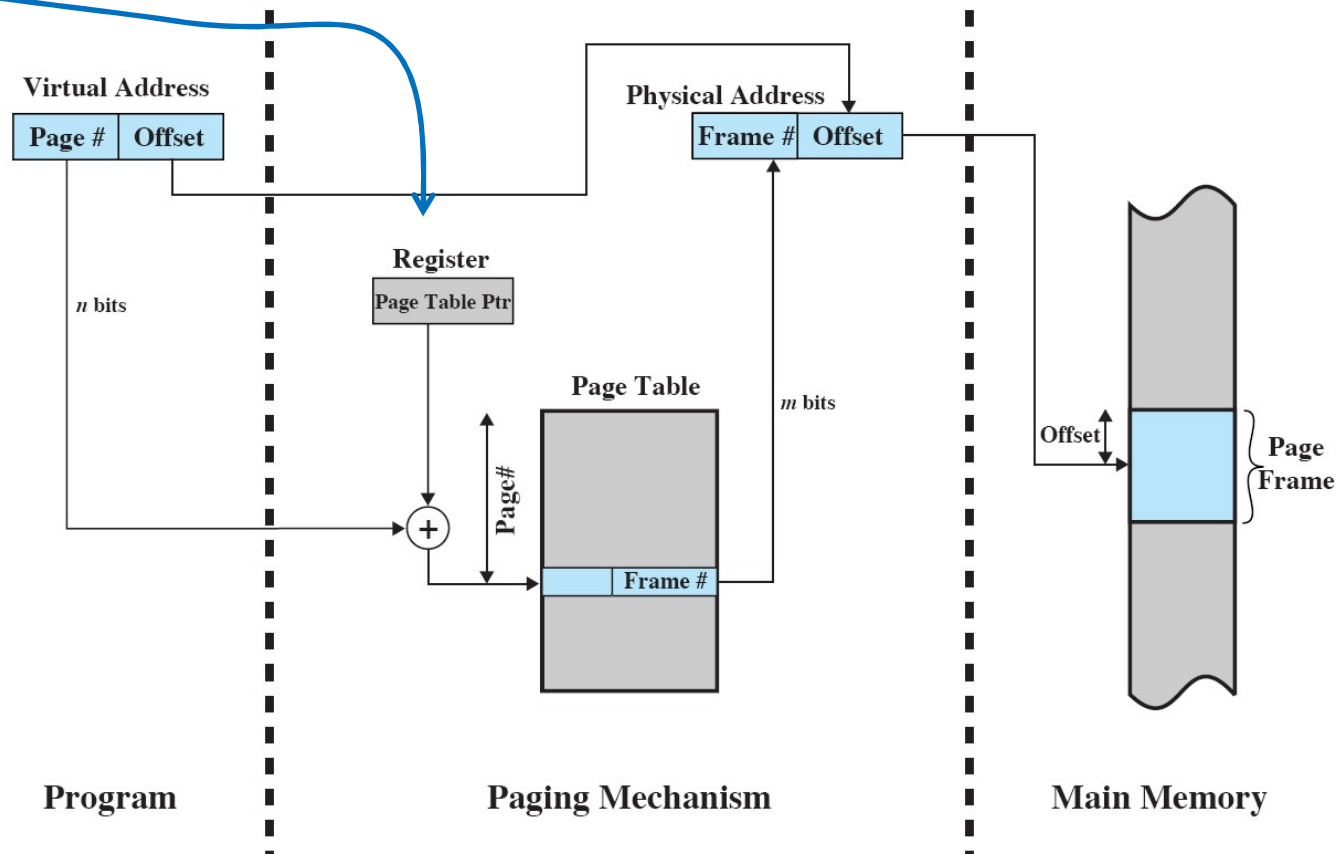
Address Translation Scheme: Example

- Consider an address of $m = 16$ and $n = 10$ bits
 - Leftmost $m-n$ bits are page number
 - Rightmost n bits are the offset
- Given a logical address 0000010111011110
 - Page number is 1 and offset 478
- Suppose this page is residing in main memory frame 6 (000110)
 - Physical address is frame number 6, offset 478

page number	page offset
p	d
$m - n$	n

Implementation of Page Table – Main Memory

- Page table is kept in main memory
 - Page-table base register (PTBR) points to the page table
 - Page-table length register (PRLR) indicates size of the page table
- Each program reference to memory → 2 memory accesses

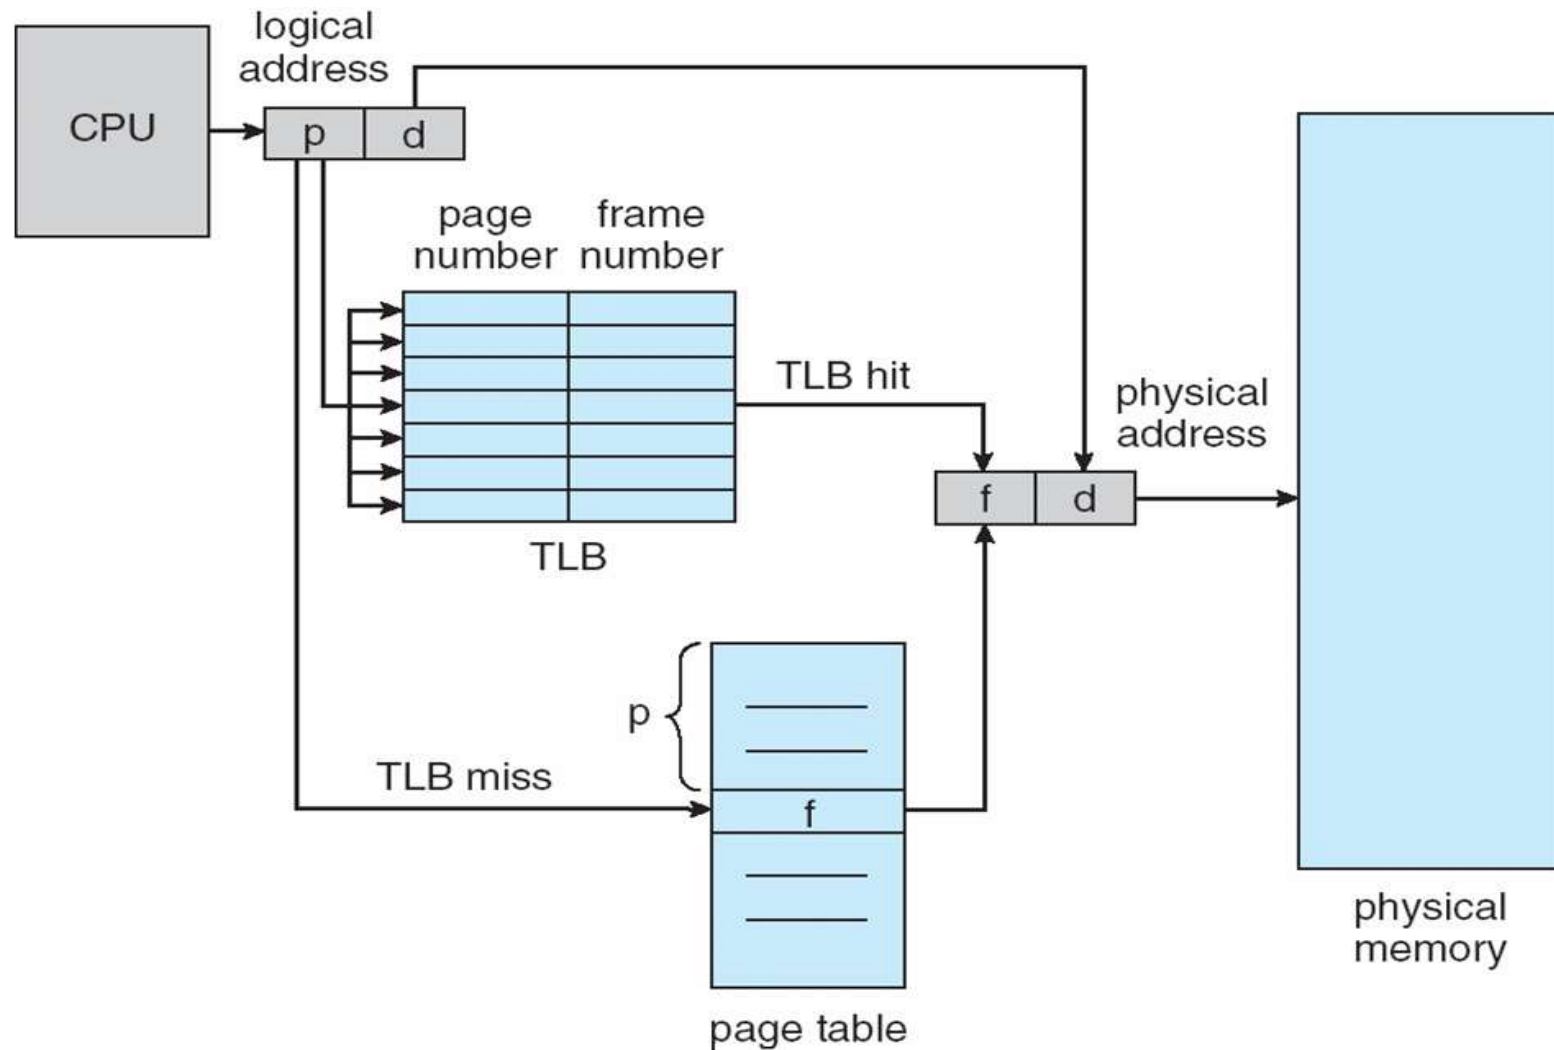


Implementation of Page Table – Associative Memory

- **Translation Lookaside Buffers (TLBs)**
 - Special fast-lookup hardware cache
 - Parallel search (cache for page table)
 - Resides usually inside the MMU
- **Address translation (P, O)**
 - If P is in associative register (hit)
 - Get frame # from TLB
 - Else get frame # from page table in m

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75

Paging Hardware With TLB



Access Time With Cache Registers

- Associative Lookup = ε time units (fraction of microsecond)
 - Can be $< 10\%$ of memory access time
- **Hit ratio** ($= \alpha$)
 - Percentage of times a page number is found in associative registers
- Effective Access Time = $(1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) = 2 + \varepsilon - \alpha$
 - Assume memory cycle time is 1 microsecond
- Example:
 - Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - Effective access time = $0.80 * 120 + 0.20 * 220 = 140 \text{ ns}$

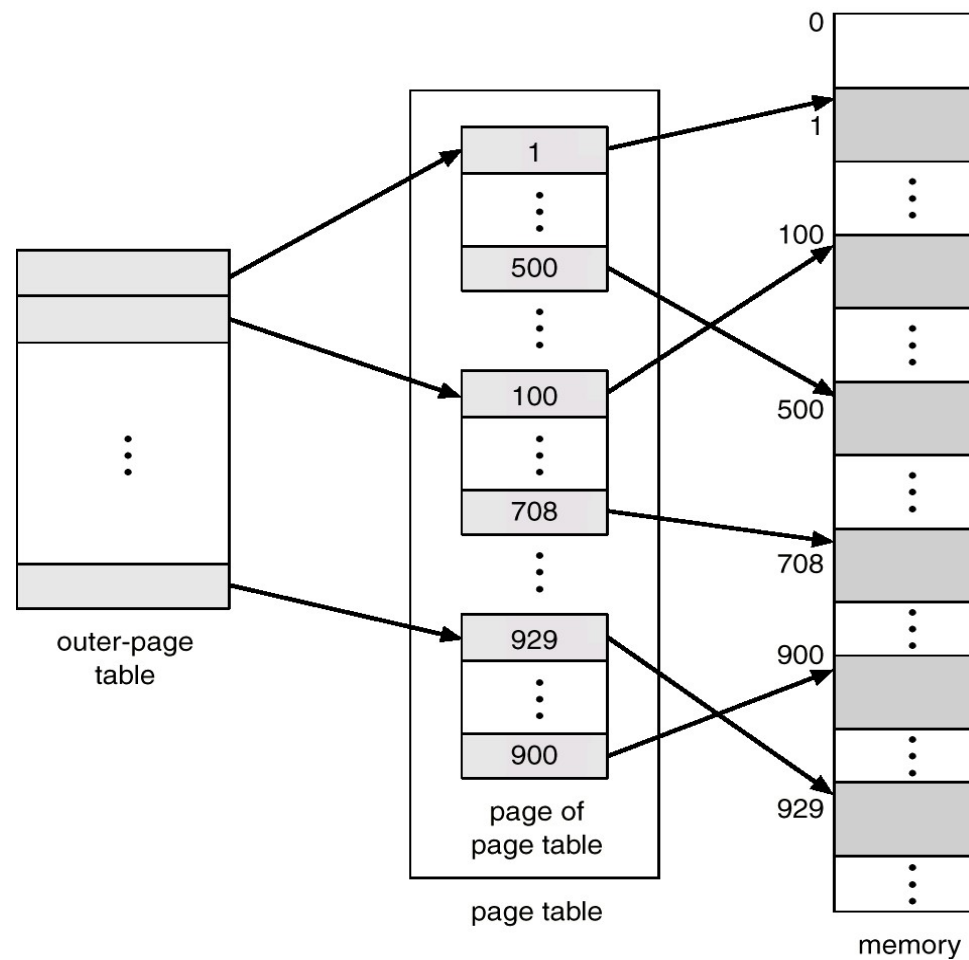
Single-Level Page Table Limitation

- Consider a system with
 - 32-bit logical address space
 - 4KB frame size
- How many entries will be in the (single-level) page table?
 - 2^{20} (1 million) entries
- Each entry of page table consists of 4 bytes
 - What will be the size of (single-level) page table?
 - 4MB

page number	page offset
p	d
m - n	n

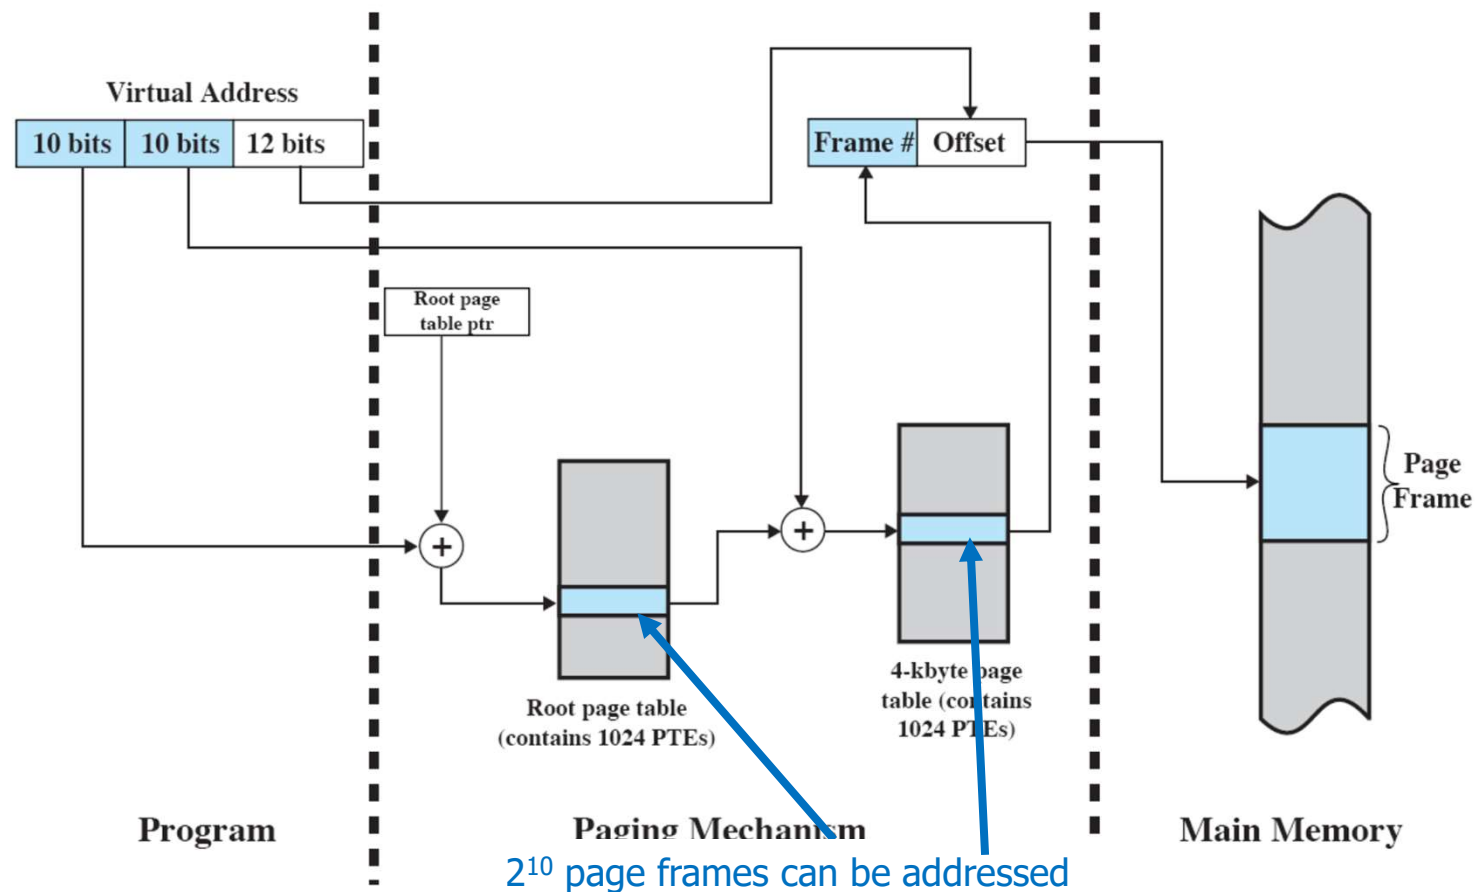
Two-Level Page-Table Scheme

- Page-table may be large, i.e., occupy several pages/frames itself



Address Translation in Two-Level Memory

- Assume byte level addressing
 - 2^{32} bytes = 2^{22} kbytes = 2^{12} Mbytes = 4Gbytes can be supported



Multi-Level Page Table

- Even two-level paging scheme might not be sufficient
- Consider a system with 64-bit logical address space
 - Frame size is 4 KB (2^{12})
- How many entries in single-level page table?
 - Page table has 2^{52} entries
- How many entries in two-level page table

outer page	inner page	page offset
p_1	p_2	d
42	10	12

Multi-Level Page Table

- How many entries in three level page table?

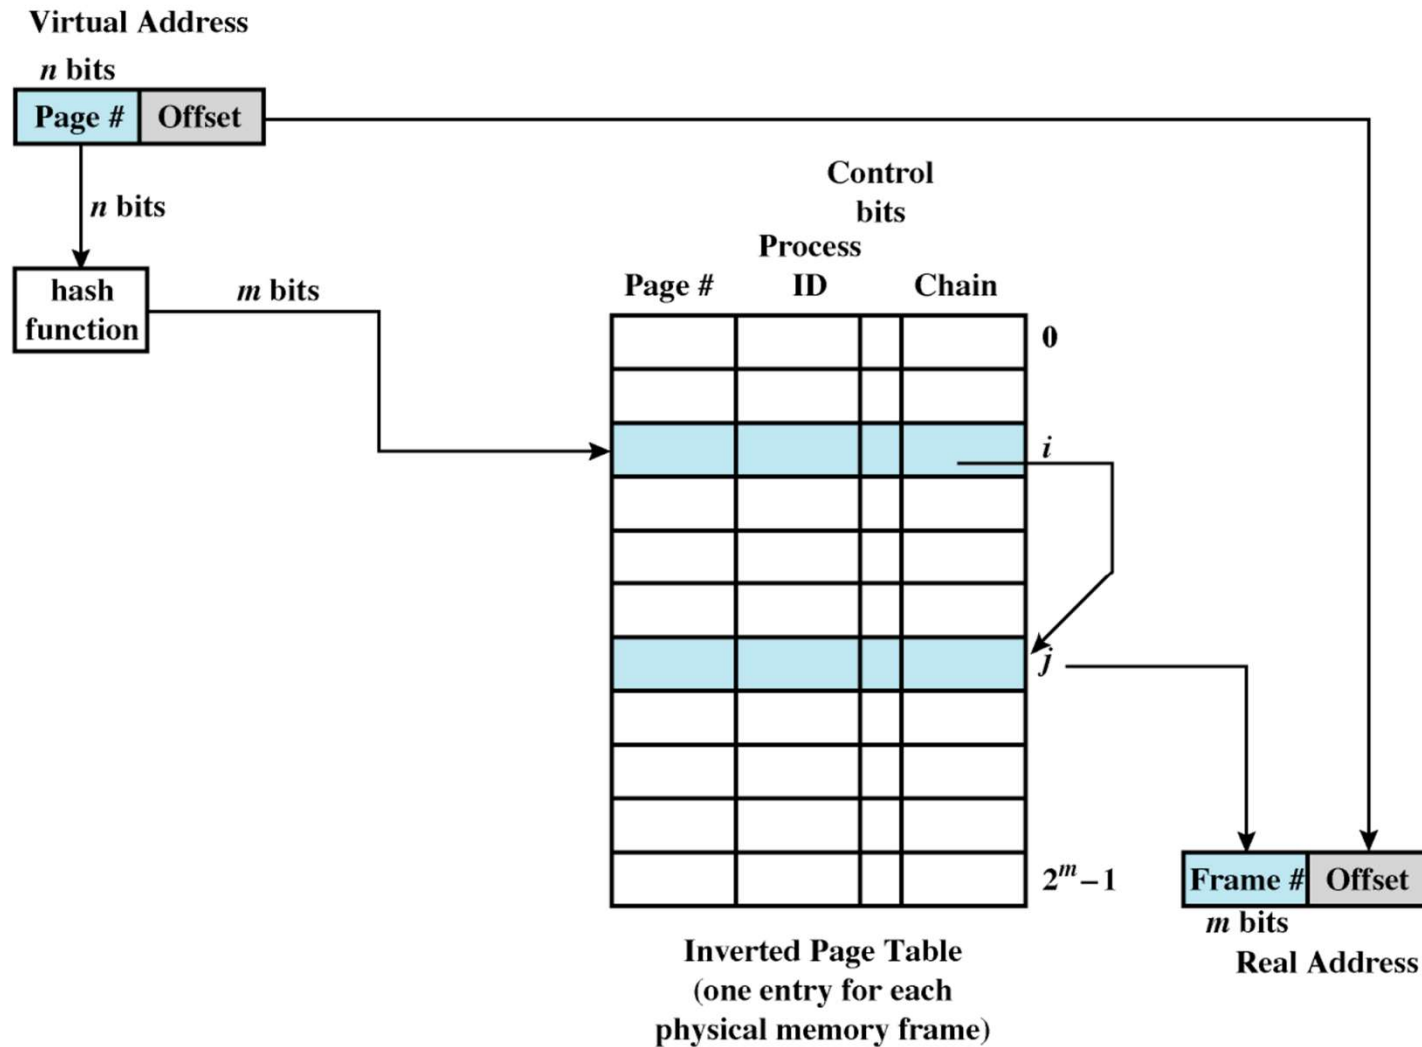
2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

- How many levels are require such that each page table fits within a single frame?

Inverted Page Table

- One entry for each page frame of real memory
- Entry consists of
 - Virtual address of the page stored in that real memory location,
 - With information about the process that owns that page
- Decreases memory needed to store each page table
- But increase in time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries

Inverted Page Table Example



Chapter Overview

- Background
- Memory allocation schemes
 - Fixed partition
 - Dynamic partition
 - Paging
- Virtual memory

Execution of a Program: Virtual Memory Concept

Main memory = cache of the disk space

- Operating system brings into main memory a few pieces of the program
- **Resident set** - portion of process that is in main memory
- Whenever an address is needed that is not in main memory
 - Generation of a **page-fault interrupt**
 - OS places the process in blocking state and issues a disk IO request
 - Another process is dispatched

Valid-Invalid (Present/Absent) Bit

- With each page table entry a valid–invalid bit is associated
 - 1 → in-memory
 - 0 → not-in-memory (Initially 0)

Frame #	valid-invalid bit
	1
	1
	1
	1
	0
⋮	
	0
	0

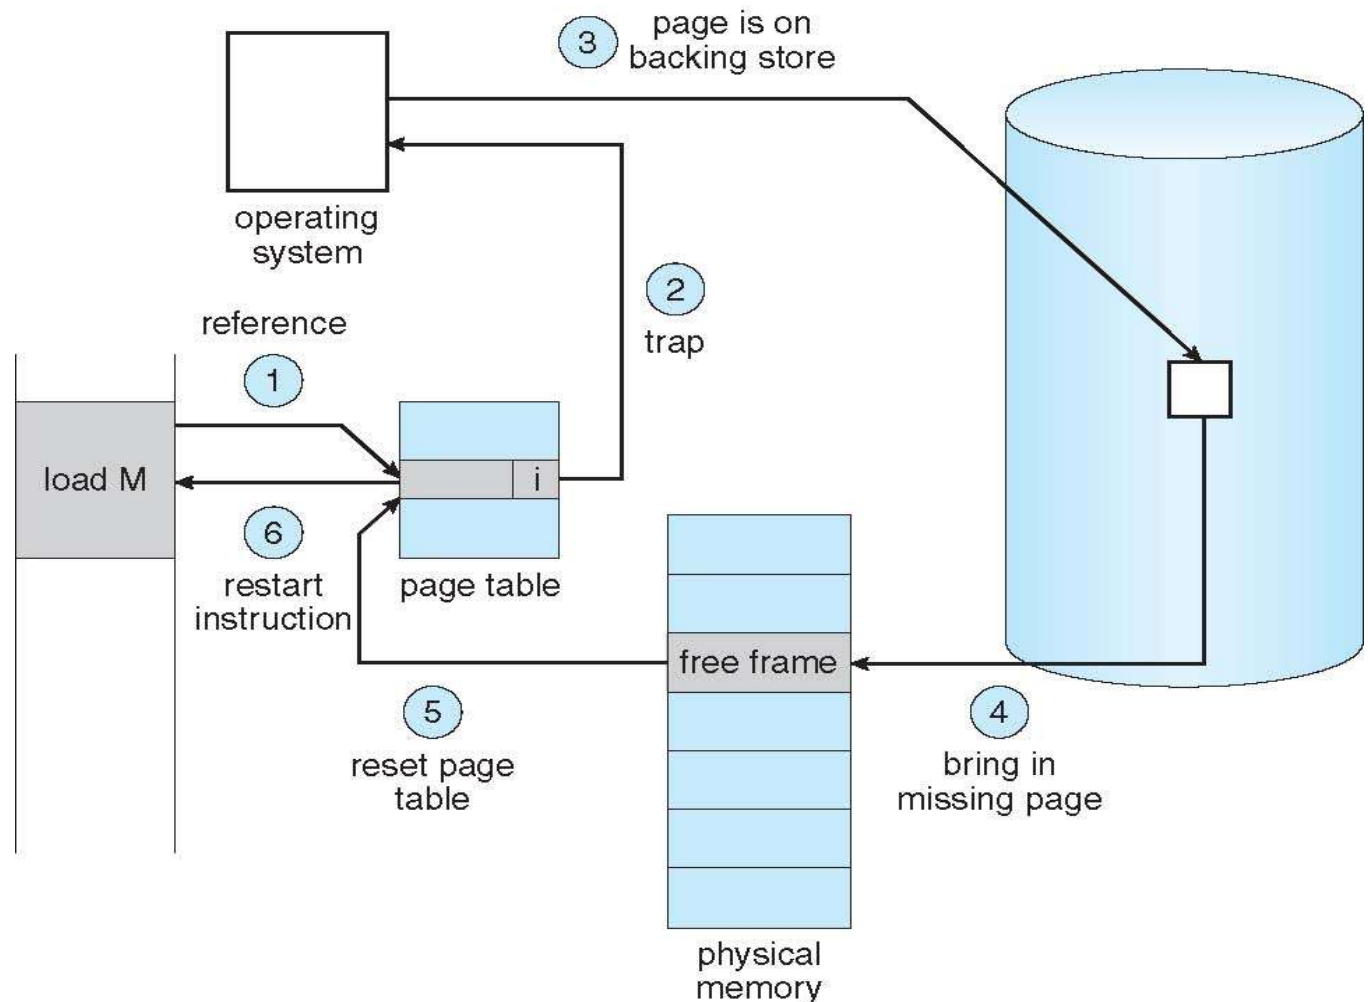
Page table

- During address translation, if valid-invalid bit in page table entry is 0
 - Page fault interrupt to OS

Page Fault and Address Translation

In response to page-fault, paging hardware issue trap to OS

- Check **validity** of memory reference
- Get empty frame (**swap out** that page?)
- **Swap in** page into frame
- **Reset** tables, validation bit
- **Restart** instruction



If There is No Free Frame?

Page replacement

- Goal: Algorithm which will result in **minimum number of page faults**
- Page fault forces choice
 - Which page must be removed
 - Make room for incoming page
- Modified page must first be saved
 - Unmodified just overwritten (use **dirty bit** to optimize writes to disk)
- Better not to choose an often used page
 - Will probably need to be brought back in soon

First-In-First-Out (FIFO) Replacement Algorithm

- Can be implemented using a circular buffer
- **Example:** Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - 3 frames

1	4	5	9 page faults
2	1	3	
3	2	4	


- 4 frames

1	5	4	10 page faults
2	1	5	
3	2		
4	3		

- **Belady's Anomaly:** More frames, sometimes more page faults
 - Problem: replaces pages that will be needed soon

Optimal Replacement Algorithm

- Replace page that will not be used for longest period of time
- 4 frames example
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



- How can we know this information?
- Standard for measuring how well other algorithms perform

Least Recently Used (LRU) Replacement Algorithm

- **Idea:** Replace the page that has not been referenced for the longest time
 - By the **principle of locality**, this should be the page least likely to be referenced in the near future
- **Implementation**
 - Tag each page with the time of last reference
 - Use a stack
- **Problem**
 - High overhead
 - **OS kernel** involvement at **every memory reference!!!**
 - If hardware support not available

LRU Algorithm


- 4 frame example
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1		5
2		
3	5	4
4	3	


8 page faults

LRU Approximations: Clock/Second Chance


- Uses **reference (use)** bit
 - Initially 0
 - When page is referenced, set to 1 by **HW**
- **To replace a page**
 - The first frame encountered with use bit 0 is replaced
 - During the search for replacement, each use bit set to 1 is changed to 0 by **OS**
- Note: if all bits set → **FIFO**



Example Clock Policy



Example Clock Policy




F = page fault occurring after the frame allocation is initially filled

Comparison of Basic Replacement Algorithms

Page address stream	2	3	2	1	5	2	4	5	3	2	5	2																																				
OPT	<table><tr><td>2</td></tr><tr><td></td></tr><tr><td></td></tr></table>	2			<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table>	2	3	1	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table> F	2	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	2	3	5	<table><tr><td>4</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table> F	4	3	5	<table><tr><td>4</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	4	3	5	<table><tr><td>4</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	4	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table> F	2	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	2	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	2	3	5
2																																																
2																																																
3																																																
2																																																
3																																																
2																																																
3																																																
1																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
LRU	<table><tr><td>2</td></tr><tr><td></td></tr><tr><td></td></tr></table>	2			<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table>	2	3	1	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>1</td></tr></table> F	2	5	1	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>1</td></tr></table>	2	5	1	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table> F	2	5	4	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table>	2	5	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table> F	3	5	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table> F	3	5	2	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table>	3	5	2	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table>	3	5	2
2																																																
2																																																
3																																																
2																																																
3																																																
2																																																
3																																																
1																																																
2																																																
5																																																
1																																																
2																																																
5																																																
1																																																
2																																																
5																																																
4																																																
2																																																
5																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
FIFO	<table><tr><td>2</td></tr><tr><td></td></tr><tr><td></td></tr></table>	2			<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table>	2	3	1	<table><tr><td>5</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	5	3	1	<table><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>1</td></tr></table> F	5	2	1	<table><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table> F	5	2	4	<table><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table>	5	2	4	<table><tr><td>3</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table> F	3	2	4	<table><tr><td>3</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table>	3	2	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table> F	3	5	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table> F	3	5	2
2																																																
2																																																
3																																																
2																																																
3																																																
2																																																
3																																																
1																																																
5																																																
3																																																
1																																																
5																																																
2																																																
1																																																
5																																																
2																																																
4																																																
5																																																
2																																																
4																																																
3																																																
2																																																
4																																																
3																																																
2																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
2																																																
CLOCK	<table><tr><td>2*</td></tr><tr><td></td></tr><tr><td></td></tr></table>	2*			<table><tr><td>2*</td></tr><tr><td>3*</td></tr><tr><td></td></tr></table>	2*	3*		<table><tr><td>2*</td></tr><tr><td>3*</td></tr><tr><td></td></tr></table>	2*	3*		<table><tr><td>2*</td></tr><tr><td>3*</td></tr><tr><td>1*</td></tr></table>	2*	3*	1*	<table><tr><td>5*</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	5*	3	1	<table><tr><td>5*</td></tr><tr><td>2*</td></tr><tr><td>1</td></tr></table> F	5*	2*	1	<table><tr><td>5*</td></tr><tr><td>2*</td></tr><tr><td>4*</td></tr></table> F	5*	2*	4*	<table><tr><td>5*</td></tr><tr><td>2*</td></tr><tr><td>4*</td></tr></table>	5*	2*	4*	<table><tr><td>3*</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table> F	3*	2	4	<table><tr><td>3*</td></tr><tr><td>2*</td></tr><tr><td>4</td></tr></table>	3*	2*	4	<table><tr><td>3*</td></tr><tr><td>2</td></tr><tr><td>5*</td></tr></table> F	3*	2	5*	<table><tr><td>3*</td></tr><tr><td>2*</td></tr><tr><td>5*</td></tr></table>	3*	2*	5*
2*																																																
2*																																																
3*																																																
2*																																																
3*																																																
2*																																																
3*																																																
1*																																																
5*																																																
3																																																
1																																																
5*																																																
2*																																																
1																																																
5*																																																
2*																																																
4*																																																
5*																																																
2*																																																
4*																																																
3*																																																
2																																																
4																																																
3*																																																
2*																																																
4																																																
3*																																																
2																																																
5*																																																
3*																																																
2*																																																
5*																																																

F = page fault occurring after the frame allocation is initially filled

Comparison of Basic Replacement Algorithms



Enhancement of Clock/Second Chance

- Clock algorithm can be made more powerful by using more bits
 - Use **reference bit** (R) and **modify bit** (M) as an ordered pair
- With R and M bits, the following four classes are possible
 - **Class 0**: Neither recently used nor modified (R=0, M=0)
 - Best page to replace
 - **Class 1**: Not recently used but modified (R=0, M=1)
 - The page will need to be written out before replacement
 - **Class 2**: Recently used but clean (R=1, M=0)
 - Probably page will be used again soon
 - **Class 3**: Recently used and modified (R=1, M=1)
 - Probably page will be need soon and disk I/O is required to write page
- Replacement strategy
 - **Replace** the first page encountered in the **lowest non empty class**

Enhancement of Clock/Second Chance

- When a process starts
 - All its page entries are marked as not referenced and not modified
- When a page is referenced
 - The R (reference) bit is set
- When the page is later written to
 - The M (modified) bit is set
- **Difference** to clock algorithm
 - **Preference** are given to **pages that have been modified** to reduce I/O

Any Question So Far?

