

Dependency Analysis

Dependencies Affect Parallelization!

- Data must be produced and consumed in the *correct order*
- Simple example of data dependence:

S_1 $PI = 3.14$

S_2 $R = 5.0$

S_3 $AREA = PI * R ** 2$

- Statement S_3 cannot be moved before either S_1 or S_2 without compromising correct results

Dependencies Affect Parallelization!

- Dependency analysis identifies the units of code that are semantically independent and hence executable in parallel.
- Modern compilers do provide some degree of optimization.

Categorizing Dependencies

- There are 3 categories of program dependencies:
 - Control Flow Dependencies
 - Data Flow Dependencies
 - Loop Dependencies

Categorizing Dependencies

- ***Control Dependency:*** How different program instructions affect each other.
- ***Data Dependency:*** How different pieces of data are related and affect each other during execution.

Control Dependencies

- Every program has a well-defined flow of control.
- This flow can be affected by several kinds of operations such as Loops and branches
- Control flow dependencies are those statements that are only executed if the program flow-of-control allows its execution.

Control Flow Dependency

Example:

```
S1    if (x > 100) {  
S2        x = 100;  
        }  
S3    y = x;
```

- The value of **y** depends on what the condition **(x > 100)** evaluates to.

S2 is control-flow dependent on S1 written:

$S1 \delta^c S2$

Data Flow Dependency


- Data flow dependencies arise from reading or writing to the same variable or resource.
- There are 4-types of data dependencies:
 - Flow, Anti, Output and Input

Flow (True) Dependence

- A statement S_j is flow dependent on S_i if and only if S_i modifies a resource that S_j requires and S_i precedes it in execution.

Example:

S1	$A = B + C;$
S2	$D = A + 2;$
S3	$E = A * 3;$



S2 and S3 are therefore flow dependent on S1 written:

$S2 \delta_f S1$

$S3 \delta_f S1$

Anti-Dependence

- A statement S_j is anti-dependent on S_i if and only if S_j modifies a resource that S_i requires and S_i precedes it in execution.

Example:

S1 $A = B + C;$

S2 $B = D / 2;$

S2 is therefore anti-dependent on S1

$S1 \delta^a S2$

Output-Dependence

- A statement S_j is output dependent on S_i if and only if S_i and S_j modify the same resource and S_i precedes S_j in execution.

Example:

S1 $X = B + C;$

S2 $X = D / 2;$

S2 is therefore output-dependent on S1

$S1 \delta^0 S2$

Input-Dependence

- A statement S_j is input dependent on S_i if and only if S_i and S_j read the same resource and S_i precedes S_j in execution.
- This type of dependence does not prohibit reordering of statements

Example:

S1 $X = B + C;$

S2 $Y = B / 2;$

S2 is therefore input-dependent on S1

$S1 \delta^i S2$

An example

```
int main(int argc, char** argv) {  
    int x,y;  
    int z=10;  
    x=TaskX();           //s1  
    y=TaskY()+z;         //s2  
    y=x-5;              //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

```
int TaskX() {  
    sleep(15);  
    return 15;}  
int TaskY() {  
    sleep(5);  
    return 5; }  
int TaskZ() {  
    sleep(5);  
    return 5; }
```

What would be the output?

How long would it take to produce output?

The output

```
ez@master: ~  
$ gcc serial.c -o serial  
$ time ./serial  
x: 15 y: 10 z: 5  
  
real    0m25.004s  
user    0m0.000s  
sys     0m0.000s  
$ █
```

How it can be improved? We know how to implement the parallel version using omp or mpi, however ...

The parallel version

- Let us assume we intend to use four processors for improved performance
- Ideally, each statement can run independently and we have speedup and efficiency of 4 and 1 respectively.
- Here is a possible (not optimized) implementation.
 - The processes 1-4, execute the individual statements, S1-S4 and send the results to process 0
 - The process 0 receives the values, updates the variables and displays the results
 - The objective is to mimic the output of serial code, just a rough implementation

The parallel version

```
switch(myrank){  
  case 1:  
    x=TaskX(); //s1  
    MPI_Send(&x,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    break;  
  case 2:  
    y=TaskY()+z; //s2  
    MPI_Send(&y,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    break;  
  case 3:  
    ...  
}
```


The parallel version

```
if (myrank == 0) {  
    MPI_Status status;  
    int recvdValue, i;  
    for (i=1; i<5; i++) {  
        MPI_Recv(&recvdValue, 1, MPI_INT,  
                 MPI_ANY_SOURCE, 787, MPI_COMM_WORLD, &status);  
        switch(status.MPI_SOURCE){  
            case 1:  
                x=recvdValue;  
                break;  
            case 2:  
                y=recvdValue;  
  
                ...  
        }  
    }  
}
```

The output

```
ez@master: ~ x ez@master
$ gcc serial.c -o serial
$ time ./serial
x: 15 y: 10 z: 5

real    0m25.002s
user    0m0.000s
sys     0m0.000s
$
$
$
$ mpicc parallel.c -o parallel
$ time mpiexec -n 5 -f machinefile ./parallel
x: 15 y: 15 z: 5

real    0m15.225s
user    0m14.412s
sys     0m0.348s
$ █
```

The output

- Few interesting things to notice
 - The time values are overall just but not the breakdown is
 - The values of the variables are incorrect when compared to the serial code
 - What can be the reason?

Dependencies

```
int main(int argc, char** argv) {  
    int x,y;  
    int z=10;  
    x=TaskX();                //s1  
    y=TaskY()+z;              //s2  
    y=x-5;                    //s3  
    z=TaskZ();                //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

Is there any dependence relation between s3 and s1?
What can be done? Can it be somehow removed?

The parallel version – v2

```
switch(myrank){  
  case 1:  
    x=TaskX();  
    MPI_Send(&x,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    y=x-5;  
    MPI_Send(&y,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    break;  
  case 2:  
    y=TaskY()+z;  
    MPI_Send(&y,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    break;  
  case 3:  
    z=TaskZ();  
    MPI_Send(&z,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    break;  
}
```

The parallel version – v2

```
if (myrank == 0) {  
    MPI_Status status;  
    int recvdValue, i;  
    for (i=1; i<4; i++) {  
        MPI_Recv(&recvdValue, 1, MPI_INT,  
                 MPI_ANY_SOURCE, 787, MPI_COMM_WORLD, &status);  
        switch(status.MPI_SOURCE){  
            case 1:  
                x=recvdValue;  
                MPI_Recv(&recvdValue, 1, MPI_INT, 1, 787,  
                         MPI_COMM_WORLD, &status);  
                y=recvdValue;  
                break;  
            ...  
        }  
    }  
}
```

How about the output?

```
$ mpicc parallel.c -o parallel  
$ time mpiexec -n 4 -f machinefile ./parallel  
x: 15 y: 15 z: 5
```

- We have decreased the level of parallelism but unfortunately, we are still not there!

Dependencies

```
int main(int argc, char** argv) {  
    int x,y;  
    int z=10;  
    x=TaskX();                //s1  
    y=TaskY()+z;              //s2  
    y=x-5;                    //s3  
    z=TaskZ();                //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

Identify the dependence relationship between s2 and s3?

The parallel version – v2

```
switch(myrank){  
  case 1:  
    x=TaskX();  
    MPI_Send(&x,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    y=x-5;  
    MPI_Send(&y,1,MPI_INT,0,787,MPI_COMM_WORLD);  
    break;  
  
    ...
```

What is our mistake then?

The output dependence

- There is an output dependence between s2 and s3.
- There are two options to fix them:
 - Further reduce the level of parallelism
 - As they are regarded as false dependencies, try to remove them!
 - We can use a technique called variable renaming

False Dependencies – Variable renaming

```
int main(int argc, char** argv) {  
    int x,y;  
    int z=10;  
    x=TaskX();           //s1  
    y=TaskY()+z; //s2  
    y=x-5;               //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

False Dependencies – Variable renaming

```
int main(int argc, char** argv) {  
    int x,y;  
    int z=10;  
    x=TaskX();           //s1  
    y=TaskY()+z; //s2  
    y=x-5;               //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

False Dependencies – Variable renaming

```
int main(int argc, char** argv) {  
    int x,y, y2;  
    int z=10;  
    x=TaskX();           //s1  
    y2=TaskY()+z;        //s2  
    y=x-5;               //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

False Dependencies – Variable renaming

```
int main(int argc, char** argv) {  
    int x,y, y2;  
    int z=10;  
    x=TaskX();           //s1  
    y2=TaskY()+z;        //s2  
    y=x-5;               //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

False Dependencies – Variable renaming

```
int main(int argc, char** argv) {  
    int x,y, y2;  
    int z=10;  
    x=TaskX();           //s1  
    y2=TaskY()+z;        //s2  
    y=x-5;               //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```

False Dependencies – Variable renaming

```
int main(int argc, char** argv) {  
    int x,y, y2;  
    int z, z2=10;  
    x=TaskX();           //s1  
    y2=TaskY()+z2;       //s2  
    y=x-5;               //s3  
    z=TaskZ();           //s4  
  
    printf("x: %d y: %d z: %d\n", x, y, z);  
    return 0;  
}
```