




Hadoop MapReduce

Department of Computer Science,
National University of Computer & Emerging Sciences,
Islamabad Campus



An excellent Hadoop Reference



Tom White



MapReduce

- To **speed up** the **processing**, we need to **run parts** of the **program** in **parallel**
- Feasible **to parallelize** the **processing**, **in practice** is **Challenging...**
- Using a **framework** like **Hadoop** to **take care of these issues** is a **great help**



MapReduce

- **MapReduce** is a **programming model** for **data processing**.
- MapReduce works by **breaking** the **processing** into **two phases**:
 - **map** phase and the **reduce** phase.
 - Each phase has **key-value pairs** as **input** and **output**,
 - Type is chosen by the programmer.

Implementation: The programmer specifies two functions: the **map function** and the **reduce function**



Example: Word counting

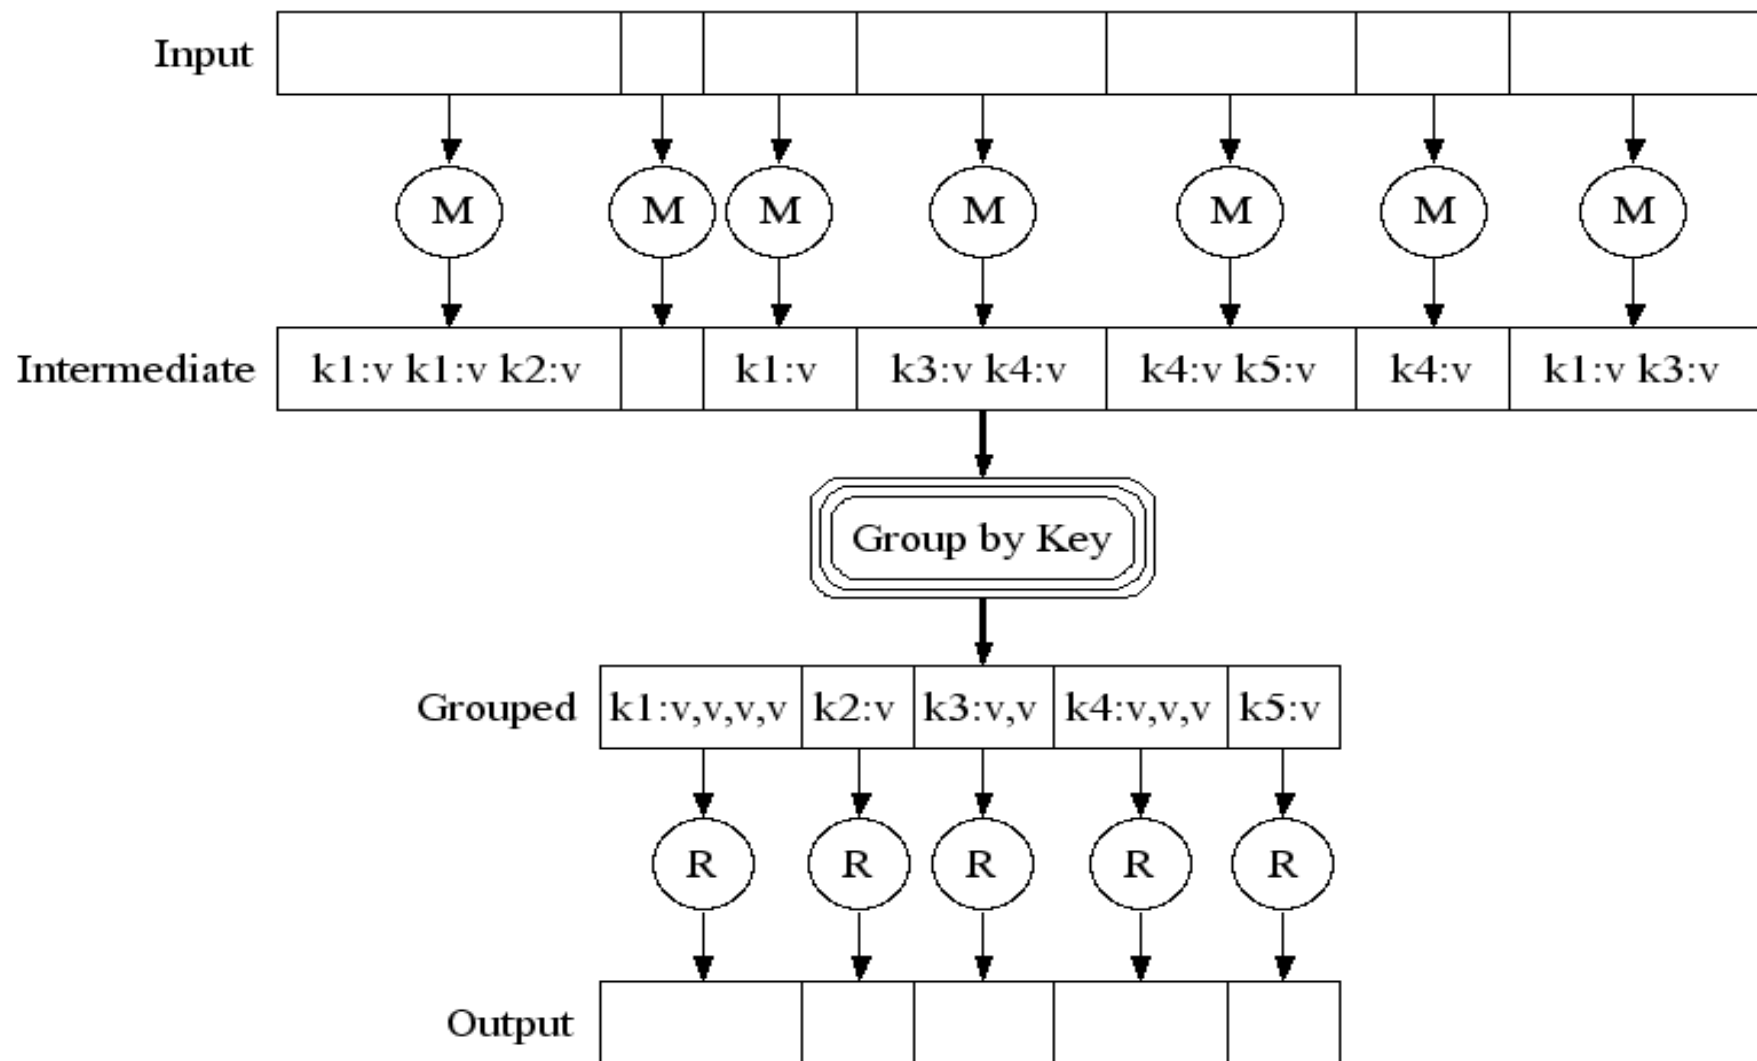
- "Consider the **problem of counting the number of occurrences of each word in a large collection of documents**"

**Divide collection
of document
among the class.**

**Sum up the
counts from all
the documents to
give final answer.**

**Each person gives count of
individual word in a document.
Repeats for assigned quota of
documents.**

(Done w/o communication)





Word Count Example

- **Mapper**

- Input: *key*: line offset, *value*: line of text
- Output: *key*: word, *value*: 1

- **Reducer**

- Input: *key*: word, *value*: set of counts
- Output: *key*: word, *value*: sum



Word Count Dataflow

The overall MapReduce word count process

Input

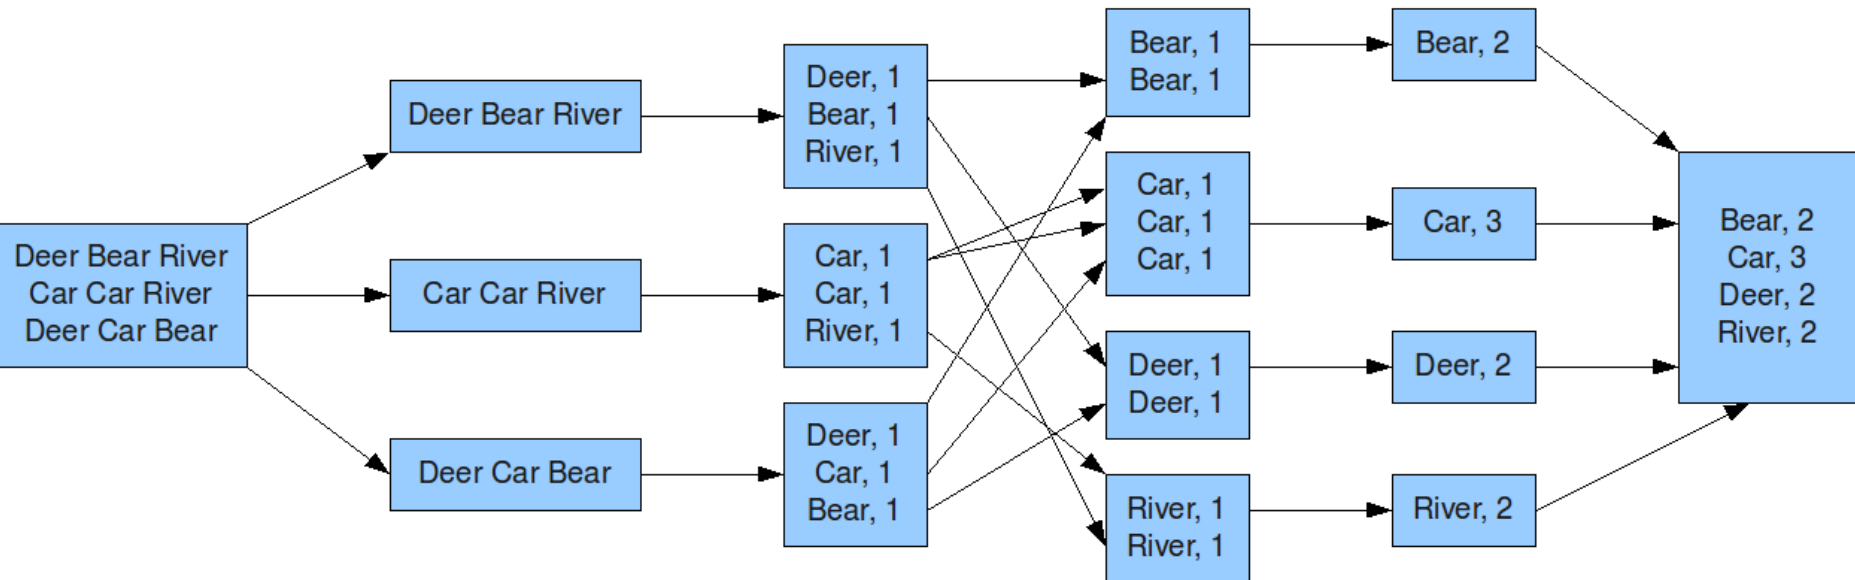
Splitting

Mapping

Shuffling

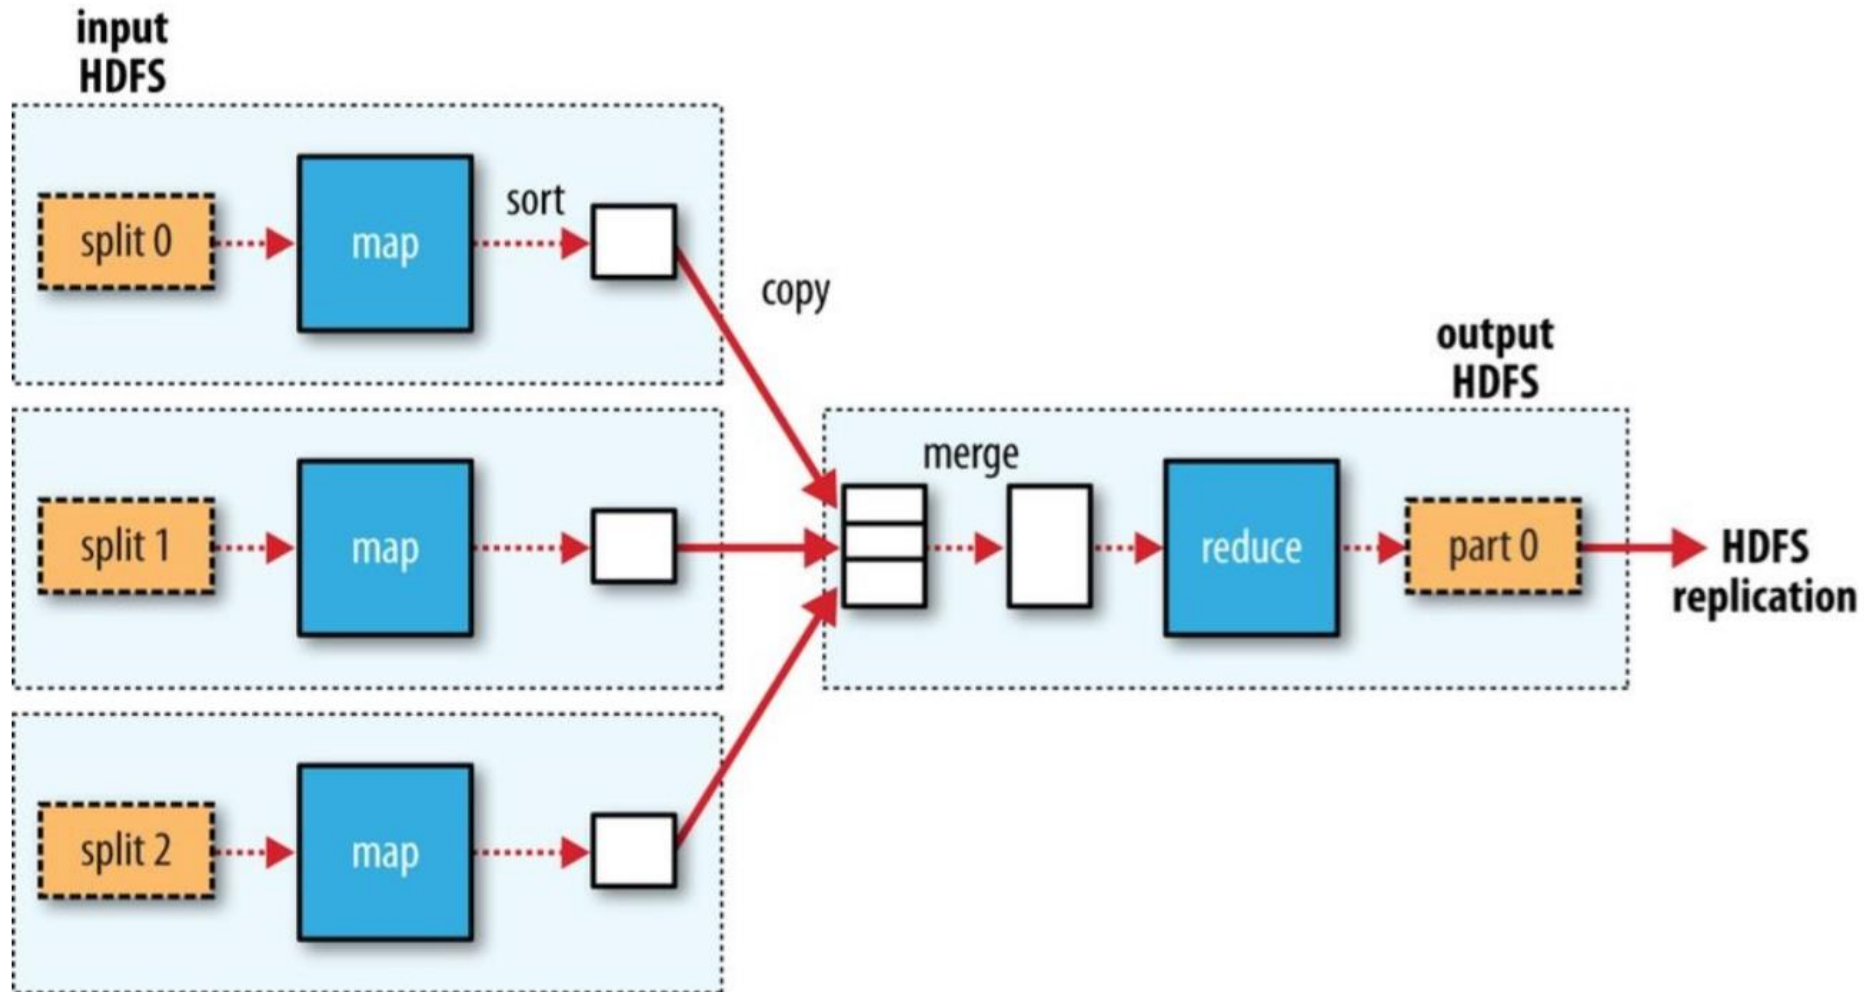
Reducing

Final result



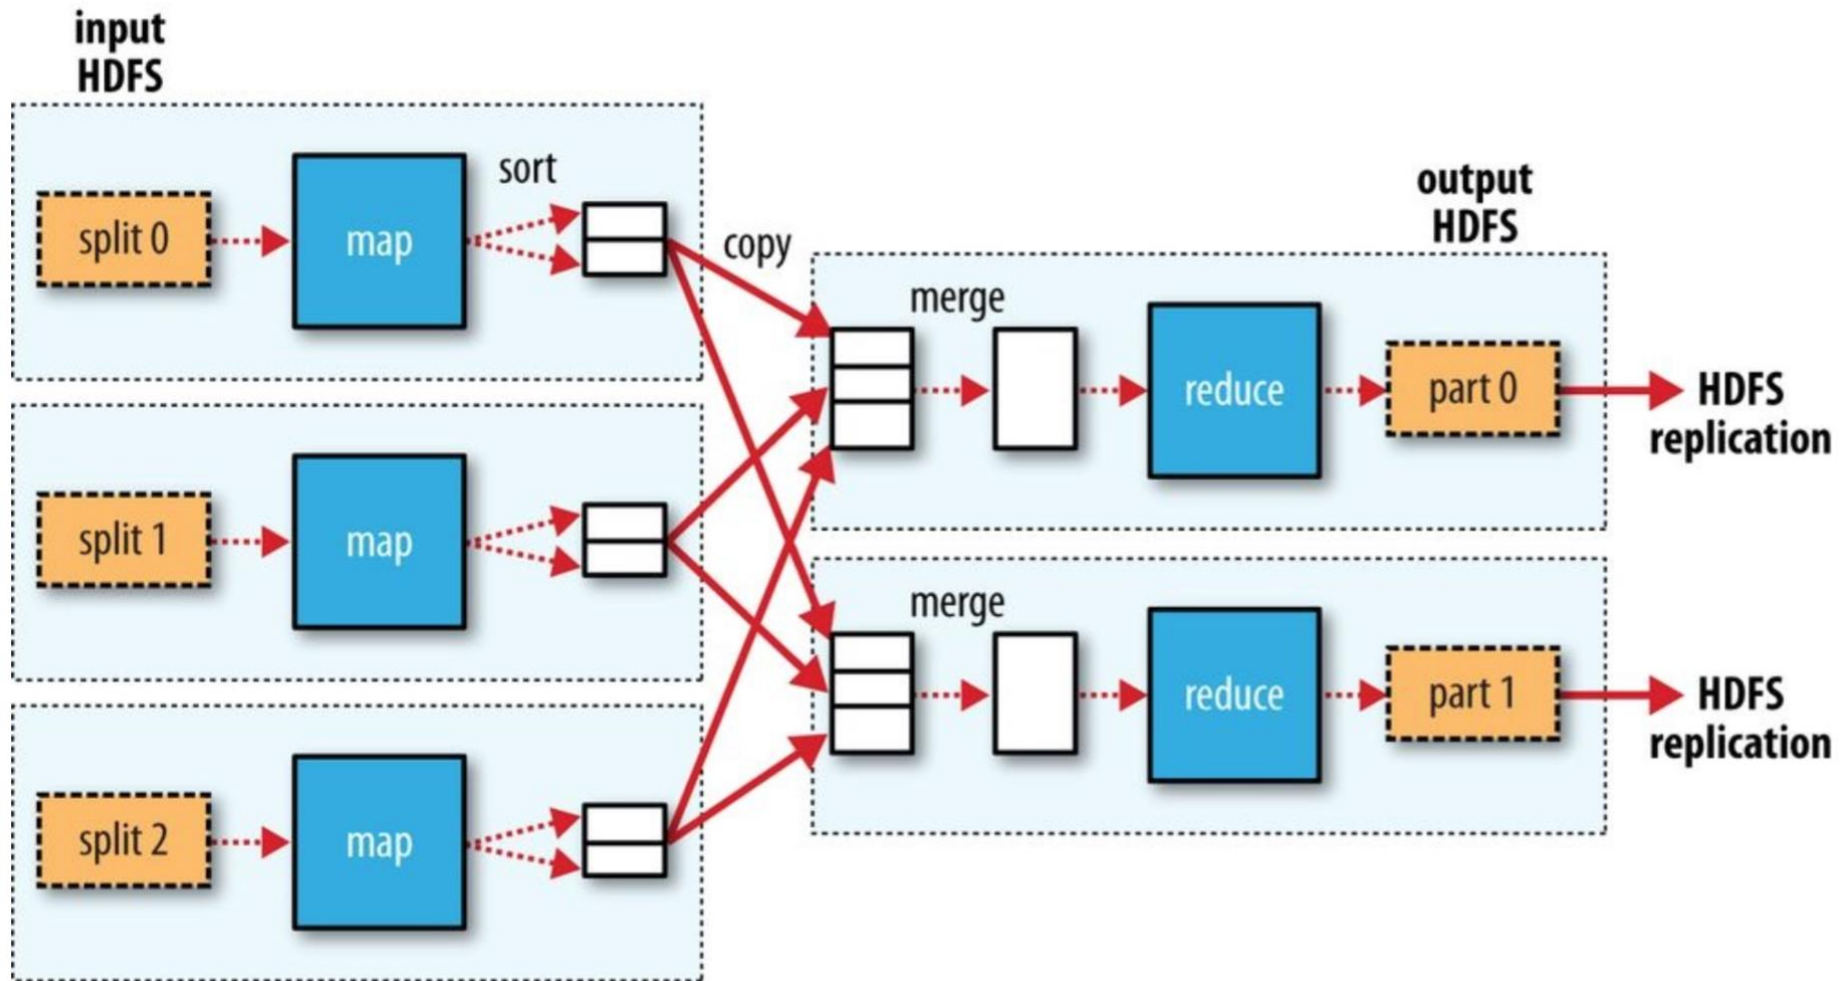


Data-flow with a single reduce task





Data-flow with multiple reduce tasks





A closer look

- **Input files:** Can have different format, usually reside on HDFS, tens of *GBs in size*
- **InputFormat:** How these input files are split up and read. The default InputFormat is the *TextInputFormat*.

InputFormat:	Description:	Key:	Value:
TextInputFormat	Default format; reads lines of text files	The byte offset of the line	The line contents
KeyValueInputFormat	Parses lines into key, val pairs	Everything up to the first tab character	The remainder of the line
SequenceFileInputFormat	A Hadoop-specific high-performance binary format	user-defined	user-defined



A closer look

- A **MapReduce job** is a **unit of work** that the **client** wants **to be performed**: it consists of:
 1. the **input data**,
 2. the **MapReduce program**,
 3. and **configuration information**.
- Hadoop **runs** the **job** by **dividing** it into **tasks**, of **which** **there are two types**:
 - **map tasks** and **reduce tasks**.
 - **scheduled** using **YARN** and **run on nodes** in the **cluster**



A closer look


- Hadoop **divides** the input to a MapReduce job **into fixed-size pieces** called **input splits**, or just **splits**.
- Hadoop **creates** one map task for **each split**, which executes user-defined **map function** for **each record in the split**.
- Splits size, parallelism, and overhead ?



Data locality optimization


- Hadoop does its best to **run the map task** on a **node** where the **input data resides** in HDFS
 - To do **data locality optimization**.
- Or → **Next best solution**: look for a **free map slot** on **another node** in the same rack
- Or → **In worst case**: an **off-rack node** is used, which results in an **inter-rack network transfer**.

Typical Hadoop Cluster





Data Locality Optimization





Reduce Tasks

- **Data Locality?**
 - Reduce tasks don't have the advantage of data locality, because:
 - Input to a **single reduce task** is normally the output from **all mappers**.
- The **reduce output** is **stored** in **HDFS** for **reliability**.



A real dataset

- **For our example:** we will **write** a **program** that **mines** **weather data**.
- Weather **sensors** collect data **every hour** at many locations across the **globe** and gather a **large volume** of log **data**
- A **good candidate** for **analysis** with **MapReduce**
 - large semi-structured data



Data Format

- The **data** we will use is from the **National Climatic Data Center**, or **NCDC**.
- The **data** is **stored** using a **line-oriented ASCII format**, in which each line is a record.



```
0057
332130 # USAF weather station identifier
99999 # WBAN weather station identifier
19500101 # observation date
0300 # observation time
4
+51317 # latitude (degrees x 1000)
+028783 # longitude (degrees x 1000)
FM-12
+0171 # elevation (meters)
99999
V020
320 # wind direction (degrees)
1 # quality code
N
0072
1
00450 # sky ceiling height (meters)
1 # quality code
C
N
010000 # visibility distance (meters)
1 # quality code
N
9
-0128 # air temperature (degrees Celsius x 10)
1 # quality code
-0139 # dew point temperature (degrees Celsius x 10)
1 # quality code
10268 # atmospheric pressure (hectopascals x 10)
1 # quality code
```



Data Format

- **Datafiles** are **organized** by **date** and **weather station**.
- There is a **directory** for **each year** from **1901 to 2001**, each containing a gzipped file for **each weather station** with its **readings** for that **year**.

```
% ls raw/1990 | head
010010-99999-1990.gz
010014-99999-1990.gz
010015-99999-1990.gz
010016-99999-1990.gz
010017-99999-1990.gz
010030-99999-1990.gz
010040-99999-1990.gz
010080-99999-1990.gz
```



Data Format

- There are **tens of thousands of weather stations**, so the **whole dataset** is made up of a **large number of relatively small files**.
- It's **generally easier and more efficient to process a smaller number of relatively large files**,
 - so the **data was preprocessed** so that each **year's readings were concatenated** into a **single file**.



Q. What's the highest recorded global temperature for each year in the dataset?



MapReduce

- **MapReduce** is a **programming model** for **data processing**.
- MapReduce works by **breaking** the **processing** into **two phases**:
 - **map** phase and the **reduce** phase.
 - **Each phase** has **key-value pairs** as **input** and **output**,
 - **Type** is chosen by the **programmer**.
 - The **programmer** also **specifies two functions**: the **map function** and the **reduce function**
 - **So? What should we do?**



How it can be done using Hadoop?

- Some sample records ...

```
00670119909999991950051507004...9999999N9+00001+9999999999...
00430119909999991950051512004...9999999N9+00221+9999999999...
00430119909999991950051518004...9999999N9-00111+9999999999...
00430126509999991949032412004...0500001N9+01111+9999999999...
00430126509999991949032418004...0500001N9+00781+9999999999...
```

- And or fields of interest

```
670119909999991950051507004...9999999N9+00001+9999999999... )
00430119909999991950051512004...9999999N9+00221+9999999999..
00430119909999991950051518004...9999999N9-00111+9999999999..
00430126509999991949032412004...0500001N9+01111+9999999999..
00430126509999991949032418004...0500001N9+00781+9999999999..
```



The *map* Phase

- The **input** to our **map phase**?
 - raw NCDC data
 - Shouldnt it be in Key/Value pair?
- The **key** is the **line offset** from the **beginning** of the **file**, but as **we have no need for this**, we ignore it.
- Each line in the **dataset** is a **text value**.



The *map* Phase

- What would the map phase do?
 - Produce some intermediate data for shuffle/reduce phase
 - So from the input line, we need to pull some values out, which?
 - We pull out the year and the air temperature, because these are the only fields we are interested in.



To **visualize** the way the **map works**, consider the following sample **lines** of **input data**

```
00670119909999991950051507004...9999999N9+00001+999999999999...
00430119909999991950051512004...9999999N9+00221+999999999999...
00430119909999991950051518004...9999999N9-00111+999999999999...
00430126509999991949032412004...0500001N9+01111+999999999999...
00430126509999991949032418004...0500001N9+00781+999999999999...
```

These **lines** are **presented** to the **map function** as the **key-value pairs**:

```
(0, 00670119909999991950051507004...9999999N9+00001+999999999999...)
(106, 00430119909999991950051512004...9999999N9+00221+999999999999...)
(212, 00430119909999991950051518004...9999999N9-00111+999999999999...)
(318, 00430126509999991949032412004...0500001N9+01111+999999999999...)
(424, 00430126509999991949032418004...0500001N9+00781+999999999999...)
```



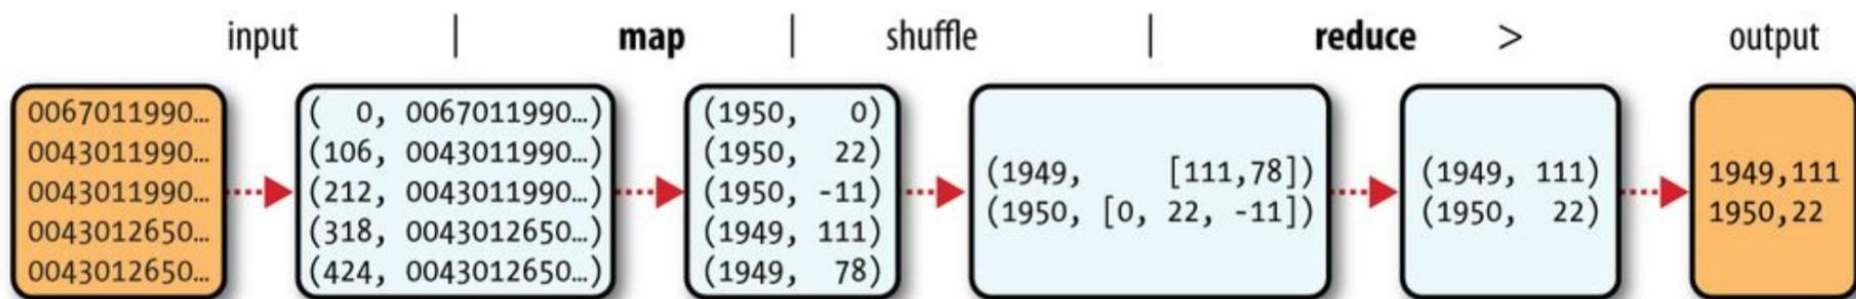
- The **map function** merely **extracts** the **year** and the **air temperature (integer)**, and **emits** them as its output:

```
(1950, 0)
(1950, 22)
(1950, -11)
(1949, 111)
(1949, 78)
```

- The **output** from the **map function** is **processed** by the **MapReduce framework** before **being sent to the reduce function**. This processing **sorts and groups** the **key-value pairs** by **key**.

```
(1949, [111, 78])
(1950, [0, 22, -11])
```

- All the **reduce** function has to do now is **iterate through** the **list** and **pick** up the **maximum reading**





Java MapReduce

- We need **three things**: a **map** function, a **reduce** function, and some **code** to run the job.
- The **map function** is represented by the **Mapper class**, which declares an abstract `map()` method.
- The **reduce function** is similarly defined using a **Reducer**


```
import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class MaxTemperatureMapper
    extends Mapper<LongWritable, Text, Text, IntWritable> {

    private static final int MISSING = 9999;

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String line = value.toString();
        String year = line.substring(15, 19);
        int airTemperature;
        if (line.charAt(87) == '+') { // parseInt doesn't like leading plus signs
            airTemperature = Integer.parseInt(line.substring(88, 92));
        } else {
            airTemperature = Integer.parseInt(line.substring(87, 92));
        }
        String quality = line.substring(92, 93);
        if (airTemperature != MISSING && quality.matches("[01459]")) {
            context.write(new Text(year), new IntWritable(airTemperature));
        }
    }
}
```




```
import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class MaxTemperatureReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int maxValue = Integer.MIN_VALUE;
        for (IntWritable value : values) {
            maxValue = Math.max(maxValue, value.get());
        }
        context.write(key, new IntWritable(maxValue));
    }
}
```

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MaxTemperature {

    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: MaxTemperature <input path> <output path>");
            System.exit(-1);
        }

        Job job = new Job();
        job.setJarByClass(MaxTemperature.class);
        job.setJobName("Max temperature");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setMapperClass(MaxTemperatureMapper.class);
        job.setReducerClass(MaxTemperatureReducer.class);

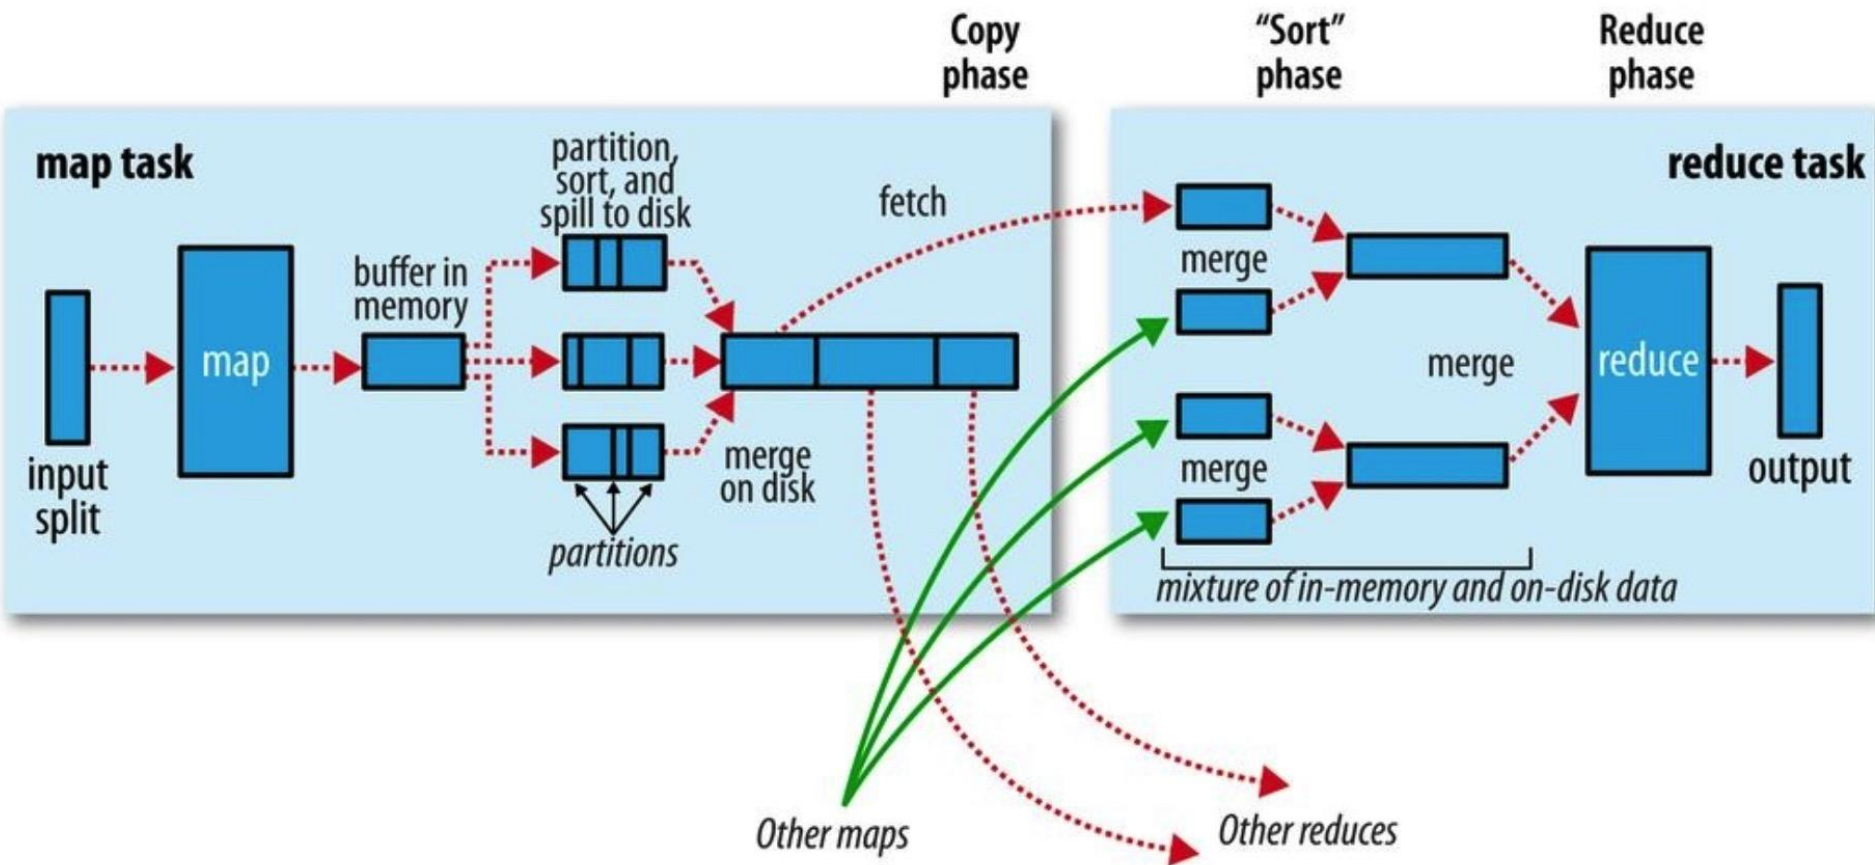
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```



Number of reduce tasks

- The **number of reduce tasks** is not governed by the size of the **input**, but instead is **specified independently**.
- When there are **multiple reducers**,
 - the **map tasks** partition their output,
 - each **creating one partition for each reduce task**.
 - the **records** for **any given key** are all in a single **partition**.





A refinement - Combiner

- Many **MapReduce** jobs are **limited** by the **bandwidth available** on the **cluster**
- It pays to **minimize** the **data transferred** between **map** and **reduce** tasks.
- Hadoop allows the **user to specify a combiner function** to be **run** on the **map output**, and the **combiner function's output** forms the **input** to the **reduce function**.



A refinement - Combiner Function

- Combiner function is an **optimization**
- Hadoop does not provide a guarantee it invocation
- Calling the combiner function *zero, one, or many times* should produce the same output from the reducer.



A refinement - Combiner Function

- Assume that for our example, **two maps** process data for year **1950** (because they were in different splits). Imagine the **first map produced** the output: **(1950, 0), (1950, 20), (1950, 10)**
- And the **second split produced**, **(1950, 25), (1950, 15)**
- The **reduce function** would be **called with a list of all the values**: **(1950, [0, 20, 10, 25, 15])**
- With output: **(1950, 25)**



A refinement - Combiner Function

- How to minimize the amount of data passed to Reduce?
- Use a **combiner function** that, just like the reduce function, finds the **maximum temperature** for each map output.

$$\max(0, 20, 10, 25, 15) = \max(\max(0, 20, 10), \max(25, 15)) = \max(20, 25) = 25$$



Some Interesting Examples



Facebook friends using MapReduce

- **Facebook** has a **list of friends**
 - **bi-directional** thing, If I'm your friend, you're mine.
 - **You** and **me** have **230 friends** in **common**.
 - When you **visit someone's profile**, you **see a list of friends** that **you** have **in common**

We're going to use mapreduce so that we can calculate everyone's common friends...



Facebook friends using MapReduce

- Assume the **friends** are **stored** as **Person->[List of Friends]**, our **friends list** is then:

$A \rightarrow B C D$
 $B \rightarrow A C D E$
 $C \rightarrow A B D E$
 $D \rightarrow A B C E$
 $E \rightarrow B C D$

- How it can be done?
- What would the Map phase do?
- and what would the reduce do?



Facebook friends - The map phase

- Each line will be an **argument** to a **mapper**.
- For **every friend** in the **list of friends**, the **mapper** will **output a key-value pair**.
 - The **key** will be a **friend** along **with the person**.
 - The **value** will be the **list of friends**.
 - The **key will be sorted** so that the **friends are in order**,
 - **Causing** all **pairs of friends** to go to the **same reducer**.



Facebook friends - The map phase

For map($A \rightarrow B C D$) :

$(A B) \rightarrow B C D$

$(A C) \rightarrow B C D$

$(A D) \rightarrow B C D$

For map($B \rightarrow A C D E$) :

$(A B) \rightarrow A C D E$

$(B C) \rightarrow A C D E$

$(B D) \rightarrow A C D E$

$(B E) \rightarrow A C D E$

*Note that **A** comes before **B** in the **key***



Facebook friends - The map phase

For map($C \rightarrow A B D E$) :

$(A C) \rightarrow A B D E$

$(B C) \rightarrow A B D E$

$(C D) \rightarrow A B D E$

$(C E) \rightarrow A B D E$

For map($D \rightarrow A B C E$) :

$(A D) \rightarrow A B C E$

$(B D) \rightarrow A B C E$

$(C D) \rightarrow A B C E$

$(D E) \rightarrow A B C E$

For map($E \rightarrow B C D$):

$(B E) \rightarrow B C D$

$(C E) \rightarrow B C D$

$(D E) \rightarrow B C D$



The Sort/Shuffle (Group) Phase

- Before we **send** these **key-value pairs** to the **reducers**, we **group them** by their **keys** and **get**:

$(A\ B) \rightarrow (A\ C\ D\ E)\ (B\ C\ D)$
 $(A\ C) \rightarrow (A\ B\ D\ E)\ (B\ C\ D)$
 $(A\ D) \rightarrow (A\ B\ C\ E)\ (B\ C\ D)$
 $(B\ C) \rightarrow (A\ B\ D\ E)\ (A\ C\ D\ E)$
 $(B\ D) \rightarrow (A\ B\ C\ E)\ (A\ C\ D\ E)$
 $(B\ E) \rightarrow (A\ C\ D\ E)\ (B\ C\ D)$
 $(C\ D) \rightarrow (A\ B\ C\ E)\ (A\ B\ D\ E)$
 $(C\ E) \rightarrow (A\ B\ D\ E)\ (B\ C\ D)$
 $(D\ E) \rightarrow (A\ B\ C\ E)\ (B\ C\ D)$



Facebook friends - The reduce phase

- Each **line** will be **passed** as an **argument** to a **reducer**.
- It will **intersect** the **lists** of **values** and **output the same key** with the **result** of the **intersection**.
- For example, **reduce((A B) -> (A C D E) (B C D))** will **output (A B) : (C D)** and **means** that **friends A** and **B** have **C and D** as **common friends**.



Facebook friends - Final outcome

(A B) \rightarrow (C D)

(A C) \rightarrow (B D)

(A D) \rightarrow (B C)

(B C) \rightarrow (A D E)

(B D) \rightarrow (A C E)

(B E) \rightarrow (C D)

(C D) \rightarrow (A B E)

(C E) \rightarrow (B D)

(D E) \rightarrow (B C)



Yahoo Search Suggestions

YAHOO!

fast nuces

fast nuces **university**

fast nuces **peshawar**

fast nuces **lahore**

fast nuces **logo**

fast nuces **entry test result**

fast nuces **isl**

fast nuces **neon**

fast nuces **islamabad**

cial Site

ation for the

ne package with

The students and staff of the university are expected to...

[Karachi](#)

Karachi Campus; Student Life; Student ... Digital Eyez is...

The National University of Computer ... this University,...

[Contact Us](#)

Contact Us. Islamabad Campus. A.K. Brohi Road, H-11/4 Tel...



Yahoo Search Suggestions

- **How does it work?**
- **Basic Idea - Related concepts appear close together**



Yahoo Search Suggestions

- **Input:** Web pages
 - 10 billions pages having 10K bytes each - 100TB of input data
- **Output:** List of words and for each word list of related words - List (Word, List(Related Words))
- Not overly complicated to do that!!



How would it Work?

- **Map** - **Tokenize words** and **output pairs**
 - <word, nextWord>
 - <word, previousWord>
- **Group by Words** to have <word, list(relatedWords)>
- **Reduce** - **Add WordCount** to output
<word, list(relatedWords), count>



Any Question?