



Hadoop YARN

Department of Computer Science,
National University of Computer & Emerging Sciences,
Islamabad Campus

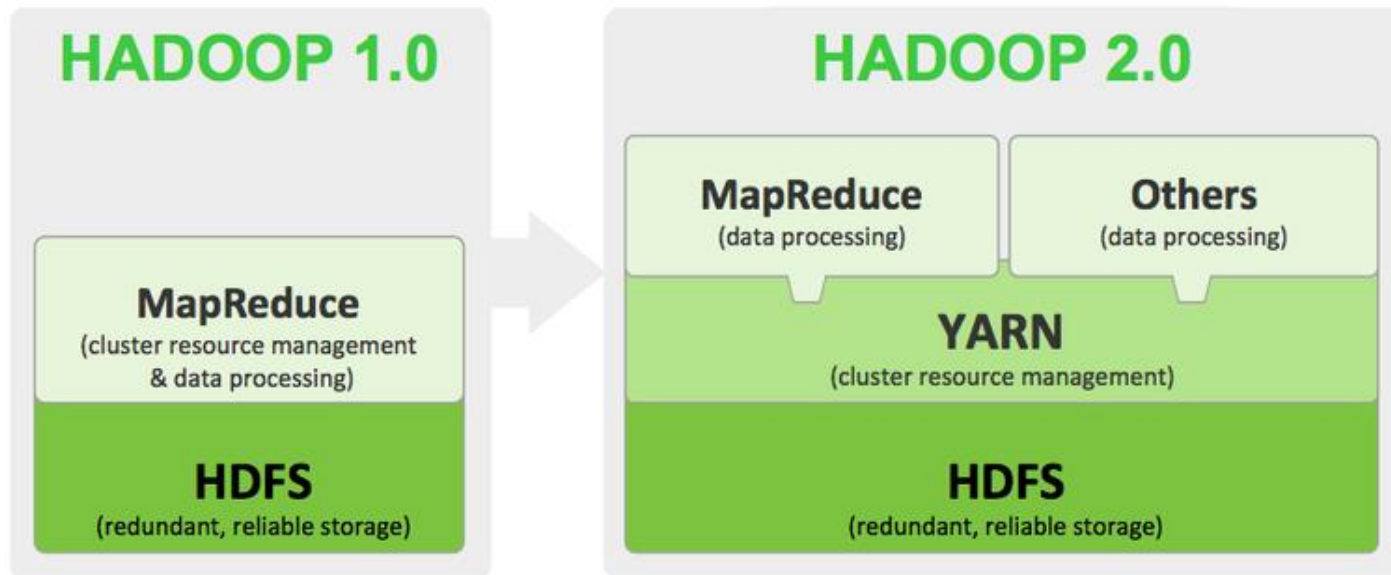


These slides are not ***always*** prepared by me and they do use content from different sources. Please do not share them 😊



Overview

- **YARN** was **introduced** in **Hadoop 2** initially, to improve the MapReduce implementation
- it is general enough to **support other distributed computing paradigms as well**






How does MapReduce1 work?

- In **MapReduce 1**, there are **two types** of **daemons** for **job execution process**:
 - The **jobtracker** **coordinates** all the **jobs** run on the system by **scheduling tasks** to **run** on **tasktrackers**.
 - **Tasktrackers** **run tasks** and **send progress reports** to the **jobtracker**
 - **If a task fails**, the **jobtracker** can **reschedule it** on a **different tasktracker**.



How does MapReduce 1 work?



Job tracker manages all the Jobs (in a cluster), 3 main responsibilities:

Schedules tasks of all jobs to nodes

Keep monitoring the tasks executions etc.

Records history of the executed jobs



The TaskTracker

- The **TaskTracker** has **simple responsibilities**:
 - **launch/teardown tasks** on **orders** from the **JobTracker**
 - **provide task-status information** to the **JobTracker** periodically



The JobTracker

- **Responsible for resource management**
 - **Manages** the **worker nodes** i.e., TaskTrackers
 - **Tracks resource consumption/availability**
- **Performs job life-cycle management**
 - **Schedules** individual **tasks** of the job
 - **Tracks progress**
 - **Provides fault-tolerance** for **tasks** etc.

Limitations?

- The **JobTracker** was **over-burdened**.
 - Resource management
 - Job and task scheduling and
 - Monitoring
 - ...
- It was **posing a limitation** in terms of
 - Scalability, Availability, Resource Utilization ...





Limitations - Scalability

- **MapReduce 1** hits **scalability bottlenecks** in the region of **4,000 nodes** and **40,000 tasks**.
- **YARN** is **designed** to **scale up to 10,000 nodes** and **100,000 tasks**



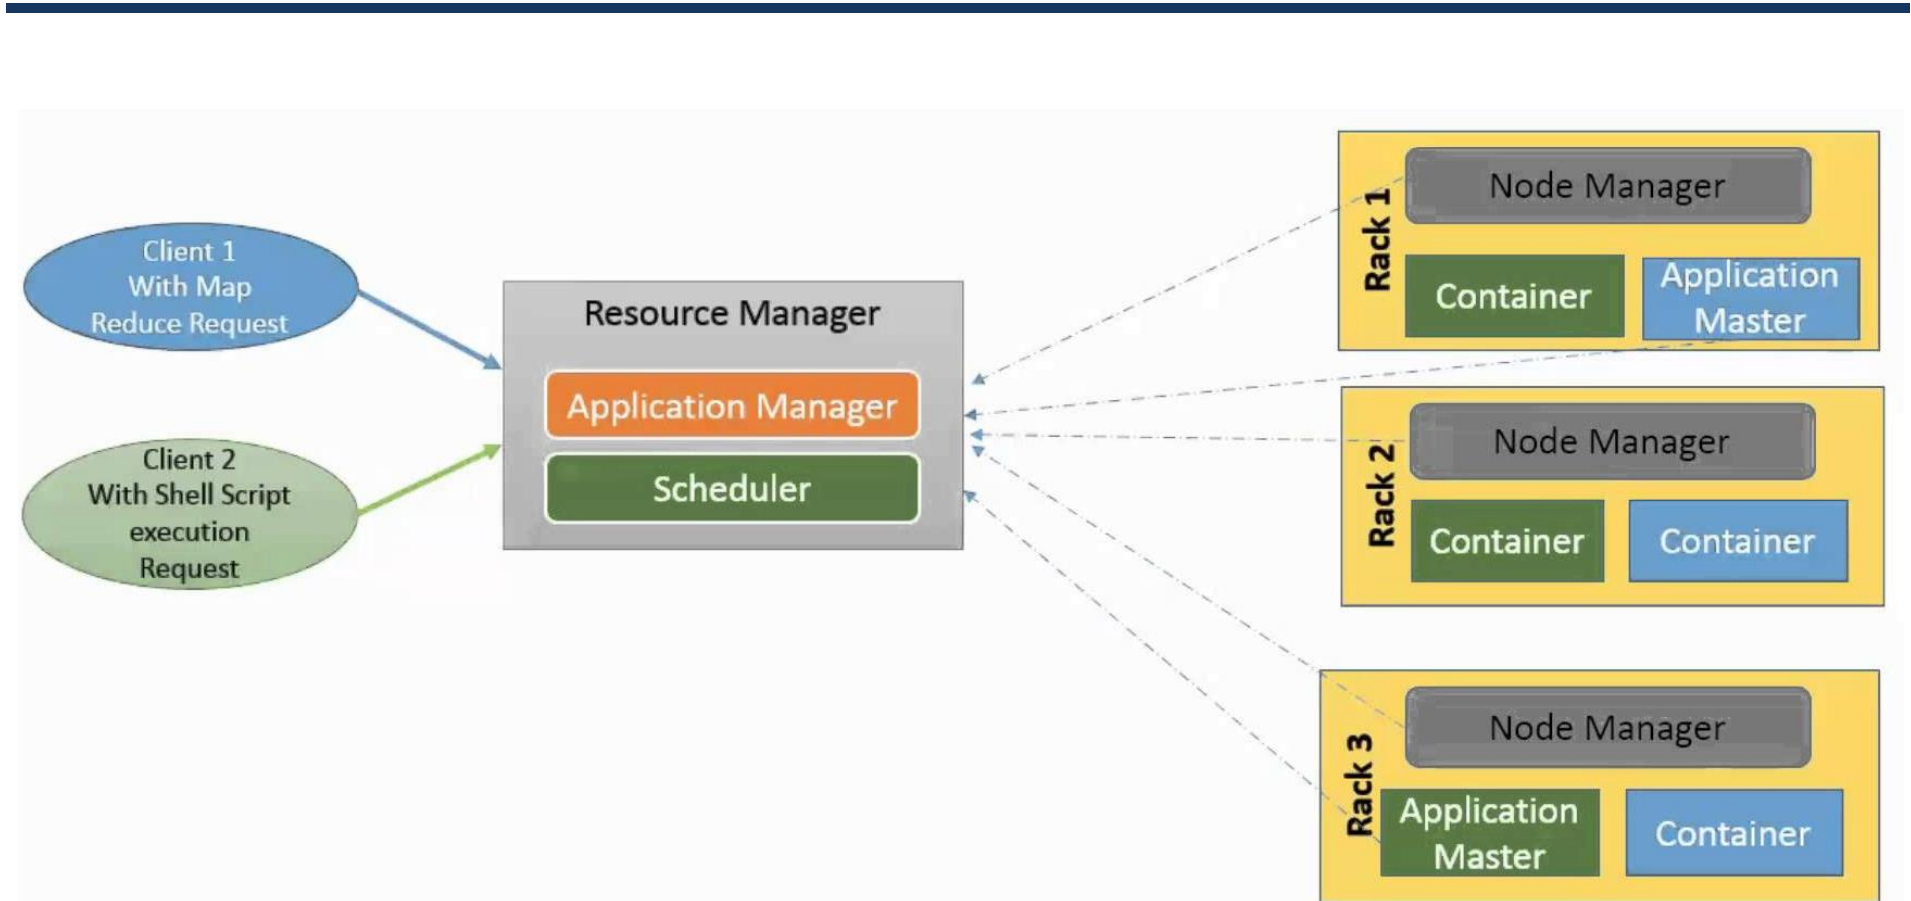
Limitations - Non MapReduce Tasks

- **Job tracker** was **tightly integrated** with **MapReduce**
- **MapReduce** works well but **provides batch processing** and **lacks real-time analysis**
- It should thus be **possible to support other workloads** in **Hadoop**



Hadoop YARN

- **Fundamental idea** - **divide and conquer**
- **Split up** the **two major responsibilities** of the **JobTracker** i.e., **resource management** and **job scheduling/monitoring**, into **separate daemons**:
 - a **global ResourceManager (one per Cluster)**
 - a **number of NodeManagers (One per node)**
 - and **per-application ApplicationMaster (AM)**.





ResourceManager

- The **ResourceManager** is the **ultimate authority** that **arbitrates (i.e., allocates) resources** among all the **applications** in the **system**.



NodeManager

- The **NodeManager** is a **per node worker** service that is **responsible** for the **execution of containers** based on the **node capacity**
- The **NodeManager** service **sends a heartbeat signal** to the **ResourceManager** to **update its health status**



YARN's ApplicationMaster

- Per application **framework-specific library** that **manages** each **instance** of an **application** that **runs within YARN**.
- Each **application** running on the **Hadoop cluster** has its own, **dedicated Application Master instance**
 - MapReduce **ApplicationMaster** to **run mapreduce jobs**
 - *Tez ApplicationsMaster, SPARK ApplicationMaster ...*
- The **Application Master** oversees the **full lifecycle** of an **application**



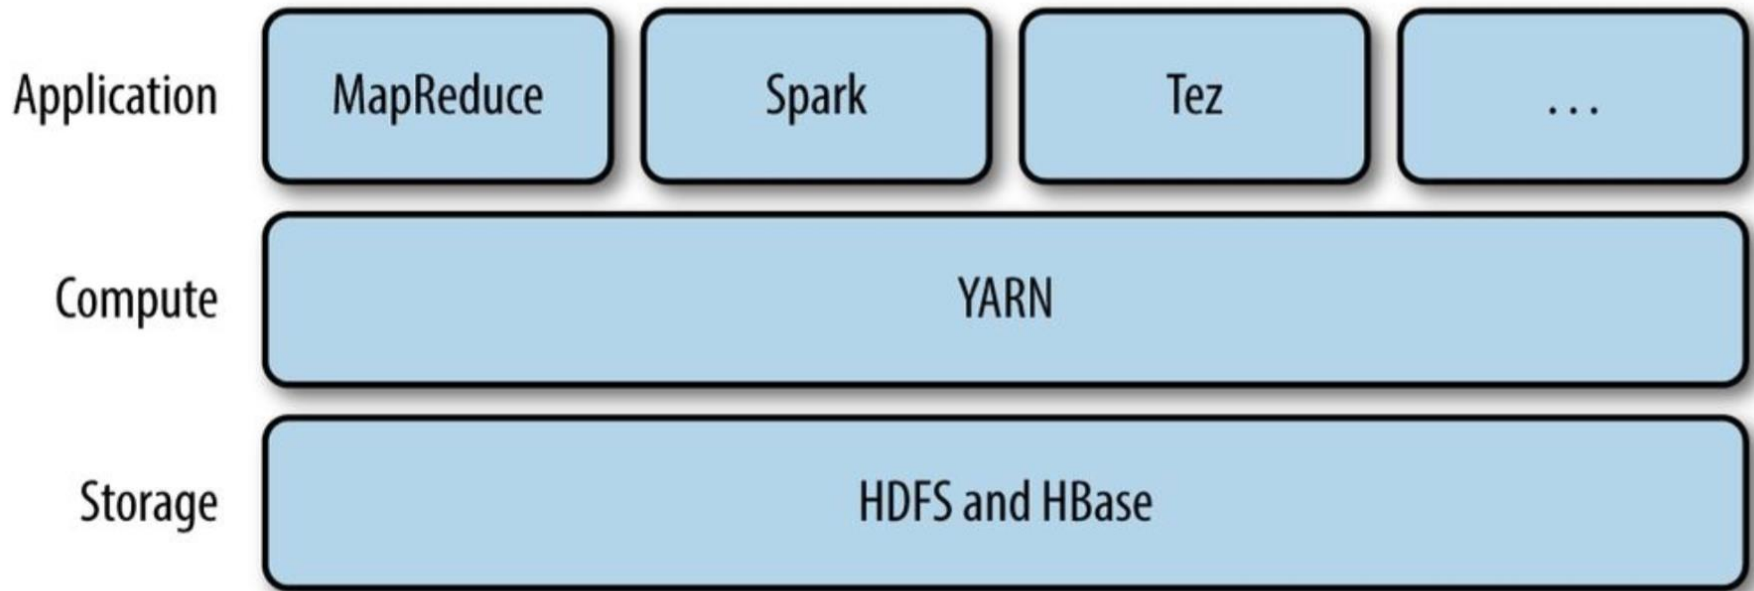
MapReduce 1 Vs. YARN

Table 4-1. A comparison of MapReduce 1 and YARN components

| MapReduce 1 | YARN |
|-------------|---|
| Jobtracker | Resource manager, application master, timeline server |
| Tasktracker | Node manager |
| Slot | Container |

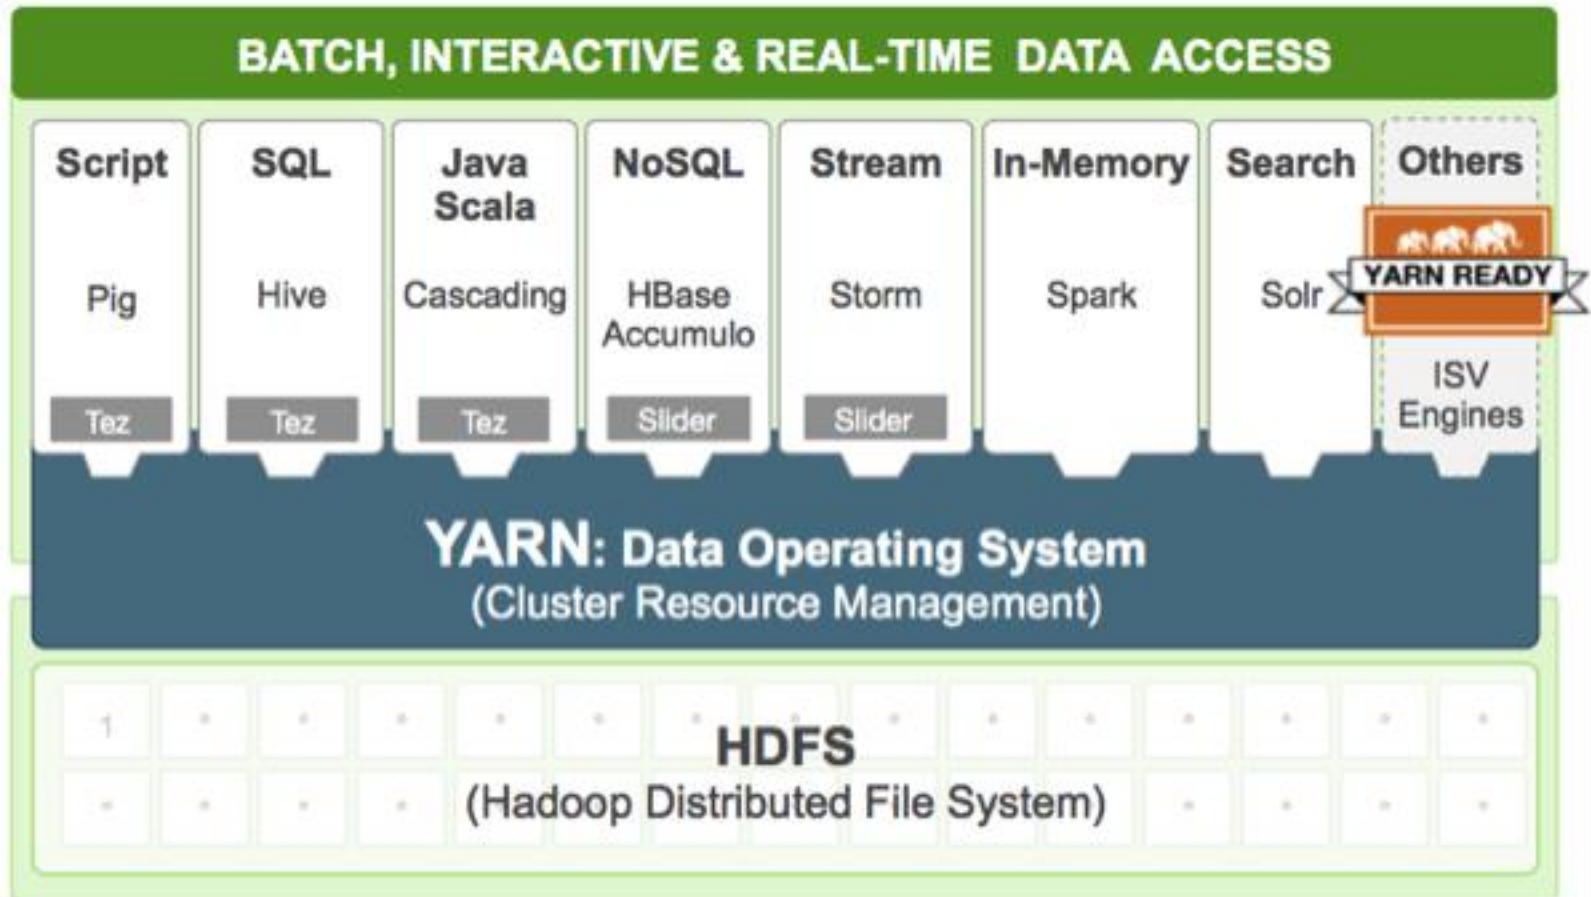


YARN Applications





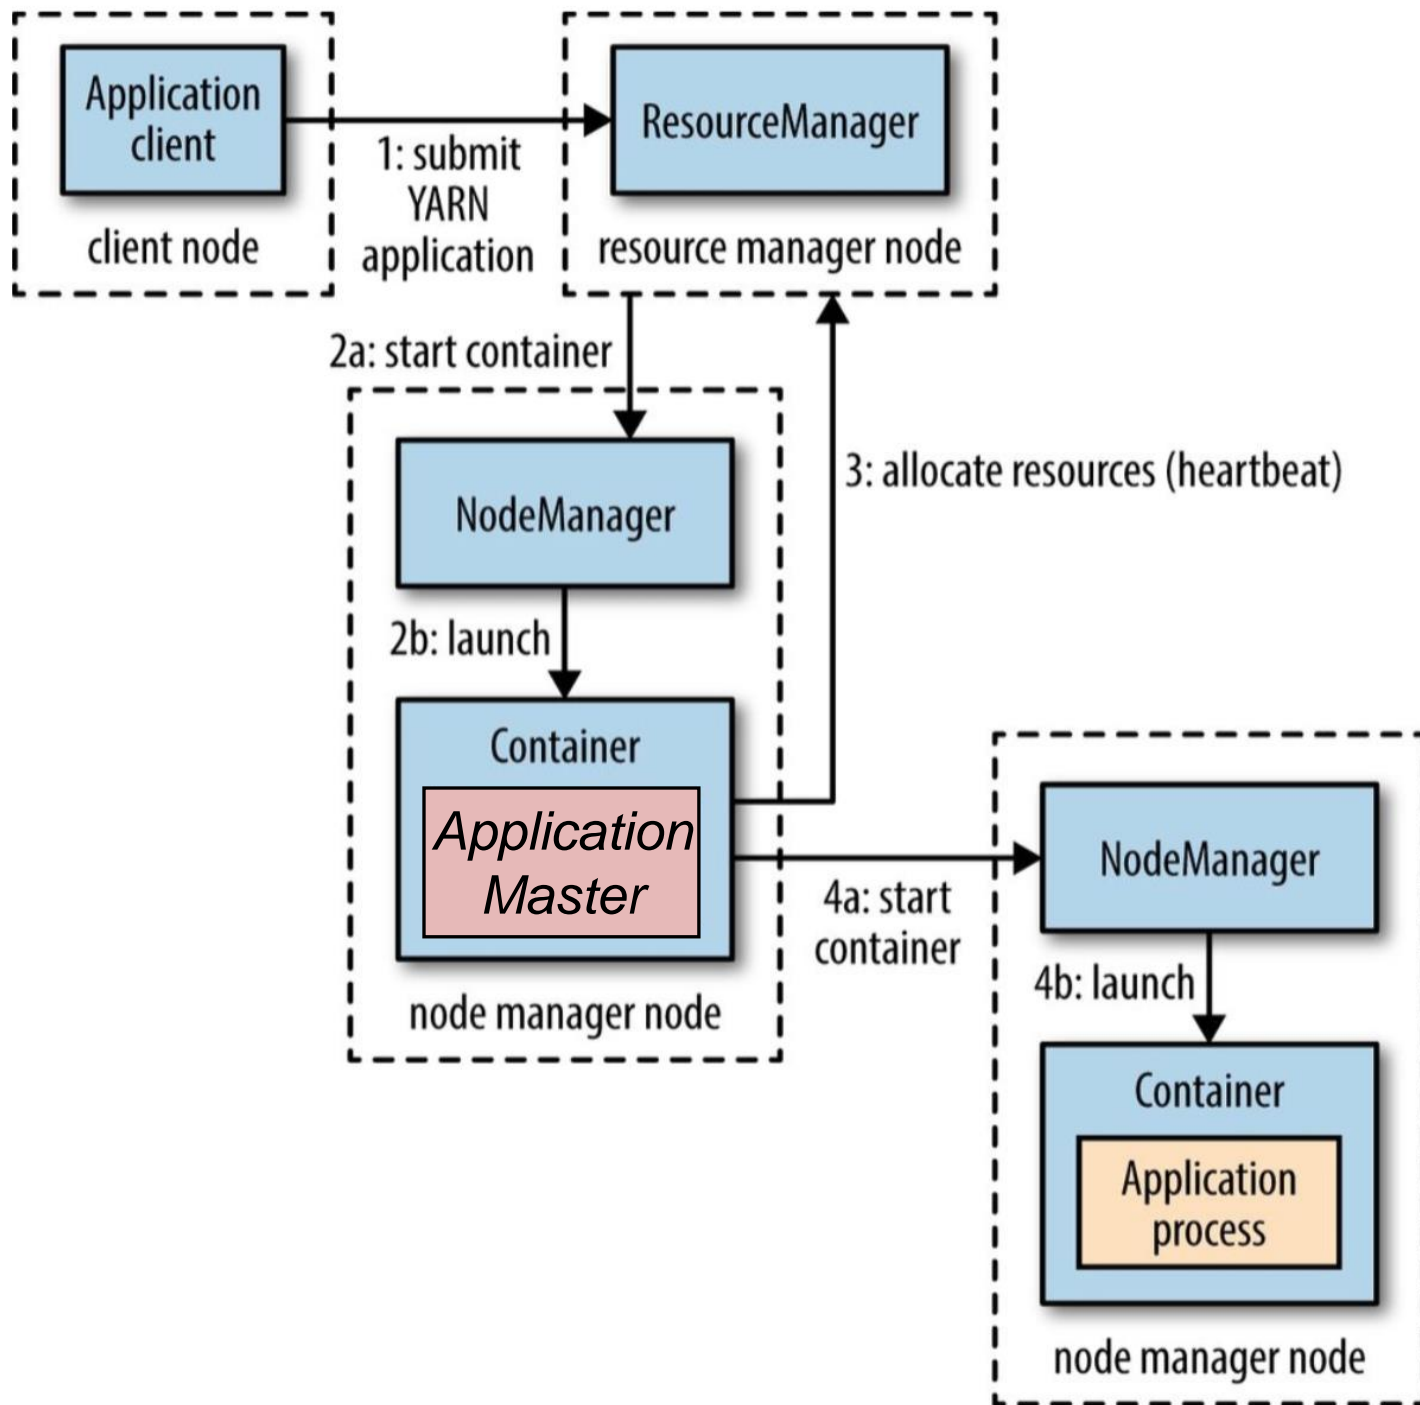
YARN Applications





Anatomy of a YARN Application Run

- To **run** an **application on YARN**, a **client contacts** the **resource manager** and asks it to **run an application master process**.
- The **resource manager** then **finds** a **node manager** that can **launch** the **application master** in a **container** (steps 2a and 2b).





Anatomy of a YARN Application Run

- **ApplicationMaster** is **framework specific** and can simply **run a computation** in the **container it is running** in
- It could **request more containers** from the **resource managers** (step 3), and **use them to run a distributed computation** (steps 4a and 4b).



Resource Requests

- A request for a **set of containers** can **express** the **amount** of **computer resources required** (*memory* and *CPU*), as well as **locality constraints**.
- **Locality constraints** can be used to **request a container** on a **specific node** or **rack**, or **anywhere on the cluster** (*off-rack*).




Scheduling in YARN

- It is the **job** of the **YARN scheduler** to **allocate resources to applications** according to some **defined policy**.
- **Three schedulers** are **available** in **YARN**: the ***FIFO***, ***Capacity***, and ***Fair Schedulers***.



YARN FIFO Scheduler

i. FIFO Scheduler




- One job execution at a time (first come first server)
- Small jobs have to wait for larger jobs to finish
- Single job utilizes full cluster resources



YARN Capacity Scheduler

ii. Capacity Scheduler




- **Capacities** for **different queues are reserved** (e.g., Queue for small, medium, or large jobs etc.)
- **Simultaneous execution of the jobs**



YARN Fair Scheduler

iii. Fair Scheduler



- Simultaneous execution of the jobs
- Dynamic resource allocation to jobs (no reserved capacity)



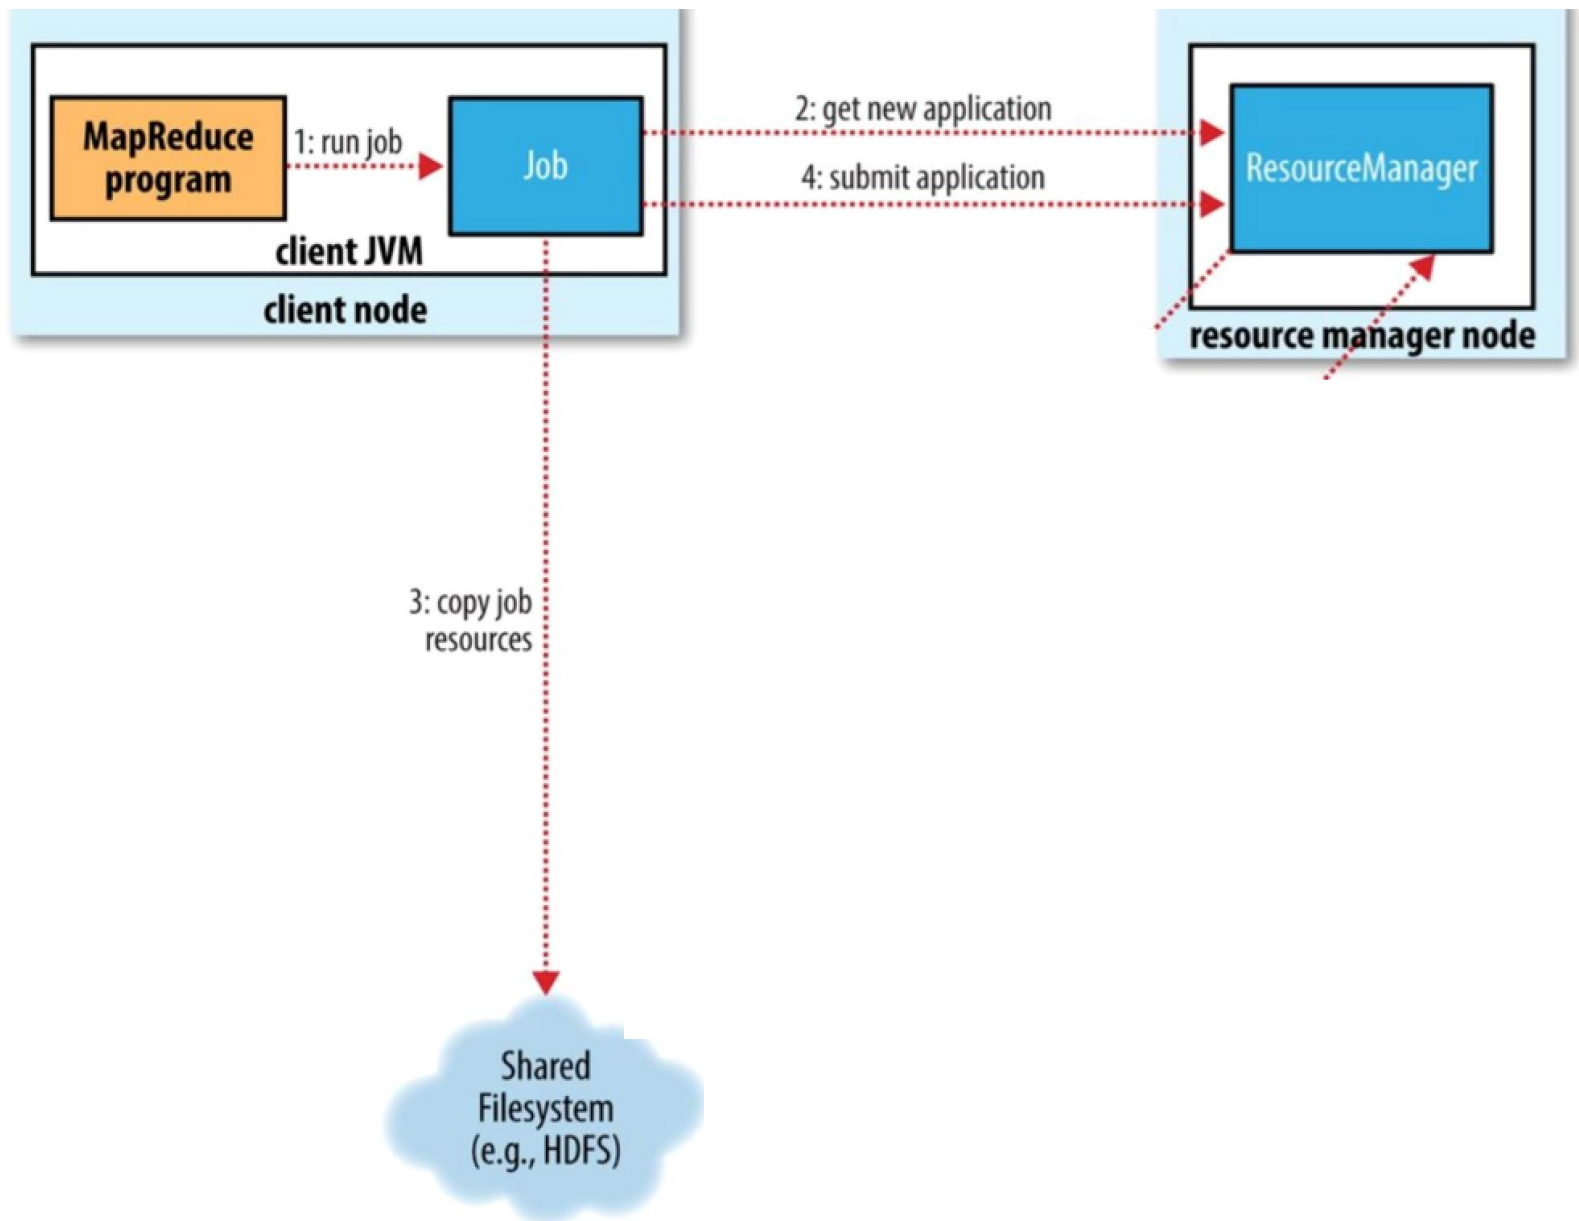
Anatomy of a MapReduce Job Run

- You can **run** a **MapReduce job** with **waitForCompletion()**, which **submits the job** and **then waits for it to finish**.



Anatomy of a MapReduce Job Run

- At the **highest level**, there are **five independent entities**:
 - The **client**, which **submits** the **MapReduce job**.
 - The **YARN resource manager**, which **coordinates** the **allocation** of compute **resources** on the cluster.
 - The **YARN node managers**, which **launch** and **monitor** the **compute containers** on **machines** in the cluster.
 - The **MapReduce application master**, which **coordinates** the **tasks** running the **MapReduce job**.
 - The **distributed filesystem**, normally **HDFS**






Job Submission

- The **job submission process** does the following:
 - Asks the **resource manager** for a **new application ID**, used for the **MapReduce job ID** (step 2).
 - Checks the **output specification** of the **job**.
 - Check input and Computes the **input splits** for the **job**.
 - Copies the **resources needed** to run the **job**, to the **shared filesystem** in a directory **named after the job ID** (step 3)
 - Submits the **job** by calling **submitApplication()** on the **resource manager** (step 4).
 - **waitForCompletion()** polls the **job's progress**



Job Initialization

- When the **resource manager** receives a call to its **submitApplication()** method, it **hands off** the **request** to the **YARN scheduler**.
- The **YARN scheduler** **allocates** a **container**, and the **resource manager** then **launches the application master's** process there, under the **node manager's** management (steps 5a and 5b).





Job Initialization

- The **application master** for **MapReduce** jobs **initializes** the **job** by **creating** a **number of bookkeeping objects** to **keep track** of the **job's progress** (step 6)
- Next, it **retrieves** the **input splits** **computed** in the **client** from the **shared filesystem** (step 7).



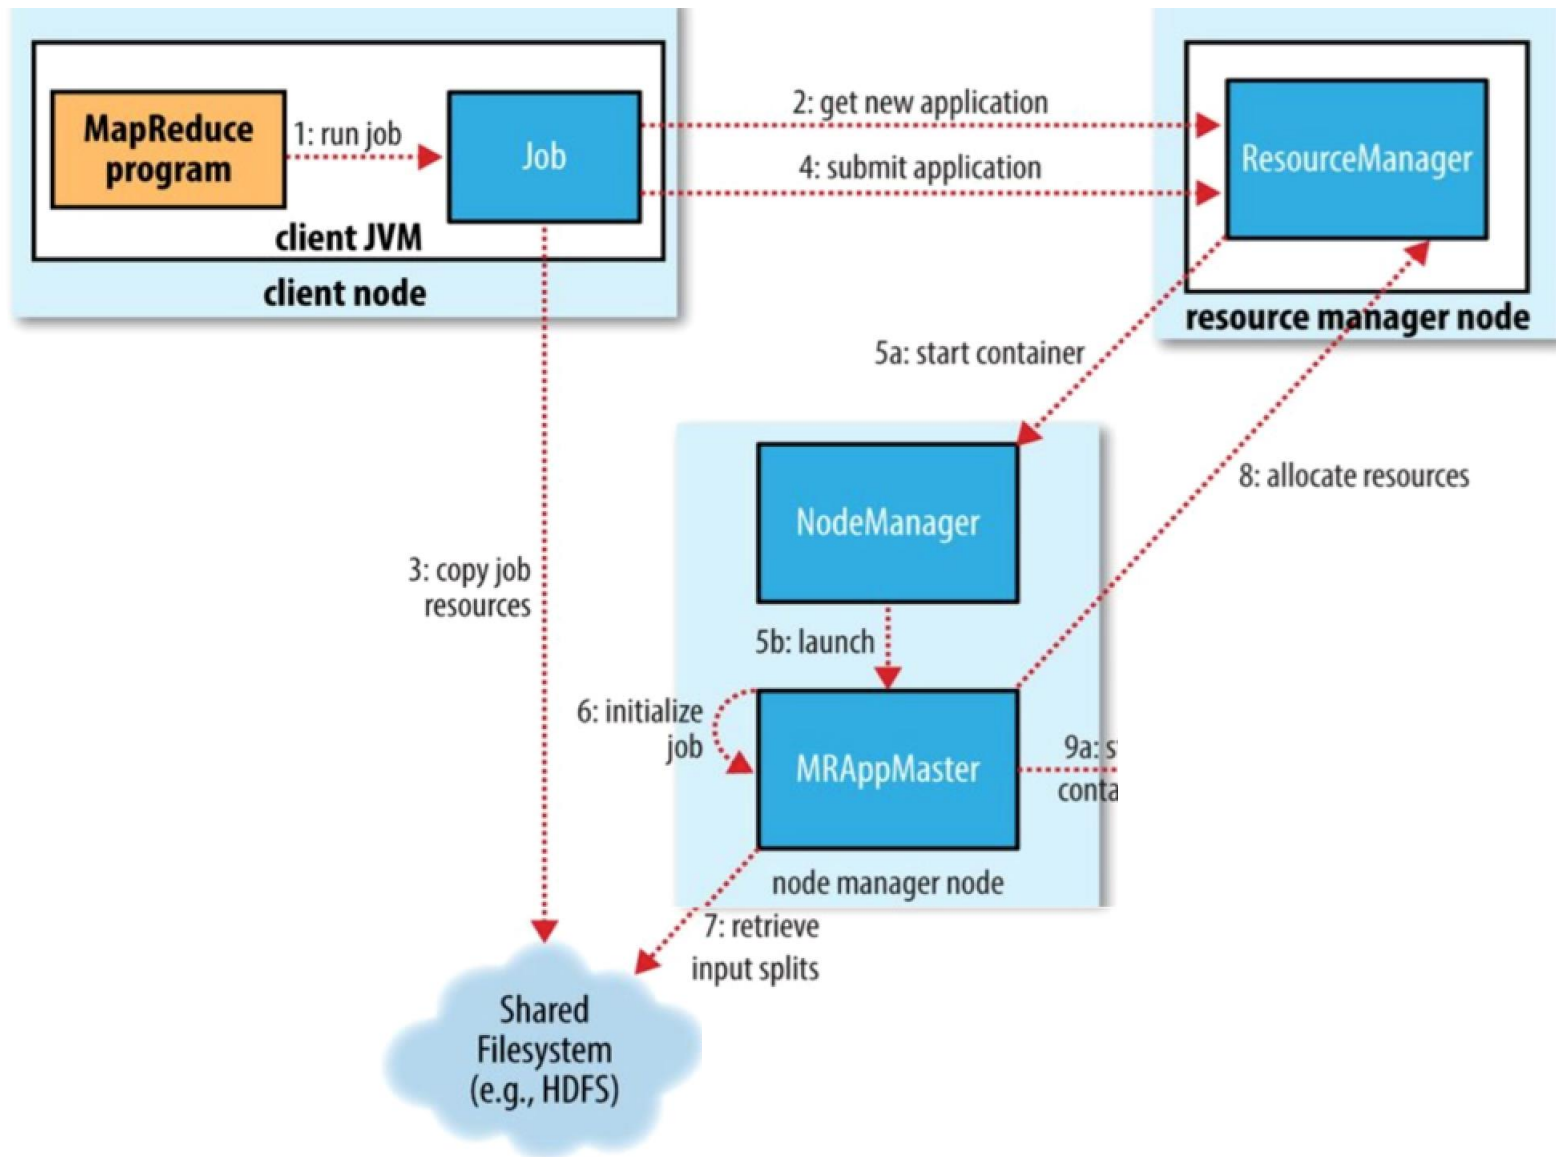
Job Initialization

- It then **creates** a **map task object** for **each split**, as well as a **number** of **reduce task objects**.
- **Tasks** are **given IDs** at this point.



Job Initialization

- The **application master** must **decide** *how to run the tasks* that **make up** the **MapReduce job**.
- **If** the **job** is **small**, the **application master** may **choose** to **run** the **tasks** in the **same JVM** as **itself**.
- **Such** a **job** is **said** to be **uberized**, or run as an uber task.
- The **application master** will also **creates** the **final output directory** the **temporary space** for the **task output**.





Task Assignment

- If the **job** does not run as an **uber task**, then the **application master** requests containers for *all* the map and reduce tasks in the **job** from the **resource manager** (step 8).
- Requests for **map tasks** are **made first** and with a higher priority than those for **reduce tasks**.
- **Reduce tasks** can run anywhere in the cluster, but requests for **map tasks** have **data locality constraints** that the scheduler *tries to honor*



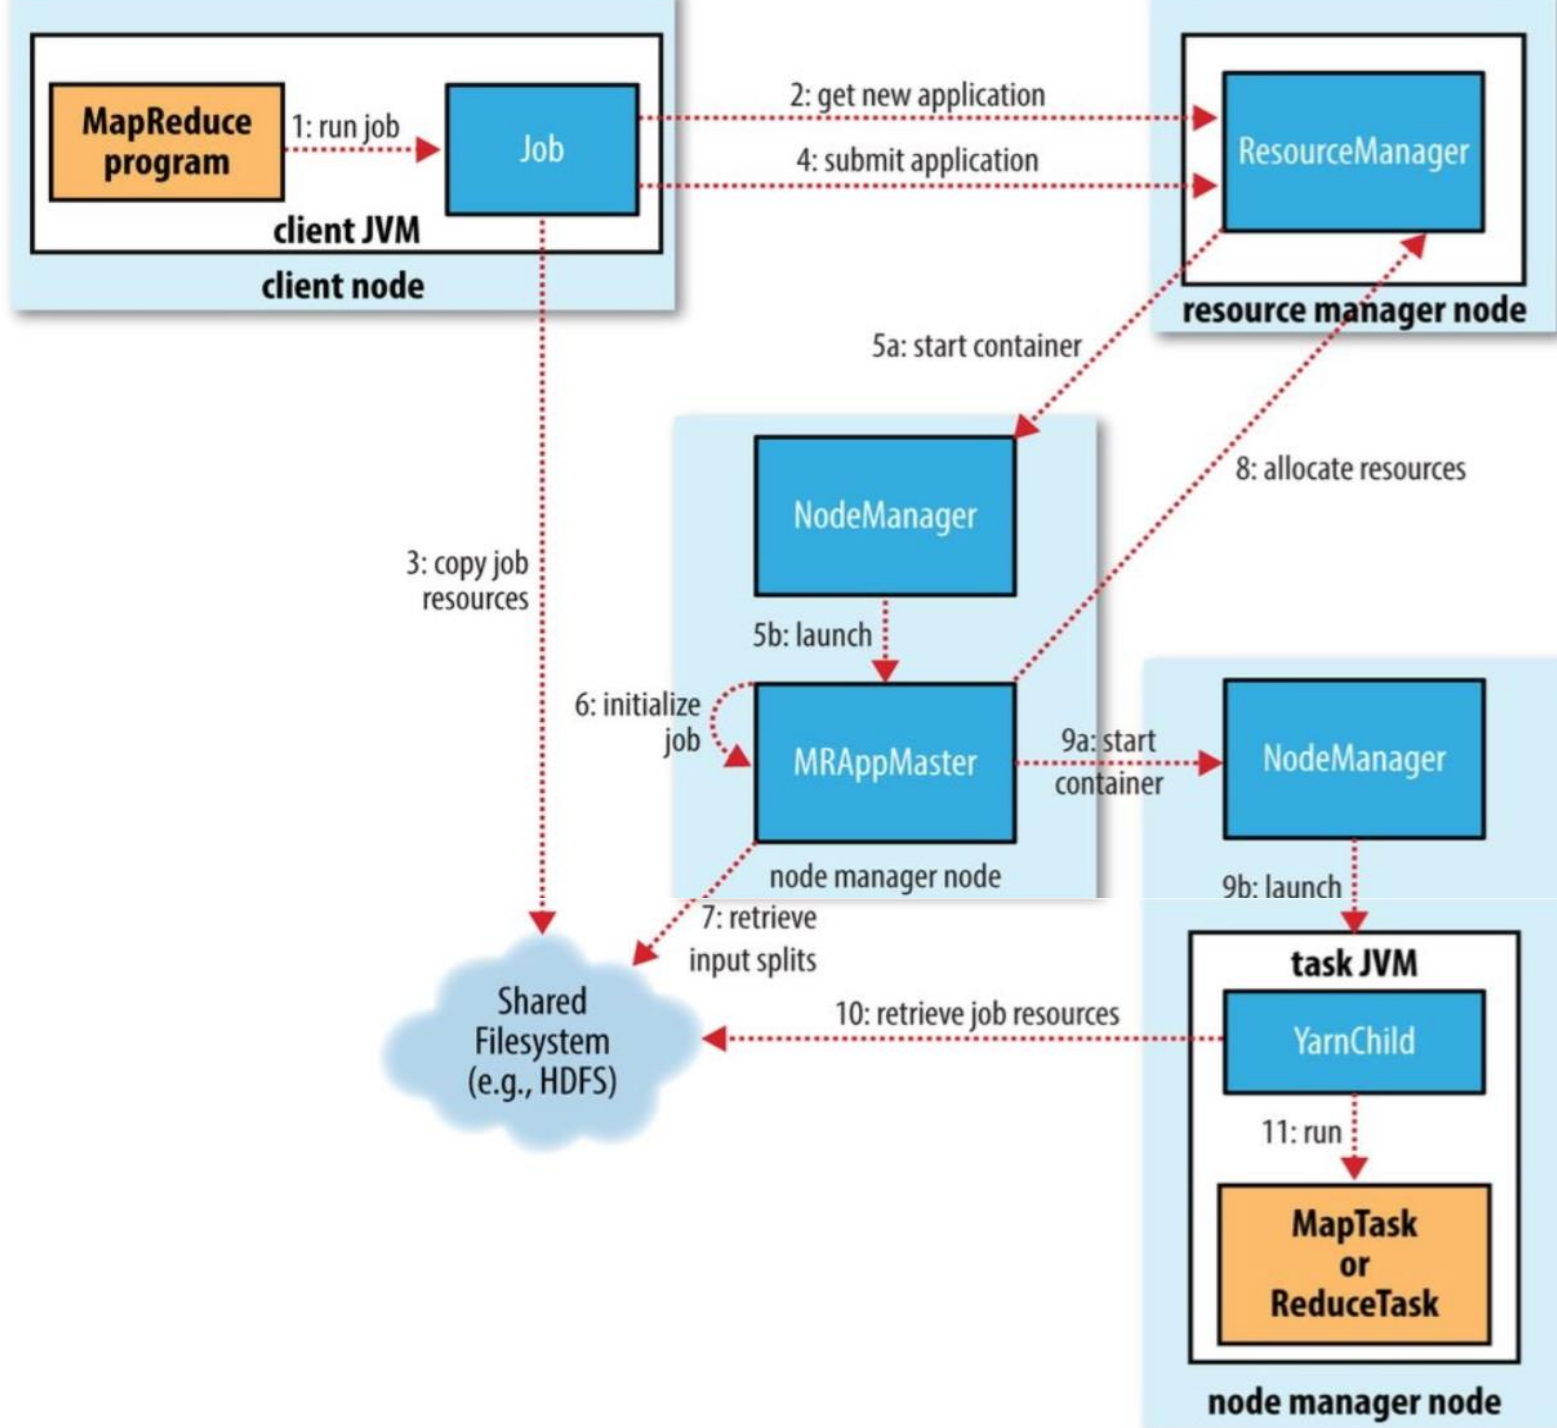
Task Execution

- On **resources assignment**, the **application master** starts the **container** by **contacting** the **node manager** (**steps 9a and 9b**).
- The **task** is **executed** by a **Java application** whose **main class** is **YarnChild**.



Task Execution

- The **YarnChild** runs in a dedicated JVM, so that **any bugs** in the **user-defined map** and **reduce functions** (or even in YarnChild) don't affect the **node manager** — by *causing it to crash or hang*.





Hadoop Installation

- **Standalone (or local) mode**, there are no daemons running and everything runs in a single process.
- **Pseudo-distributed mode**, Hadoop daemons run on the local machine, thus simulating a cluster on a small scale.
- **Fully distributed mode**, the Hadoop daemons run on a cluster of machines.



Hadoop Installation

- **Standalone (or local) mode**, there are no daemons running and everything runs in a single process.
- **Pseudo-distributed mode**, Hadoop daemons run on the local machine, thus simulating a cluster on a small scale.
- **Fully distributed mode**, the Hadoop daemons run on a cluster of machines.



YARN on a Single Node

- *Assuming you have already configured HDFS, follow the steps ...*



YARN on a Single Node

- **Step 1: Configure parameters as follows:**

etc/hadoop/mapred-site.xml:

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
  <property>
    <name>mapreduce.application.classpath</name>
    <value>$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/*:$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/lib/*</value>
  </property>
</configuration>
```

https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html#YARN_on_a_Single_Node



YARN on a Single Node

- **Step 1: Configure parameters as follows:**

etc/hadoop/yarn-site.xml:

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.env-whitelist</name>
    <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_
DIR,CLASSPATH_PREPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HO
ME</value>
  </property>
</configuration>
```



YARN on a Single Node

- **Step 2: Start ResourceManager daemon and NodeManager daemon:**

\$ start-yarn.sh

***Browse the web interface** for the **ResourceManager**; by default it is available at: <http://localhost:8088/>*



Any Questions ?