

National University of Computer & Emerging Sciences


Lecture 08 MPI - Point-to-Point Communication

MPI Fundamentals

- A **communicator** defines a group of processes that have the ability to communicate with one another.
- In this group of processes, each is assigned a unique **rank**, and they explicitly communicate with one another by their ranks.

Communicator `MPI_COMM_WORLD`

- All processes of an MPI program are members of the default **communicator `MPI_COMM_WORLD`**.
- `MPI_COMM_WORLD` is a predefined
- Each process has its own **rank** in a communicator:
 - starting with 0
 - ending with (size-1)



Lecture 08: Point-to-Point Communication

```
#include <mpi.h>
#include <stdio.h>
```

MPI header file

MPI_Comm_size returns the size of a communicator. Here, MPI_COMM_WORLD encloses all of the processes, so this call should return the amount of processes that were requested for the job.

```
int world_size;
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

// Get the rank of the process
int world_rank;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
```

MPI environment initialization, all of MPI's global and internal variables are constructed. For example, a communicator is formed around all of the processes that were spawned, and unique ranks are assigned to each process.

MPI_Comm_rank returns the rank of a process in a communicator. Ranks are incremental starting from zero and are primarily used for identification purposes during send/receive.

MPI_Get_processor_name obtains the actual name of the processor on which the process is executing.

```
" out of %d processors\n",
processor_name, world_rank, world_size);
```

```
// Finalize the MPI environment.
MPI_Finalize();
```

clean up the MPI environment

```
}
```

Running MPICH Program

- Regular applications
`./mpi_hello_world`
- MPI applications (running with 16 processes)
`mpiexec -n 16 ./mpi_hello_world`

Messages


- A message contains a number of elements of some particular datatype.

Example: message with 5 integers

2345	654	96574	-12	7676
------	-----	-------	-----	------

Point-to-Point Communication

- Communication between two processes.
- Source process sends message to destination process.
- Communication takes place within a communicator, e.g., MPI_COMM_WORLD.
- Processes are identified by their ranks in the communicator.



Point to Point Communication

- Communication is done using send and receive operations among processes.
 - To send a message, sender provides the rank of the process and a unique *tag* to identify the message.
 - The receiver can then receive a message with a given tag (or it may not even care about the tag), and then handle the data accordingly.
 - Two basic (and simple) functions, MPI_Send and MPI_Recv

Data Communication in MPI

- Communication requires the following information:
 - Sender has to know:
 - Whom to send the data to (receiver's process rank)
 - What kind of data to send (100 integers or 200 characters, etc)
 - A user-defined “tag” for the message (think of it as an email subject; allows the receiver to distinguish different messages)
 - Receiver “might” have to know:
 - Who is sending the data (OK if the receiver does not know; in this case sender rank will be `MPI_ANY_SOURCE`, meaning anyone can send)
 - What kind of data is being received (partial information is OK: I might receive *up to* 1000 integers)
 - What the user-defined “tag” of the message is (OK if the receiver does not know; in this case tag will be `MPI_ANY_TAG`)

MPI Send

```
MPI_Send( void* data, int count, MPI_Datatype datatype, int dest,  
          int tag, MPI_Comm comm)
```

- The target process is specified by `dest` and `comm`.
- The data, its type and amount is described by (`data`, `count`, `datatype`).
- `tag` is a user-defined “identifier” for the message
- When this function returns, the data has been delivered to the system and the buffer can be reused.
 - The message may not have been received by the target process.

MPI Receive

```
MPI_Recv( void* data, int count, MPI_Datatype datatype, int  
         source, int tag, MPI_Comm comm, MPI_Status* status)
```

- The sending process is specified by `source` and `comm`.
- The receiving data buffer and its type is defined by `data` and `datatype`.
- Receiving fewer than `count` occurrences of `datatype` is OK, but receiving more is an error.
- Waits until a matching (on `source`, `tag`, `comm`) message is received from the system, and the buffer can be used.

Elementary MPI datatypes

- MPI Datatype is very similar to a C datatype
 - `int` → `MPI_INT`
 - `double` → `MPI_DOUBLE`
 - `char` → `MPI_CHAR`
- More complex datatypes are also possible:
 - E.g., you can create a structure datatype that comprises of other datatypes → a char, an int and a double.
- The “count” in `MPI_SEND` and `MPI_RECV` refers to how many datatype elements should be communicated

Simple Communication in MPI

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char ** argv)
{
    int rank;
    char data[] = "Hello World";

    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if (rank == 0)
        MPI_Send(data, 11, MPI_CHAR, 1, 0, MPI_COMM_WORLD);
    else if (rank == 1)
        MPI_Recv(data, 11, MPI_CHAR, 0, 0, MPI_COMM_WORLD,
                 MPI_STATUS_IGNORE);
    printf("I am a slave and I received %s message
    from my master", data);

    MPI_Finalize();
    return 0;
}
```

So far so good but ...

- Let's review the **MPI_RECV** function

MPI_RECV(buf, count, datatype, source, tag, comm, status)

- The receiver may not know the rank of source, the tag the receiver used and the *count* of elements sent by receiver.
- What can the receiver do?

Wildcards


- Receiver can wildcard.
- To receive from any source — source = MPI_ANY_SOURCE
- To receive from any tag — tag = MPI_ANY_TAG
- The count variable is the upper-bound on the number of elements received, it doesn't specify *how many were actually received*.

```
MPI_RECV(buf, SOME_MAX_NUMBER, datatype,  
MPI_ANY_SOURCE, MPI_ANY_TAG, comm, status)
```

The MPI_Status structure

- Envelope information is returned from MPI_RECV in *status*.

status.MPI_SOURCE
status.MPI_TAG
count via MPI_Get_count()



The MPI_Status structure

- **status** contains further information:
 - Who sent the message (can be used if you used **MPI_ANY_SOURCE**)
 - How much data was actually received
 - What tag was used with the message (can be used if you used **MPI_ANY_TAG**)
 - **MPI_STATUS_IGNORE** can be used if we don't need any additional information

The MPI_Status structure

- The MPI_Recv operation takes the address of an MPI_Status structure as an argument (which can be ignored with MPI_STATUS_IGNORE).
- If we pass an MPI_Status structure to the MPI_Recv function, it will be populated with additional information about the receive operation after it completes.

The MPI_Status structure

- The three primary pieces of information include:
 - **The rank of the sender.** The rank of the sender is stored in the MPI_SOURCE element of the structure. That is, if we declare an MPI_Status stat variable, the rank can be accessed with stat.MPI_SOURCE.
 - **The tag of the message.** The tag of the message can be accessed by the MPI_TAG element of the structure (similar to MPI_SOURCE).
 - **The length of the message.** The length of the message does not have a predefined element in the status structure. Instead, we have to find out the length of the message with MPI_Get_count.

The MPI_Status structure

```
MPI_Get_count(  
    MPI_Status* status,  
    MPI_Datatype datatype,  
    int* count)
```

- In MPI_Get_count, the user passes the MPI_Status structure, the datatype of the message, and count is returned. The count variable is the total number of datatype elements that were received.

```
const int MAX_NUMBERS = 100;
int numbers[MAX_NUMBERS];
int number_amount;
if (world_rank == 0) {
    // Pick a random amount of integers to send to process one
    srand(time(NULL));
    number_amount = (rand() / (float)RAND_MAX) * MAX_NUMBERS;

    // Send the amount of integers to process one
    MPI_Send(numbers, number_amount, MPI_INT, 1, 0, MPI_COMM_WORLD);
    printf("0 sent %d numbers to 1\n", number_amount);
} else if (world_rank == 1) {
    MPI_Status status;
    // Receive at most MAX_NUMBERS from process zero
    MPI_Recv(numbers, MAX_NUMBERS, MPI_INT, 0, 0, MPI_COMM_WORLD,
              &status);

    // After receiving the message, check the status to determine
    // how many numbers were actually received
    MPI_Get_count(&status, MPI_INT, &number_amount);

    // Print off the amount of numbers, and also print additional
    // information in the status object
    printf("1 received %d numbers from 0. Message source = %d, "
           "tag = %d\n",
           number_amount, status.MPI_SOURCE, status.MPI_TAG);
}
```